

xv6 is a re-implementation of Dennis Ritchie's and Ken Thompson's Unix Version 6 (v6). xv6 loosely follows the structure and style of v6, but is implemented for a modern x86-based multiprocessor using ANSI C.

ACKNOWLEDGMENTS

xv6 is inspired by John Lions's Commentary on UNIX 6th Edition (Peer to Peer Communications; ISBN: 1-57398-013-7; 1st edition (June 14, 2000)). See also <http://pdos.csail.mit.edu/6.828/2014/xv6.html>, which provides pointers to on-line resources for v6.

xv6 borrows code from the following sources:

JOS (asm.h, elf.h, mmu.h, bootasm.S, ide.c, console.c, and others)
Plan 9 (entryother.S, mp.h, mp.c, lapic.c)
FreeBSD (ioapic.c)
NetBSD (console.c)

The following people have made contributions:

Russ Cox (context switching, locking)
Cliff Frey (MP)
Xiao Yu (MP)
Nickolai Zeldovich
Austin Clements

In addition, we are grateful for the bug reports and patches contributed by Silas Boyd-Wickizer, Peter Froehlich, Shivam Handa, Anders Kaseorg, Eddie Kohler, Yandong Mao, Hitoshi Mitake, Carmi Merimovich, Joel Nider, Greg Price, Eldar Sehayek, Yongming Shen, Stephen Tu, and Zouchangwei.

The code in the files that constitute xv6 is
Copyright 2006-2014 Frans Kaashoek, Robert Morris, and Russ Cox.

ERROR REPORTS

If you spot errors or have suggestions for improvement, please send email to Frans Kaashoek and Robert Morris (kaashoek,rtm@csail.mit.edu).

BUILDING AND RUNNING XV6

To build xv6 on an x86 ELF machine (like Linux or FreeBSD), run "make". On non-x86 or non-ELF machines (like OS X, even on x86), you will need to install a cross-compiler gcc suite capable of producing x86 ELF binaries. See <http://pdos.csail.mit.edu/6.828/2014/tools.html>. Then run "make TOOLPREFIX=i386-jos-elf-".

To run xv6, install the QEMU PC simulators. To run in QEMU, run "make qemu".

To create a typeset version of the code, run "make xv6.pdf". This requires the "mpage" utility. See <http://www.mesa.nl/pub/mpage/>.

The numbers to the left of the file names in the table are sheet numbers. The source code has been printed in a double column format with fifty lines per column, giving one hundred lines per sheet (or page). Thus there is a convenient relationship between line numbers and sheet numbers.

# basic headers	32 vectors.pl	68 mp.h
01 types.h	33 trapasm.S	69 mp.c
01 param.h	33 trap.c	71 lapic.c
02 memlayout.h	35 syscall.h	73 ioapic.c
02 date.h	35 syscall.c	74 picirq.c
03 defs.h	38 sysproc.c	75 kbd.h
05 x86.h	40 halt.c	77 kbd.c
07 asm.h		77 console.c
07 mmu.h	# file system	80 timer.c
10 elf.h	41 buf.h	81 uart.c
	41 fcntl.h	
# entering xv6	42 stat.h	# user-level
10 entry.S	42 fs.h	82 initcode.S
11 entryother.S	43 file.h	82 usys.S
12 main.c	44 ide.c	83 init.c
	46 bio.c	83 sh.c
# locks	47 log.c	
14 spinlock.h	50 fs.c	# bootloader
14 spinlock.c	57 file.c	90 bootasm.S
	59 sysfile.c	91 bootmain.c
# processes	63 exec.c	
16 vm.c		# add student files her
20 proc.h	# pipes	92 print_mode.c
22 proc.c	65 pipe.c	93 date.c
30 swtch.S		93 uproc.h
31 kalloc.c	# string operations	94 testgiduid.c
	66 string.c	94 ps.c
# system calls		95 time.c
32 traps.h	# low-level hardware	96 testSched.c

The source listing is preceded by a cross-reference that lists every defined constant, struct, global variable, and function in xv6. Each entry gives, on the same line as the name, the line number (or, in a few cases, numbers) where the name is defined. Successive lines in an entry list the line numbers where the name is used. For example, this entry:

```
swtch 2658
      0374 2428 2466 2657 2658
```

indicates that swtch is defined on line 2658 and is mentioned on five lines on sheets 03, 24, and 26.

```

acquire 1474
  0432 1474 1478 2296 2323
  2330 2346 2367 2384 2410
  2444 2522 2563 2591 2651
  2663 2692 2739 2776 2804
  2849 2864 2891 2908 3019
  3034 3175 3192 3400 3872
  3892 4507 4546 4665 4731
  4880 4907 4924 4981 5229
  5265 5282 5311 5327 5337
  5729 5754 5768 6563 6583
  6605 7810 7944 7990 8026
allocproc 2317
  2317 2424 2486
allocuvmm 1853
  0478 1853 1867 2465 6396
  6408
alltraps 3304
  3259 3267 3280 3285 3303
  3304
ALT 7560
  7560 7588 7590
argfd 5919
  5919 5956 5971 5983 5994
  6006
argint 3595
  0451 3595 3608 3624 3834
  3856 3870 3926 3939 3981
  4008 4010 5924 5971 5983
  6187 6260 6261 6307
argptr 3604
  0452 3604 3914 3983 5971
  5983 6006 6333
argstr 3621
  0453 3621 6018 6085 6187
  6236 6259 6278 6307
__attribute__ 1360
  0324 0414 1259 1360
BACK 8362
  8362 8477 8733 8989
backcmd 8400 8727
  8400 8414 8478 8727 8729
  8842 8955 8990
BACKSPACE 7873
  7873 7890 7922 7954 7960
ballocc 5054
  5054 5074 5375 5383 5387
BBLOCK 4310
  4310 5061 5085
B_BUSY 4109
  4109 4539 4671 4672 4685
  4688 4717 4728 4740
  B_DIRTY 4111
    4111 4493 4516 4521 4541
    4561 4685 4719 4989
  begin_op 4878
    0388 2558 4878 5783 5857
    6021 6088 6190 6235 6258
    6277 6370
  bfree 5079
    5079 5414 5424 5427
  bget 4661
    4661 4693 4706
  binit 4639
    0316 1281 4639
  bmap 5368
    5172 5368 5394 5469 5496
  bootmain 9117
    9063 9117
  BPB 4307
    4307 4310 5060 5062 5086
  bread 4702
    0317 4702 4827 4828 4840
    4856 4938 4939 5032 5043
    5061 5085 5188 5209 5289
    5384 5420 5469 5496
  brelse 4726
    0318 4726 4729 4831 4832
    4847 4864 4942 4943 5034
    5046 5067 5072 5092 5194
    5197 5218 5297 5390 5426
    5472 5500
  BSIZE 4255
    4107 4255 4273 4301 4307
    4481 4495 4517 4808 4829
    4940 5044 5469 5470 5471
    5492 5496 5497 5498
  BUDGET 2066
    2066 2301 2526 2699 2711
    2790
  buf 4100
    0300 0317 0318 0319 0360
    0387 2006 2009 2018 2020
    4100 4104 4105 4106 4412
    4428 4431 4475 4504 4535
    4537 4540 4627 4631 4635
    4641 4648 4660 4663 4701
    4704 4715 4726 4755 4827
    4828 4840 4841 4847 4856
    4857 4863 4864 4938 4939

```

```

  4972 5019 5030 5041 5057
  5081 5184 5206 5276 5371
  5409 5455 5482 7779 7790
  7794 7797 7931 7952 7966
  8000 8021 8028 8487 8490
  8491 8492 8606 8618 8620
  8623 8624 8625 8629 8630
  8635
  B_VALID 4110
    4110 4520 4541 4561 4707
  bwrite 4715
    0319 4715 4718 4830 4863
    4941
  bzero 5039
    5039 5068
  C 7581 7937
    7581 7629 7654 7655 7656
    7657 7658 7660 7937 7947
    7950 7957 7968 8001
  CAPSLOCK 7562
    7562 7595 7736
  cgaputc 7878
    7878 7926
  clearpteu 1929
    0487 1929 1935 6410
  cli 0607
    0607 0609 1176 1560 7860
    7917 9012
  cmd 8366
    8366 8378 8387 8388 8393
    8394 8402 8407 8411 8420
    8423 8428 8436 8442 8446
    8454 8478 8480 8569 8581
    8585 8586 8663 8666 8668
    8669 8670 8671 8674 8675
    8677 8679 8680 8681 8682
    8683 8684 8685 8686 8687
    8700 8701 8703 8705 8706
    8707 8708 8709 8710 8713
    8714 8716 8718 8719 8720
    8721 8722 8723 8726 8727
    8729 8731 8732 8733 8734
    8735 8812 8813 8814 8815
    8817 8821 8824 8830 8831
    8834 8837 8839 8842 8846
    8848 8850 8853 8855 8858
    8860 8863 8864 8875 8878
    8881 8885 8900 8903 8908
    8912 8913 8916 8921 8922
    8928 8937 8938 8944 8945
  8951 8952 8961 8964 8966
  8972 8973 8978 8984 8990
  8991 8994
  CMOS_PORT 7235
    7235 7249 7250 7288
  CMOS_RETURN 7236
    7236 7291
  CMOS_STATA 7275
    7275 7323
  CMOS_STATB 7276
    7276 7316
  CMOS_UIP 7277
    7277 7323
  COM1 8113
    8113 8123 8126 8127 8128
    8129 8130 8131 8134 8140
    8141 8157 8159 8167 8169
  commit 4951
    4803 4923 4951
  CONSOLE 4387
    4387 8040 8041
  consoleinit 8036
    0321 1277 8036
  consoleintr 7940
    0323 7748 7940 8175
  consoleread 7983
    7983 8041
  consolewrite 8021
    8021 8040
  consputc 7914
    7766 7797 7818 7836 7839
    7843 7844 7914 7954 7960
    7967 8028
  context 2110
    0301 0429 2071 2110 2129
    2362 2363 2364 2365 2667
    2743 2795 2990
  CONV 7332
    7332 7333 7334 7335 7336
    7337 7338 7339
  copyout 2004
    0486 2004 6418 6429
  copyuvm 1953
    0483 1953 1964 1966 2490
  countForever 9611
    9611 9639 9643
  cprintf 7802
    0322 1274 1314 1867 2292
    2305 2706 2939 2941 2942
    2943 2945 2947 2949 2971

```

```

2974 2976 2987 2992 2994
3424 3432 3437 3762 3765
3902 5172 7019 7039 7211
7412 7802 7862 7863 7864
7867
cpu 2069
0363 1274 1314 1316 1328
1406 1466 1487 1508 1546
1561 1562 1570 1572 1618
1631 1637 1776 1777 1778
1779 2069 2079 2083 2094
2667 2743 2768 2774 2795
2796 3399 3424 3425 3432
3433 3437 3439 6913 6914
7211 7862
cpunum 7201
0378 1338 1624 7201 7423
7432
CR0_PE 0777
0777 1185 1209 9043
CR0_PG 0787
0787 1100 1209
CR0_WP 0783
0783 1100 1209
CR4_PSE 0789
0789 1093 1202
create 6135
6135 6155 6168 6172 6193
6236 6262
CRTPORT 7874
7874 7883 7884 7885 7886
7906 7907 7908 7909
CTL 7559
7559 7585 7589 7735
DAY 7282
7282 7305
deallocuvn 1882
0479 1868 1882 1916 2468
DEVSPACE 0204
0204 1732 1745
devsw 4380
4380 4385 5458 5460 5485
5487 5711 8040 8041
dinode 4277
4277 4301 5185 5189 5207
5210 5277 5290
dirent 4315
4315 5524 5555 6066 6081
dirlink 5552
0340 5531 5552 5567 5575

```

```

6041 6167 6171 6172
dirlookup 5521
0341 5521 5527 5559 5644
6100 6145
DIRSIZ 4313
4313 4317 5515 5572 5608
5609 5661 6015 6082 6139
dobuiltin 8581
8581 8630
DPL_USER 0829
0829 1627 1628 2431 2432
3373 3447 3456
EOESC 7566
7566 7720 7724 7725 7727
7730
elfhdr 1005
1005 6365 9119 9124
ELF_MAGIC 1002
1002 6381 9130
ELF_PROG_LOAD 1036
1036 6392
end_op 4903
0389 2560 4903 5785 5862
6023 6030 6048 6057 6090
6124 6130 6195 6200 6206
6215 6219 6237 6241 6263
6267 6279 6285 6290 6372
6402 6455
entry 1090
1011 1086 1089 1090 3252
3253 6442 6821 9121 9145
9146
EOI 7116
7116 7186 7225
ERROR 7137
7137 7179
ESR 7119
7119 7182 7183
exec 6360
0327 3709 6323 6360 8268
8329 8330 8431 8432 9570
EXEC 8358
8358 8427 8670 8965
execcmd 8370 8664
8370 8415 8428 8664 8666
8921 8927 8928 8956 8966
exit 2538
0408 2538 2580 3389 3393
3448 3457 3704 3819 8216
8219 8261 8326 8331 8421

```

```

8430 8440 8483 8638 8645
9311 9315 9435 9442 9466
9498 9505 9572 9577 9584
9591 9644
EXTMEM 0202
0202 0208 1729
fdalloc 5938
5938 5958 6211 6338
fetchint 3567
0454 3567 3597 6314
fetchstr 3579
0455 3579 3626 6320
file 4350
0302 0330 0331 0332 0334
0335 0336 0401 2132 4350
5020 5708 5714 5724 5727
5730 5751 5752 5764 5766
5802 5815 5835 5913 5919
5922 5938 5953 5967 5979
5992 6003 6184 6330 6506
6521 7760 8108 8379 8438
8439 8675 8683 8872
filealloc 5725
0330 5725 6211 6527
fileclose 5764
0331 2553 5764 5770 5997
6213 6341 6342 6554 6556
filedup 5752
0332 2514 5752 5756 5960
fileinit 5718
0333 1282 5718
fileread 5815
0334 5815 5830 5973
filestat 5802
0335 5802 6008
filewrite 5835
0336 5835 5867 5872 5985
FL_IF 0760
0760 1562 1568 2435 2772
7208
fork 2480
0409 2480 3703 3813 8260
8323 8325 8655 8657 9563
9637
fork1 8651
8405 8447 8457 8464 8479
8634 8651
forkret 2813
2223 2365 2813
freerange 3151

```

```

3111 3134 3140 3151
freevm 1910
0480 1910 1915 1978 2604
6445 6452
FSSIZE 0162
0162 4479
gatedesc 0951
0573 0576 0951 3361
getbuiltin 8551
8551 8576
getcallerpcs 1526
0433 1488 1526 2990 7865
getcmd 8487
8487 8618
getprocs 3003
0422 3003 3732 3985 8288
9463
gettoken 8756
8756 8841 8845 8857 8870
8871 8907 8911 8933
growproc 2459
0410 2459 3859
havedisk1 4430
4430 4464 4543
holding 1544
0434 1477 1504 1544 2240
2254 2266 2382 2766
HOURS 7281
7281 7304
ialloc 5181
0342 5181 5199 6154 6155
IBLOCK 4304
4304 5188 5209 5289
I_BUSY 4375
4375 5283 5285 5308 5312
5330 5332
ICRHI 7130
7130 7189 7257 7269
ICRLO 7120
7120 7190 7191 7258 7260
7270
ID 7113
7113 7149 7216
IDE_BSY 4415
4415 4439
IDE_CMD_READ 4420
4420 4497
IDE_CMD_WRITE 4421
4421 4494
IDE_DF 4417

```

```

4417 4441
IDE_DRDY 4416
4416 4439
IDE_ERR 4418
4418 4441
ideinit 4451
0358 1283 4451
ideintr 4502
0359 3408 4502
idelock 4427
4427 4455 4507 4509 4528
4546 4562 4565
iderw 4535
0360 4535 4540 4542 4544
4708 4720
idestart 4475
4431 4475 4478 4484 4526
4558
idewait 4435
4435 4458 4486 4516
idtinit 3379
0462 1315 3379
idup 5263
0343 2515 5263 5631
iget 5225
5176 5195 5225 5245 5539
5629
iinit 5168
0344 2824 5168
ilock 5274
0345 5274 5280 5300 5634
5805 5824 5858 6027 6040
6053 6094 6102 6143 6147
6157 6203 6282 6375 7995
8015 8030
inb 0503
0503 4439 4463 7054 7291
7714 7717 7884 7886 8134
8140 8141 8157 8167 8169
9023 9031 9154
INITGID 2055
2055 2452
initlock 1462
0435 1462 2231 3132 3375
4455 4643 4812 5170 5720
6535 8038
initlog 4806
0386 2825 4806 4809
INITUID 2054
2054 2451
inituvm 1803
0481 1803 1808 2428
inode 4362
0303 0340 0341 0342 0343
0345 0346 0347 0348 0349
0351 0352 0353 0354 0355
0482 1818 2133 4356 4362
4381 4382 5023 5164 5176
5180 5204 5224 5227 5233
5262 5263 5274 5306 5325
5352 5368 5406 5437 5452
5479 5520 5521 5552 5556
5623 5626 5658 5665 6016
6063 6080 6134 6138 6185
6233 6253 6275 6366 7983
8021
INPUT_BUF 7929
7929 7931 7952 7964 7966
7968 8000
insl 0512
0512 0514 4517 9173
install_trans 4822
4822 4871 4956
INT_DISABLED 7369
7369 7417
ioapic 7377
7007 7029 7030 7374 7377
7386 7387 7393 7394 7408
IOAPIC 7358
7358 7408
ioapicenable 7423
0363 4457 7423 8045 8143
ioapicid 6917
0364 6917 7030 7047 7411
7412
ioapicinit 7401
0365 1276 7401 7412
ioapicread 7384
7384 7409 7410
ioapicwrite 7391
7391 7417 7418 7431 7432
IO_PIC1 7457
7457 7470 7485 7494 7497
7502 7512 7526 7527
IO_PIC2 7458
7458 7471 7486 7515 7516
7517 7520 7529 7530
IO_TIMER1 8059
8059 8068 8078 8079
IPB 4301

```

```

4301 4304 5189 5210 5290
iput 5325
0346 2559 5325 5331 5355
5560 5652 5784 6046 6289
IRQ_COM1 3233
3233 3418 8142 8143
IRQ_ERROR 3235
3235 7179
IRQ_IDE 3234
3234 3407 3411 4456 4457
IRQ_KBD 3232
3232 3414 8044 8045
IRQ_SLAVE 7460
7460 7464 7502 7517
IRQ_SPURIOUS 3236
3236 3423 7159
IRQ_TIMER 3231
3231 3398 3452 7166 8080
isdirempty 6063
6063 6070 6106
ismp 6915
0392 1284 6915 7012 7020
7040 7043 7405 7425
itrunc 5406
5023 5334 5406
iunlock 5306
0347 5306 5309 5354 5641
5807 5827 5861 6036 6218
6288 7988 8025
iunlockput 5352
0348 5352 5636 5645 5648
6029 6042 6045 6056 6107
6118 6122 6129 6146 6150
6174 6205 6214 6240 6266
6284 6401 6454
iupdate 5204
0349 5204 5336 5432 5505
6035 6055 6116 6121 6161
6165
I_INVALID 4376
4376 5288 5298 5328
kalloc 3187
0368 1344 1663 1742 1809
1865 1969 2344 3187 6529
KBDATAP 7554
7554 7717
kbdgetc 7706
7706 7748
kbdintr 7746
0374 3415 7746
KBS_DIB 7553
7553 7715
KBSTATP 7552
7552 7714
KERNBASE 0207
0207 0208 0212 0213 0217
0218 0220 0221 1365 1533
1729 1858 1916
KERNLINK 0208
0208 1730
KEY_DEL 7578
7578 7619 7641 7665
KEY_DN 7572
7572 7615 7637 7661
KEY_END 7570
7570 7618 7640 7664
KEY_HOME 7569
7569 7618 7640 7664
KEY_INS 7577
7577 7619 7641 7665
KEY_LF 7573
7573 7617 7639 7663
KEY_PGDN 7576
7576 7616 7638 7662
KEY_PGUP 7575
7575 7616 7638 7662
KEY_RT 7574
7574 7617 7639 7663
KEY_UP 7571
7571 7615 7637 7661
kfree 3164
0369 1898 1900 1920 1923
2491 2602 3156 3164 3169
6552 6573
kill 2904
0411 2904 3438 3708 3836
8267
kinit1 3130
0370 1269 3130
kinit2 3138
0371 1287 3138
KSTACKSIZE 0151
0151 1104 1113 1345 1779
2351
kvmalloc 1757
0474 1270 1757
lapiceoi 7222
0380 3405 3409 3416 3420
3426 7222
lapicinit 7153

```

```

0381 1272 1306 7153
lapicstartap 7241
0382 1349 7241
lapicw 7146
7146 7159 7165 7166 7167
7170 7171 7176 7179 7182
7183 7186 7189 7190 7195
7225 7257 7258 7260 7269
7270
lcr3 0640
0640 1768 1783
lgdt 0562
0562 0570 1183 1633 9041
lidt 0576
0576 0584 3381
LINT0 7135
7135 7170
LINT1 7136
7136 7171
LIST 8361
8361 8445 8720 8983
listcmd 8391 8714
8391 8416 8446 8714 8716
8846 8957 8984
loadgs 0601
0601 1634
loadvm 1818
0482 1818 1824 1827 6398
log 4787 4800
4787 4800 4812 4814 4815
4816 4826 4827 4828 4840
4843 4844 4845 4856 4859
4860 4861 4872 4880 4882
4883 4884 4886 4888 4889
4907 4908 4909 4910 4911
4913 4916 4918 4924 4925
4926 4927 4937 4938 4939
4953 4957 4976 4978 4981
4982 4983 4986 4987 4988
4990
logheader 4782
4782 4794 4808 4809 4841
4857
LOGSIZE 0160
0160 4784 4884 4976 5850
log_write 4972
0387 4972 4979 5045 5066
5091 5193 5217 5388 5499
ltr 0588
0588 0590 1780

```

```

makeint 8514
8514 8535 8541
mappages 1679
1679 1748 1811 1872 1972
MAXARG 0158
0158 6303 6364 6415
MAXARGS 8364
8364 8372 8373 8940
MAXFILE 4274
4274 5492
MAXOPBLOCKS 0159
0159 0160 0161 4884
memcmp 6665
0441 6665 6945 6988 7326
memmove 6681
0442 1335 1812 1971 2018
4829 4940 5033 5216 5296
5471 5498 5609 5611 6681
6704 7901
memset 6654
0443 1666 1744 1810 1871
2364 2430 3172 5044 5191
6111 6310 6654 7903 8490
8669 8680 8706 8719 8732
microdelay 7231
0383 7231 7259 7261 7271
7289 8158
min 5022
5022 5470 5497
MINS 7280
7280 7303
MONTH 7283
7283 7306
mp 6802
6802 6908 6937 6944 6945
6946 6955 6960 6964 6965
6968 6969 6980 6983 6985
6987 6994 7004 7010 7050
mpbcpu 6920
0393 6920
MPBUS 6852
6852 7033
mpconf 6813
6813 6979 6982 6987 7005
mpconfig 6980
6980 7010
mpenter 1302
1302 1346
mpinit 7001
0394 1271 7001 7019 7039

```

```

mpioapic 6839
6839 7007 7029 7031
MPIOAPIC 6853
6853 7028
MPIOINTR 6854
6854 7034
MPLINTR 6855
6855 7035
mpmain 1312
1259 1290 1307 1312
mpproc 6828
6828 7006 7017 7026
MPPROC 6851
6851 7016
mpsearch 6956
6956 6985
mpsearch1 6938
6938 6964 6968 6971
multiboot_header 1075
1074 1075
namecmp 5513
0350 5513 5534 6097
namei 5659
0351 2440 5659 6022 6199
6278 6371
nameiparent 5666
0352 5624 5639 5651 5666
6038 6089 6141
namex 5624
5624 5662 5668
NBUF 0161
0161 4631 4648
ncpu 6916
1274 1337 2084 4457 6916
7018 7019 7023 7024 7025
7045
NCPU 0152
0152 2083 6913
NDEV 0156
0156 5458 5485 5711
NDIRECT 4272
4272 4274 4283 4373 5373
5378 5382 5383 5412 5419
5420 5427 5428
NELEM 0490
0490 1747 2983 3758 6312
nextpid 2222
2222 2340
NFILE 0154
0154 5714 5730

```

```

NINDIRECT 4273
4273 4274 5380 5422
NINODE 0155
0155 5164 5233
NO 7556
7556 7602 7605 7607 7608
7609 7610 7612 7624 7627
7629 7630 7631 7632 7634
7652 7653 7655 7656 7657
7658
NOFILE 0153
0153 2132 2512 2551 5926
5942
NPENTRIES 0871
0871 1361 1917
NPROC 0150
0150 2212 2298 2324 2413
2569 2595 2652 2695 2875
2909 2980 3017 3018 3020
NPENTRIES 0872
0872 1894
NSEGS 2051
1611 2051 2073
nulterminate 8952
8815 8830 8952 8973 8979
8980 8985 8986 8991
numChildren 9608
9608 9636
NUMLOCK 7563
7563 7596
NUM_READY_LISTS 2058
2058 2062 2214 2418
O_CREATE 4153
4153 6192 8878 8881
O_RDONLY 4150
4150 6204 8875
O_RDWR 4152
4152 6225 8314 8316 8610
outb 0521
0521 4461 4470 4487 4488
4489 4490 4491 4492 4494
4497 7053 7054 7249 7250
7288 7470 7471 7485 7486
7494 7497 7502 7512 7515
7516 7517 7520 7526 7527
7529 7530 7883 7885 7906
7907 7908 7909 8077 8078
8079 8123 8126 8127 8128
8129 8130 8131 8159 9028
9036 9164 9165 9166 9167

```

```

    9168 9169
outs1 0533
    0533 0535 4495
outw 0527
    0527 1219 1221 3904 9069
    9071
O_WRONGLY 4151
    4151 6224 6225 8878 8881
P2V 0218
    0218 1269 1287 6962 7251
    7875
panic 7855 8642
    0324 1478 1505 1569 1571
    1690 1746 1782 1808 1824
    1827 1898 1915 1935 1964
    1966 2241 2255 2267 2290
    2383 2427 2544 2580 2729
    2767 2769 2771 2773 2837
    2840 3169 3434 4478 4480
    4484 4540 4542 4544 4693
    4718 4729 4809 4910 4977
    4979 5074 5089 5199 5245
    5280 5300 5309 5331 5394
    5527 5531 5567 5575 5756
    5770 5830 5867 5872 6070
    6105 6113 6155 6168 6172
    7813 7855 7862 7896 8406
    8425 8456 8642 8657 8828
    8872 8906 8910 8936 8941
panicked 7768
    7768 7868 7916
parseblock 8901
    8901 8906 8925
parsecmd 8818
    8407 8635 8818
parseexec 8917
    8814 8855 8917
parseline 8835
    8812 8824 8835 8846 8908
parsepipe 8851
    8813 8839 8851 8858
parseredirs 8864
    8864 8912 8931 8942
PCINT 7134
    7134 7176
pde_t 0103
    0103 0476 0477 0478 0479
    0480 0481 0482 0483 0486
    0487 1260 1320 1361 1610
    1654 1656 1679 1736 1739
    1742 1803 1818 1853 1882
    1910 1929 1952 1953 1955
    1984 2004 2123 6368
PDX 0862
    0862 1659
PDXSHIFT 0877
    0862 0868 0877 1365
peek 8801
    8801 8825 8840 8844 8856
    8869 8905 8909 8924 8932
PGROUNDDOWN 0880
    0880 1684 1685 2011
PGROUNDUP 0879
    0879 1863 1890 3154 6407
PGSIZE 0873
    0873 0879 0880 1360 1666
    1694 1695 1744 1807 1810
    1811 1823 1825 1829 1832
    1864 1871 1872 1891 1894
    1962 1971 1972 2015 2021
    2429 2436 3155 3168 3172
    6408 6410
PHYSTOP 0203
    0203 1287 1731 1745 1746
    3168
picenable 7475
    0398 4456 7475 8044 8080
    8142
picinit 7482
    0399 1275 7482
picsetmask 7467
    7467 7477 7533
pinit 2229
    0412 1279 2229
pipe 6511
    0304 0402 0403 0404 3706
    4355 5781 5822 5842 6511
    6523 6529 6535 6539 6543
    6561 6579 6601 8263 8455
    8456
PIPE 8360
    8360 8453 8707 8977
pipealloc 6521
    0401 6335 6521
pipeclose 6561
    0402 5781 6561
pipecmd 8385 8701
    8385 8417 8454 8701 8703
    8858 8958 8978
piperead 6601

```

```

    0403 5822 6601
PIPESIZE 6509
    6509 6513 6585 6593 6616
pipewrite 6579
    0404 5842 6579
popcli 1566
    0438 1521 1566 1569 1571
    1784
popq 2238
    2238 2241 2331 2708 2726
print_elapsed 2935
    2935 2988
printint 7776
    7776 7826 7830
PrioCount 9607
    9607 9617 9624
PRIORITY_HIGH 2060
    2060 2062 2291 2292 2418
    2445 2525 2527 2696 2705
    2721
PRIORITY_LOW 2062
    2062 2291 2292 2705 2721
    2788
proc 2121
    0305 0407 0484 1255 1458
    1606 1638 1773 1779 2080
    2095 2121 2127 2146 2206
    2212 2214 2215 2220 2237
    2238 2243 2252 2264 2275
    2297 2298 2316 2319 2324
    2406 2413 2463 2465 2468
    2471 2472 2483 2490 2499
    2500 2501 2505 2506 2513
    2514 2515 2517 2540 2543
    2552 2553 2554 2559 2561
    2566 2569 2570 2578 2588
    2595 2596 2619 2625 2641
    2652 2659 2667 2672 2684
    2695 2735 2743 2752 2770
    2777 2779 2788 2789 2790
    2792 2793 2795 2805 2836
    2854 2855 2859 2873 2875
    2906 2909 2935 2966 2980
    3016 3020 3355 3388 3390
    3392 3430 3438 3439 3441
    3447 3452 3456 3555 3569
    3583 3586 3597 3610 3757
    3759 3762 3766 3767 3807
    3842 3858 3875 3930 3943
    3954 3961 3968 3969 3970
    3971 3972 4407 5016 5631
    5911 5926 5943 5944 5996
    6289 6291 6340 6354 6436
    6439 6440 6441 6442 6443
    6444 6504 6586 6607 6911
    7006 7017 7018 7019 7022
    7763 7993 8110
procdump 2955
    0413 2955 7978
proghdr 1024
    1024 6367 9120 9134
PTE_ADDR 0894
    0894 1661 1828 1896 1919
    1967 1993
PTE_FLAGS 0895
    0895 1968
PTE_P 0883
    0883 1363 1365 1660 1670
    1689 1691 1895 1918 1965
    1989
PTE_PS 0890
    0890 1363 1365
pte_t 0898
    0898 1653 1657 1661 1663
    1682 1821 1884 1931 1956
    1986
PTE_U 0885
    0885 1670 1811 1872 1936
    1991
PTE_W 0884
    0884 1363 1365 1670 1729
    1731 1732 1811 1872
PTX 0865
    0865 1672
PTXSHIFT 0876
    0865 0868 0876
pushcli 1555
    0437 1476 1555 1775
pushfreeq 2252
    2252 2255 2347 2415 2495
    2607
pushreadyq 2264
    2264 2267 2527 2712 2793
    2882 2919
rcr2 0632
    0632 3433 3440
readeflags 0594
    0594 1559 1568 2772 7208
read_head 4838
    4838 4870

```

```

readi 5452
  0353 1833 5452 5530 5566
  5825 6069 6070 6379 6390
readsb 5028
  0339 4813 5028 5084 5171
readsect 9160
  9160 9195
readseg 9179
  9114 9127 9138 9179
recover_from_log 4868
  4802 4817 4868
REDIR 8359
  8359 8435 8681 8971
redircmd 8376 8675
  8376 8418 8436 8675 8677
  8875 8878 8881 8959 8972
REG_ID 7360
  7360 7410
REG_TABLE 7362
  7362 7417 7418 7431 7432
REG_VER 7361
  7361 7409
release 1502
  0436 1502 1505 2302 2306
  2327 2334 2341 2348 2369
  2386 2390 2421 2447 2529
  2613 2620 2665 2674 2741
  2754 2781 2807 2817 2850
  2863 2893 2923 2927 3037
  3045 3180 3197 3402 3876
  3881 3894 4509 4528 4565
  4673 4689 4743 4889 4918
  4927 4990 5236 5255 5267
  5286 5314 5333 5342 5733
  5737 5758 5772 5778 6572
  6575 6587 6596 6608 6619
  7851 7976 7994 8014 8029
ROOTDEV 0157
  0157 2824 2825 5629
ROOTINO 4254
  4254 5629
rtcdat 0250
  0250 0306 0377 3913 7300
  7311 7313 9308
run 3114
  2962 3011 3114 3115 3121
  3166 3176 3189
runcmd 8411
  8411 8425 8442 8448 8450
  8462 8469 8480 8635

```

```

RUNNING 2118
  2118 2661 2698 2737 2770
  2962 3011 3452
safestrcpy 6732
  0444 2439 2517 3038 3040
  6436 6732
sb 5024
  0339 4304 4310 4811 4813
  4814 4815 5024 5028 5033
  5060 5061 5062 5084 5085
  5171 5172 5173 5187 5188
  5209 5289 7314 7316 7318
sched 2762
  0415 2579 2762 2767 2769
  2771 2773 2806 2856
scheduler 2639 2682
  0414 1317 2071 2639 2667
  2682 2743 2795
SCROLLLOCK 7564
  7564 7597
SECS 7279
  7279 7302
SECTOR_SIZE 4414
  4414 4481
SECTSIZE 9112
  9112 9173 9186 9189 9194
SEG 0819
  0819 1625 1626 1627 1628
  1631
SEG16 0823
  0823 1776
SEG_ASM 0710
  0710 1228 1229 9079 9080
segdesc 0802
  0559 0562 0802 0819 0823
  1611 2073
seginit 1616
  0473 1273 1305 1616
SEG_KCODE 0791
  0791 1188 1625 3372 3373
  9049
SEG_KCPU 0793
  0793 1631 1634 3316
SEG_KDATA 0792
  0792 1192 1626 1778 3313
  9053
SEG_NULLASM 0704
  0704 1227 9078
SEG_TSS 0796
  0796 1776 1777 1780

```

```

SEG_UCODE 0794
  0794 1627 2431
SEG_UDATA 0795
  0795 1628 2432
setbuiltin 8526
  8526 8575
SETGATE 0971
  0971 3372 3373
setpriority 2287
  0425 2287 3735 4012 8289
  9618 9625
setupkvm 1737
  0476 1737 1759 1960 2426
  6384
SHIFT 7558
  7558 7586 7587 7735
skipelem 5595
  5595 5633
sleep 2834
  0416 2625 2834 2837 2840
  2960 3009 3715 3879 4562
  4676 4883 4886 5284 6591
  6611 7998 8279
spinlock 1401
  0307 0416 0432 0434 0435
  0436 0465 1401 1459 1462
  1474 1502 1544 2207 2211
  2834 3109 3119 3358 3363
  4410 4427 4625 4630 4753
  4788 5017 5163 5709 5713
  6507 6512 7758 7771 8106
STA_R 0719 0836
  0719 0836 1228 1625 1627
  9079
start 1175 8208 9011
  1174 1175 1205 1213 1215
  4789 4814 4827 4840 4856
  4938 5172 8207 8208 9010
  9011 9062 9619
startothers 1324
  1258 1286 1324
stat 4204
  0308 0335 0354 4204 5014
  5437 5802 5909 6004 8303
  9203
stati 5437
  0354 5437 5806
STA_W 0718 0835
  0718 0835 1229 1626 1628
  1631 9080

```

```

STA_X 0715 0832
  0715 0832 1228 1625 1627
  9079
sti 0613
  0613 0615 1573 2645 2688
stosb 0542
  0542 0544 6660 9140
stosl 0551
  0551 0553 6658
strlen 6751
  0445 6417 6418 6751 8530
  8533 8539 8553 8585 8623
  8823
STRMAX 9350
  9350 9359 9361
strncmp 6708 8504
  0446 5515 6708 8504 8531
  8532 8534 8538 8540 8554
  8555 8559 8585
strncpy 6718
  0447 5572 6718
STS_IG32 0850
  0850 0977
STS_T32A 0847
  0847 1776
STS_TG32 0851
  0851 0977
sum 6926
  6926 6928 6930 6932 6933
  6945 6992
superblock 4262
  0309 0339 4262 4811 5024
  5028
SVR 7117
  7117 7159
switchkvm 1766
  0485 1304 1760 1766 2668
  2744
switchvum 1773
  0484 1773 1782 2472 2660
  2736 6444
swtch 3058
  0429 2667 2743 2795 3057
  3058
syscall 3753
  0456 3391 3557 3753
SYSCALL 8253 8260 8261 8262 8263 82
  8260 8261 8262 8263 8264
  8265 8266 8267 8268 8269
  8270 8271 8272 8273 8274

```

```

8275 8276 8277 8278 8279
8280 8281 8282 8283 8284
8285 8286 8287 8288 8289
sys_chdir 6272
3629 3673 6272
SYS_chdir 3509
3509 3510 3673 3711
sys_close 5989
3630 3685 5989
SYS_close 3521
3521 3522 3685 3723
sys_date 3911
3651 3687 3911
SYS_date 3524
3524 3525 3687 3725
sys_dup 5951
3631 3674 5951
SYS_dup 3510
3510 3511 3674 3712
sys_exec 6301
3632 3671 6301
SYS_exec 3507
3507 3508 3671 3709 8212
sys_exit 3817
3633 3666 3817
SYS_exit 3502
3502 3503 3666 3704 8217
sys_fork 3811
3634 3665 3811
SYS_fork 3501
3501 3502 3665 3703
sys_fstat 6001
3635 3672 6001
SYS_fstat 3508
3508 3509 3672 3710
sys_getgid 3959
3654 3690 3959
SYS_getgid 3526
3526 3527 3690 3727
sys_getpid 3840
3636 3675 3840
SYS_getpid 3511
3511 3512 3675 3713
sys_getppid 3966
3655 3691 3966
SYS_getppid 3527
3527 3528 3691 3729
sys_getprocs 3977
3658 3694 3977
SYS_getprocs 3530
3530 3531 3694 3732
sys_getuid 3952
3653 3689 3952
SYS_getuid 3525
3525 3526 3689 3728
SYS_halt 3522
3522 3524 3686 3724
sys_kill 3830
3637 3670 3830
SYS_kill 3506
3506 3507 3670 3708
sys_link 6013
3638 3683 6013
SYS_link 3519
3519 3520 3683 3721
sys_mkdir 6230
3639 3684 6230
SYS_mkdir 3520
3520 3521 3684 3722
sys_mknod 6251
3640 3681 6251
SYS_mknod 3517
3517 3518 3681 3719
sys_open 6180
3641 3679 6180
SYS_open 3515
3515 3516 3679 3717
sys_pipe 6327
3642 3668 6327
SYS_pipe 3504
3504 3505 3668 3706
sys_read 5965
3643 3669 5965
SYS_read 3505
3505 3506 3669 3707
sys_sbrk 3851
3644 3676 3851
SYS_sbrk 3512
3512 3513 3676 3714
sys_setgid 3936
3657 3693 3936
SYS_setgid 3529
3529 3530 3693 3730
sys_setpriority 4004
3661 3697 4004
SYS_setpriority 3531
3531 3697 3735
sys_setuid 3923
3656 3692 3923
SYS_setuid 3528

```

```

3528 3529 3692 3731
sys_sleep 3865
3645 3677 3865
SYS_sleep 3513
3513 3514 3677 3715
sys_unlink 6078
3646 3682 6078
SYS_unlink 3518
3518 3519 3682 3720
sys_uptime 3888
3649 3678 3888
SYS_uptime 3514
3514 3515 3678 3716
sys_wait 3824
3647 3667 3824
SYS_wait 3503
3503 3504 3667 3705
sys_write 5977
3648 3680 5977
SYS_write 3516
3516 3517 3680 3718
taskstate 0901
0901 2072
TDCR 7141
7141 7165
T_DEV 4202
4202 5457 5484 6262 9208
T_DIR 4200
4200 5526 5635 6028 6106
6114 6163 6204 6236 6283
9206
testgiduid 9406
9406 9434
T_FILE 4201
4201 6148 6193 9207
ticks 3364
0463 2368 2385 2389 2664
2740 2777 2779 2938 3035
3364 3401 3403 3873 3874
3879 3893
tickslock 3363
0465 2367 2369 2384 2386
2390 2663 2665 2739 2741
2776 2781 3034 3037 3363
3375 3400 3402 3872 3876
3879 3881 3892 3894
TICKS_TO_PROMOTE 2064
2064 2389
TIICR 7139
7139 7167
TIMER 7131
7131 7166
TIMER_16BIT 8071
8071 8077
TIMER_DIV 8066
8066 8078 8079
TIMER_FREQ 8065
8065 8066
timerinit 8074
0459 1285 8074
TIMER_MODE 8068
8068 8077
TIMER_RATEGEN 8070
8070 8077
TIMER_SEL0 8069
8069 8077
timetopromote 2380
2380 2383 2693
T_IRQ0 3229
3229 3398 3407 3411 3414
3418 3422 3423 3452 7159
7166 7179 7417 7431 7497
7516
TPR 7115
7115 7195
trap 3385
3252 3254 3322 3385 3432
3434 3437
trapframe 0652
0652 2128 2355 3385
trapret 3327
2224 2360 3326 3327
T_SYSCALL 3226
3226 3373 3387 8213 8218
8257
tvinit 3367
0464 1280 3367
uart 8115
8115 8136 8155 8165
uartgetc 8163
8163 8175
uartinit 8118
0468 1278 8118
uartintr 8173
0469 3419 8173
uartputc 8151
0470 7923 7925 8147 8151
uproc 9352
0311 0422 2208 3003 3808
3980 9352 9451 9461 9462

```


userinit 2404	6569 6590 6595 6618 7970
0417 1288 2404 2427	wakeup1 2871
uva2ka 1984	2226 2566 2573 2871 2892
0477 1984 2012	walkpgdir 1654
V2P 0217	1654 1687 1826 1892 1933
0217 1730 1731	1963 1988
V2P_WO 0220	write_head 4854
0220 1086 1096	4854 4873 4955 4958
VER 7114	writei 5479
7114 7175	0355 5479 5574 5859 6112
wait 2586	6113
0418 2586 3705 3826 8262	write_log 4933
8333 8449 8473 8474 8636	4933 4954
9565	xchg 0619
waitdisk 9151	0619 1316 1483 1519
9151 9163 9172	YEAR 7284
wakeup 2889	7284 7307
0419 2889 3403 4522 4741	yield 2802
4916 4926 5313 5339 6566	0420 2802 3453

```
0100 typedef unsigned int    uint;
0101 typedef unsigned short ushort;
0102 typedef unsigned char    uchar;
0103 typedef uint pde_t;
0104
0105
0106
0107
0108
0109
0110
0111
0112
0113
0114
0115
0116
0117
0118
0119
0120
0121
0122
0123
0124
0125
0126
0127
0128
0129
0130
0131
0132
0133
0134
0135
0136
0137
0138
0139
0140
0141
0142
0143
0144
0145
0146
0147
0148
0149
```

```
0150 #define NPROC          64 // maximum number of processes
0151 #define KSTACKSIZE 4096 // size of per-process kernel stack
0152 #define NCPU           8 // maximum number of CPUs
0153 #define NOFILE         16 // open files per process
0154 #define NFILE          100 // open files per system
0155 #define NINODE          50 // maximum number of active i-nodes
0156 #define NDEV           10 // maximum major device number
0157 #define ROOTDEV         1 // device number of file system root disk
0158 #define MAXARG         32 // max exec arguments
0159 #define MAXOPBLOCKS    10 // max # of blocks any FS op writes
0160 #define LOGSIZE        (MAXOPBLOCKS*3) // max data blocks in on-disk log
0161 #define NBUF           (MAXOPBLOCKS*3) // size of disk block cache
0162 #define FSSIZE         1000 // size of file system in blocks
0163
0164
0165
0166
0167
0168
0169
0170
0171
0172
0173
0174
0175
0176
0177
0178
0179
0180
0181
0182
0183
0184
0185
0186
0187
0188
0189
0190
0191
0192
0193
0194
0195
0196
0197
0198
0199
```

```
0200 // Memory layout
0201
0202 #define EXTMEM 0x100000 // Start of extended memory
0203 #define PHYSTOP 0xE000000 // Top physical memory
0204 #define DEVSPACE 0xFE000000 // Other devices are at high addresses
0205
0206 // Key addresses for address space layout (see kmap in vm.c for layout)
0207 #define KERNBASE 0x80000000 // First kernel virtual address
0208 #define KERNLINK (KERNBASE+EXTMEM) // Address where kernel is linked
0209
0210 #ifndef __ASSEMBLER__
0211
0212 static inline uint v2p(void *a) { return ((uint) (a)) - KERNBASE; }
0213 static inline void *p2v(uint a) { return (void *) ((a) + KERNBASE); }
0214
0215 #endif
0216
0217 #define V2P(a) (((uint) (a)) - KERNBASE)
0218 #define P2V(a) (((void *) (a)) + KERNBASE)
0219
0220 #define V2P_WO(x) ((x) - KERNBASE) // same as V2P, but without casts
0221 #define P2V_WO(x) ((x) + KERNBASE) // same as P2V, but without casts
0222
0223
0224
0225
0226
0227
0228
0229
0230
0231
0232
0233
0234
0235
0236
0237
0238
0239
0240
0241
0242
0243
0244
0245
0246
0247
0248
0249
```

```
0250 struct rtcdate {
0251     uint second;
0252     uint minute;
0253     uint hour;
0254     uint day;
0255     uint month;
0256     uint year;
0257 };
0258
0259
0260
0261
0262
0263
0264
0265
0266
0267
0268
0269
0270
0271
0272
0273
0274
0275
0276
0277
0278
0279
0280
0281
0282
0283
0284
0285
0286
0287
0288
0289
0290
0291
0292
0293
0294
0295
0296
0297
0298
0299
```

```

0300 struct buf;
0301 struct context;
0302 struct file;
0303 struct inode;
0304 struct pipe;
0305 struct proc;
0306 struct rtcdate;
0307 struct spinlock;
0308 struct stat;
0309 struct superblock;
0310 #ifdef CS333_P2
0311 struct uproc;
0312 #endif
0313
0314
0315 // bio.c
0316 void          binit(void);
0317 struct buf*   bread(uint, uint);
0318 void          brelse(struct buf*);
0319 void          bwrite(struct buf*);
0320 // console.c
0321 void          consoleinit(void);
0322 void          cprintf(char*, ...);
0323 void          consoleintr(int (*)(void));
0324 void          panic(char*) __attribute__((noreturn));
0325
0326 // exec.c
0327 int           exec(char*, char**);
0328
0329 // file.c
0330 struct file*  filealloc(void);
0331 void          fileclose(struct file*);
0332 struct file*  filedup(struct file*);
0333 void          fileinit(void);
0334 int           fileread(struct file*, char*, int n);
0335 int           filestat(struct file*, struct stat*);
0336 int           filewrite(struct file*, char*, int n);
0337
0338 // fs.c
0339 void          readsb(int dev, struct superblock *sb);
0340 int           dirlink(struct inode*, char*, uint);
0341 struct inode* dirlookup(struct inode*, char*, uint*);
0342 struct inode* ialloc(uint, short);
0343 struct inode* idup(struct inode*);
0344 void          iinit(int dev);
0345 void          ilock(struct inode*);
0346 void          iput(struct inode*);
0347 void          iunlock(struct inode*);
0348 void          iunlockput(struct inode*);
0349 void          iupdate(struct inode*);

```

```

0350 int           namecmp(const char*, const char*);
0351 struct inode* namei(char*);
0352 struct inode* nameiparent(char*, char*);
0353 int           readi(struct inode*, char*, uint, uint);
0354 void          stati(struct inode*, struct stat*);
0355 int           writei(struct inode*, char*, uint, uint);
0356
0357 // ide.c
0358 void          ideinit(void);
0359 void          ideintr(void);
0360 void          iderw(struct buf*);
0361
0362 // ioapic.c
0363 void          ioapicenable(int irq, int cpu);
0364 extern uchar  ioapicid;
0365 void          ioapicinit(void);
0366
0367 // kalloc.c
0368 char*         kalloc(void);
0369 void          kfree(char*);
0370 void          kinit1(void*, void*);
0371 void          kinit2(void*, void*);
0372
0373 // kbd.c
0374 void          kbdintr(void);
0375
0376 // lapic.c
0377 void          cmostime(struct rtcdate *r);
0378 int           cpunum(void);
0379 extern volatile uint* lapic;
0380 void          lapiceoi(void);
0381 void          lapicinit(void);
0382 void          lapicstartap(uchar, uint);
0383 void          microdelay(int);
0384
0385 // log.c
0386 void          initlog(int dev);
0387 void          log_write(struct buf*);
0388 void          begin_op();
0389 void          end_op();
0390
0391 // mp.c
0392 extern int     ismp;
0393 int           mpbcpu(void);
0394 void          mpinit(void);
0395 void          mpstartthem(void);
0396
0397 // picirq.c
0398 void          picenable(int);
0399 void          picinit(void);

```

```

0400 // pipe.c
0401 int      pipealloc(struct file**, struct file**);
0402 void      pipeclose(struct pipe*, int);
0403 int      piperead(struct pipe*, char*, int);
0404 int      pipewrite(struct pipe*, char*, int);
0405
0406 // proc.c
0407 struct proc* copyproc(struct proc*);
0408 void      exit(void);
0409 int      fork(void);
0410 int      growproc(int);
0411 int      kill(int);
0412 void      pinit(void);
0413 void      procdump(void);
0414 void      scheduler(void) __attribute__((noreturn));
0415 void      sched(void);
0416 void      sleep(void*, struct spinlock*);
0417 void      userinit(void);
0418 int      wait(void);
0419 void      wakeup(void*);
0420 void      yield(void);
0421 #ifdef CS333_P2
0422 int
0423 #endif
0424 #ifdef CS333_P3
0425 int
0426 #endif
0427
0428 // swtch.S
0429 void      swtch(struct context**, struct context*);
0430
0431 // spinlock.c
0432 void      acquire(struct spinlock*);
0433 void      getcallerpcs(void*, uint*);
0434 int      holding(struct spinlock*);
0435 void      initlock(struct spinlock*, char*);
0436 void      release(struct spinlock*);
0437 void      pushcli(void);
0438 void      popcli(void);
0439
0440 // string.c
0441 int      memcmp(const void*, const void*, uint);
0442 void*     memmove(void*, const void*, uint);
0443 void*     memset(void*, int, uint);
0444 char*     safestrcpy(char*, const char*, int);
0445 int      strlen(const char*);
0446 int      strncmp(const char*, const char*, uint);
0447 char*     strncpy(char*, const char*, int);
0448
0449

```

```

0450 // syscall.c
0451 int      argint(int, int*);
0452 int      argptr(int, char**, int);
0453 int      argstr(int, char**);
0454 int      fetchint(uint, int*);
0455 int      fetchstr(uint, char**);
0456 void      syscall(void);
0457
0458 // timer.c
0459 void      timerinit(void);
0460
0461 // trap.c
0462 void      idtinit(void);
0463 extern uint ticks;
0464 void      tvinit(void);
0465 extern struct spinlock tickslock;
0466
0467 // uart.c
0468 void      uartinit(void);
0469 void      uartintr(void);
0470 void      uartputc(int);
0471
0472 // vm.c
0473 void      seginit(void);
0474 void      kvmalloc(void);
0475 void      vmenable(void);
0476 pde_t*     setupkvm(void);
0477 char*     uva2ka(pde_t*, char*);
0478 int      allocvm(pde_t*, uint, uint);
0479 int      deallocvm(pde_t*, uint, uint);
0480 void      freevm(pde_t*);
0481 void      initvm(pde_t*, char*, uint);
0482 int      loadvm(pde_t*, char*, struct inode*, uint, uint);
0483 pde_t*     copyvm(pde_t*, uint);
0484 void      switchvm(struct proc*);
0485 void      switchkvm(void);
0486 int      copyout(pde_t*, uint, void*, uint);
0487 void      clearpteu(pde_t *pgdir, char *uva);
0488
0489 // number of elements in fixed-size array
0490 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))
0491
0492
0493
0494
0495
0496
0497
0498
0499

```

```

0500 // Routines to let C code use special x86 instructions.
0501
0502 static inline uchar
0503 inb(ushort port)
0504 {
0505     uchar data;
0506
0507     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
0508     return data;
0509 }
0510
0511 static inline void
0512 insl(int port, void *addr, int cnt)
0513 {
0514     asm volatile("cld; rep insl" :
0515                 "=D" (addr), "=c" (cnt) :
0516                 "d" (port), "0" (addr), "1" (cnt) :
0517                 "memory", "cc");
0518 }
0519
0520 static inline void
0521 outb(ushort port, uchar data)
0522 {
0523     asm volatile("out %0,%1" : : "a" (data), "d" (port));
0524 }
0525
0526 static inline void
0527 outw(ushort port, ushort data)
0528 {
0529     asm volatile("out %0,%1" : : "a" (data), "d" (port));
0530 }
0531
0532 static inline void
0533 outsl(int port, const void *addr, int cnt)
0534 {
0535     asm volatile("cld; rep outsl" :
0536                 "=S" (addr), "=c" (cnt) :
0537                 "d" (port), "0" (addr), "1" (cnt) :
0538                 "cc");
0539 }
0540
0541 static inline void
0542 stosb(void *addr, int data, int cnt)
0543 {
0544     asm volatile("cld; rep stosb" :
0545                 "=D" (addr), "=c" (cnt) :
0546                 "0" (addr), "1" (cnt), "a" (data) :
0547                 "memory", "cc");
0548 }
0549

```

```

0550 static inline void
0551 stosl(void *addr, int data, int cnt)
0552 {
0553     asm volatile("cld; rep stosl" :
0554                 "=D" (addr), "=c" (cnt) :
0555                 "0" (addr), "1" (cnt), "a" (data) :
0556                 "memory", "cc");
0557 }
0558
0559 struct segdesc;
0560
0561 static inline void
0562 lgdt(struct segdesc *p, int size)
0563 {
0564     volatile ushort pd[3];
0565
0566     pd[0] = size-1;
0567     pd[1] = (uint)p;
0568     pd[2] = (uint)p >> 16;
0569
0570     asm volatile("lgdt (%0)" : : "r" (pd));
0571 }
0572
0573 struct gatedesc;
0574
0575 static inline void
0576 lidt(struct gatedesc *p, int size)
0577 {
0578     volatile ushort pd[3];
0579
0580     pd[0] = size-1;
0581     pd[1] = (uint)p;
0582     pd[2] = (uint)p >> 16;
0583
0584     asm volatile("lidt (%0)" : : "r" (pd));
0585 }
0586
0587 static inline void
0588 ltr(ushort sel)
0589 {
0590     asm volatile("ltr %0" : : "r" (sel));
0591 }
0592
0593 static inline uint
0594 readeflags(void)
0595 {
0596     uint eflags;
0597     asm volatile("pushfl; popl %0" : "=r" (eflags));
0598     return eflags;
0599 }

```

```

0600 static inline void
0601 loadgs(ushort v)
0602 {
0603     asm volatile("movw %0, %%gs" : : "r" (v));
0604 }
0605
0606 static inline void
0607 cli(void)
0608 {
0609     asm volatile("cli");
0610 }
0611
0612 static inline void
0613 sti(void)
0614 {
0615     asm volatile("sti");
0616 }
0617
0618 static inline uint
0619 xchg(volatile uint *addr, uint newval)
0620 {
0621     uint result;
0622
0623     // The + in "+m" denotes a read-modify-write operand.
0624     asm volatile("lock; xchgl %0, %1" :
0625                 "+m" (*addr), "=a" (result) :
0626                 "l" (newval) :
0627                 "cc");
0628     return result;
0629 }
0630
0631 static inline uint
0632 rcr2(void)
0633 {
0634     uint val;
0635     asm volatile("movl %%cr2,%0" : "=r" (val));
0636     return val;
0637 }
0638
0639 static inline void
0640 lcr3(uint val)
0641 {
0642     asm volatile("movl %0,%%cr3" : : "r" (val));
0643 }
0644
0645
0646
0647
0648
0649

```

```

0650 // Layout of the trap frame built on the stack by the
0651 // hardware and by trapasm.S, and passed to trap().
0652 struct trapframe {
0653     // registers as pushed by pusha
0654     uint edi;
0655     uint esi;
0656     uint ebp;
0657     uint oesp;      // useless & ignored
0658     uint ebx;
0659     uint edx;
0660     uint ecx;
0661     uint eax;
0662
0663     // rest of trap frame
0664     ushort gs;
0665     ushort padding1;
0666     ushort fs;
0667     ushort padding2;
0668     ushort es;
0669     ushort padding3;
0670     ushort ds;
0671     ushort padding4;
0672     uint trapno;
0673
0674     // below here defined by x86 hardware
0675     uint err;
0676     uint eip;
0677     ushort cs;
0678     ushort padding5;
0679     uint eflags;
0680
0681     // below here only when crossing rings, such as from user to kernel
0682     uint esp;
0683     ushort ss;
0684     ushort padding6;
0685 };
0686
0687
0688
0689
0690
0691
0692
0693
0694
0695
0696
0697
0698
0699

```

```

0700 //
0701 // assembler macros to create x86 segments
0702 //
0703
0704 #define SEG_NULLASM \
0705     .word 0, 0; \
0706     .byte 0, 0, 0, 0
0707
0708 // The 0xC0 means the limit is in 4096-byte units
0709 // and (for executable segments) 32-bit mode.
0710 #define SEG_ASM(type,base,lim) \
0711     .word (((lim) >> 12) & 0xffff), ((base) & 0xffff); \
0712     .byte (((base) >> 16) & 0xff), (0x90 | (type)), \
0713     (0xC0 | (((lim) >> 28) & 0xf)), (((base) >> 24) & 0xff)
0714
0715 #define STA_X 0x8 // Executable segment
0716 #define STA_E 0x4 // Expand down (non-executable segments)
0717 #define STA_C 0x4 // Conforming code segment (executable only)
0718 #define STA_W 0x2 // Writeable (non-executable segments)
0719 #define STA_R 0x2 // Readable (executable segments)
0720 #define STA_A 0x1 // Accessed
0721
0722
0723
0724
0725
0726
0727
0728
0729
0730
0731
0732
0733
0734
0735
0736
0737
0738
0739
0740
0741
0742
0743
0744
0745
0746
0747
0748
0749

```

```

0750 // This file contains definitions for the
0751 // x86 memory management unit (MMU).
0752
0753 // Eflags register
0754 #define FL_CF 0x00000001 // Carry Flag
0755 #define FL_PF 0x00000004 // Parity Flag
0756 #define FL_AF 0x00000010 // Auxiliary carry Flag
0757 #define FL_ZF 0x00000040 // Zero Flag
0758 #define FL_SF 0x00000080 // Sign Flag
0759 #define FL_TF 0x00000100 // Trap Flag
0760 #define FL_IF 0x00000200 // Interrupt Enable
0761 #define FL_DF 0x00000400 // Direction Flag
0762 #define FL_OF 0x00000800 // Overflow Flag
0763 #define FL_IOPL_MASK 0x00003000 // I/O Privilege Level bitmask
0764 #define FL_IOPL_0 0x00000000 // IOPL == 0
0765 #define FL_IOPL_1 0x00001000 // IOPL == 1
0766 #define FL_IOPL_2 0x00002000 // IOPL == 2
0767 #define FL_IOPL_3 0x00003000 // IOPL == 3
0768 #define FL_NT 0x00004000 // Nested Task
0769 #define FL_RF 0x00010000 // Resume Flag
0770 #define FL_VM 0x00020000 // Virtual 8086 mode
0771 #define FL_AC 0x00040000 // Alignment Check
0772 #define FL_VIF 0x00080000 // Virtual Interrupt Flag
0773 #define FL_VIP 0x00100000 // Virtual Interrupt Pending
0774 #define FL_ID 0x00200000 // ID flag
0775
0776 // Control Register flags
0777 #define CR0_PE 0x00000001 // Protection Enable
0778 #define CR0_MP 0x00000002 // Monitor coProcessor
0779 #define CR0_EM 0x00000004 // Emulation
0780 #define CR0_TS 0x00000008 // Task Switched
0781 #define CR0_ET 0x00000010 // Extension Type
0782 #define CR0_NE 0x00000020 // Numeric Error
0783 #define CR0_WP 0x00010000 // Write Protect
0784 #define CR0_AM 0x00040000 // Alignment Mask
0785 #define CR0_NW 0x00080000 // Not Writethrough
0786 #define CR0_CD 0x00100000 // Cache Disable
0787 #define CR0_PG 0x00200000 // Paging
0788
0789 #define CR4_PSE 0x00000010 // Page size extension
0790
0791 #define SEG_KCODE 1 // kernel code
0792 #define SEG_KDATA 2 // kernel data+stack
0793 #define SEG_KCPU 3 // kernel per-cpu data
0794 #define SEG_UCODE 4 // user code
0795 #define SEG_UDATA 5 // user data+stack
0796 #define SEG_TSS 6 // this process's task state
0797
0798
0799

```



```

0800 #ifndef __ASSEMBLER__
0801 // Segment Descriptor
0802 struct segdesc {
0803     uint lim_15_0 : 16; // Low bits of segment limit
0804     uint base_15_0 : 16; // Low bits of segment base address
0805     uint base_23_16 : 8; // Middle bits of segment base address
0806     uint type : 4;       // Segment type (see STS_constants)
0807     uint s : 1;         // 0 = system, 1 = application
0808     uint dpl : 2;       // Descriptor Privilege Level
0809     uint p : 1;         // Present
0810     uint lim_19_16 : 4; // High bits of segment limit
0811     uint avl : 1;       // Unused (available for software use)
0812     uint rsv1 : 1;      // Reserved
0813     uint db : 1;        // 0 = 16-bit segment, 1 = 32-bit segment
0814     uint g : 1;         // Granularity: limit scaled by 4K when set
0815     uint base_31_24 : 8; // High bits of segment base address
0816 };
0817
0818 // Normal segment
0819 #define SEG(type, base, lim, dpl) (struct segdesc) \
0820 { ((lim) >> 12) & 0xffff, (uint)(base) & 0xffff, \
0821   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
0822   (uint)(lim) >> 28, 0, 0, 1, 1, (uint)(base) >> 24 }
0823 #define SEG16(type, base, lim, dpl) (struct segdesc) \
0824 { (lim) & 0xffff, (uint)(base) & 0xffff, \
0825   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
0826   (uint)(lim) >> 16, 0, 0, 1, 0, (uint)(base) >> 24 }
0827 #endif
0828
0829 #define DPL_USER 0x3 // User DPL
0830
0831 // Application segment type bits
0832 #define STA_X 0x8 // Executable segment
0833 #define STA_E 0x4 // Expand down (non-executable segments)
0834 #define STA_C 0x4 // Conforming code segment (executable only)
0835 #define STA_W 0x2 // Writeable (non-executable segments)
0836 #define STA_R 0x2 // Readable (executable segments)
0837 #define STA_A 0x1 // Accessed
0838
0839 // System segment type bits
0840 #define STS_T16A 0x1 // Available 16-bit TSS
0841 #define STS_LDT 0x2 // Local Descriptor Table
0842 #define STS_T16B 0x3 // Busy 16-bit TSS
0843 #define STS_CG16 0x4 // 16-bit Call Gate
0844 #define STS_TG 0x5 // Task Gate / Coum Transmissions
0845 #define STS_IG16 0x6 // 16-bit Interrupt Gate
0846 #define STS_TG16 0x7 // 16-bit Trap Gate
0847 #define STS_T32A 0x9 // Available 32-bit TSS
0848 #define STS_T32B 0xB // Busy 32-bit TSS
0849 #define STS_CG32 0xC // 32-bit Call Gate

```

```

0850 #define STS_IG32 0xE // 32-bit Interrupt Gate
0851 #define STS_TG32 0xF // 32-bit Trap Gate
0852
0853 // A virtual address 'la' has a three-part structure as follows:
0854 //
0855 // +-----10-----+-----10-----+-----12-----+
0856 // | Page Directory | Page Table | Offset within Page |
0857 // | Index | Index | |
0858 // +-----+-----+-----+
0859 // \--- PDX(va) --/ \--- PTX(va) --/
0860
0861 // page directory index
0862 #define PDX(va) (((uint)(va) >> PDXSHIFT) & 0x3FF)
0863
0864 // page table index
0865 #define PTX(va) (((uint)(va) >> PTXSHIFT) & 0x3FF)
0866
0867 // construct virtual address from indexes and offset
0868 #define PGADDR(d, t, o) ((uint)((d) << PDXSHIFT | (t) << PTXSHIFT | (o)))
0869
0870 // Page directory and page table constants.
0871 #define NPENTRIES 1024 // # directory entries per page directory
0872 #define NPTENTRIES 1024 // # PTEs per page table
0873 #define PGSIZE 4096 // bytes mapped by a page
0874
0875 #define PGSHIFT 12 // log2(PGSIZE)
0876 #define PTXSHIFT 12 // offset of PTX in a linear address
0877 #define PDXSHIFT 22 // offset of PDX in a linear address
0878
0879 #define PGROUNDUP(sz) (((sz)+PGSIZE-1) & ~(PGSIZE-1))
0880 #define PGROUNDDOWN(a) (((a)) & ~(PGSIZE-1))
0881
0882 // Page table/directory entry flags.
0883 #define PTE_P 0x001 // Present
0884 #define PTE_W 0x002 // Writeable
0885 #define PTE_U 0x004 // User
0886 #define PTE_PWT 0x008 // Write-Through
0887 #define PTE_PCD 0x010 // Cache-Disable
0888 #define PTE_A 0x020 // Accessed
0889 #define PTE_D 0x040 // Dirty
0890 #define PTE_PS 0x080 // Page Size
0891 #define PTE_MBZ 0x180 // Bits must be zero
0892
0893 // Address in page table or page directory entry
0894 #define PTE_ADDR(pte) ((uint)(pte) & ~0xFFF)
0895 #define PTE_FLAGS(pte) ((uint)(pte) & 0xFFF)
0896
0897 #ifndef __ASSEMBLER__
0898 typedef uint pte_t;
0899

```

```

0900 // Task state segment format
0901 struct taskstate {
0902     uint link;           // Old ts selector
0903     uint esp0;           // Stack pointers and segment selectors
0904     ushort ss0;          // after an increase in privilege level
0905     ushort padding1;
0906     uint *esp1;
0907     ushort ssl;
0908     ushort padding2;
0909     uint *esp2;
0910     ushort ss2;
0911     ushort padding3;
0912     void *cr3;           // Page directory base
0913     uint *eip;           // Saved state from last task switch
0914     uint eflags;
0915     uint eax;            // More saved state (registers)
0916     uint ecx;
0917     uint edx;
0918     uint ebx;
0919     uint *esp;
0920     uint *ebp;
0921     uint esi;
0922     uint edi;
0923     ushort es;           // Even more saved state (segment selectors)
0924     ushort padding4;
0925     ushort cs;
0926     ushort padding5;
0927     ushort ss;
0928     ushort padding6;
0929     ushort ds;
0930     ushort padding7;
0931     ushort fs;
0932     ushort padding8;
0933     ushort gs;
0934     ushort padding9;
0935     ushort ldt;
0936     ushort padding10;
0937     ushort t;            // Trap on task switch
0938     ushort iomb;         // I/O map base address
0939 };
0940
0941
0942
0943
0944
0945
0946
0947
0948
0949

```

```

0950 // Gate descriptors for interrupts and traps
0951 struct gatedesc {
0952     uint off_15_0 : 16;  // low 16 bits of offset in segment
0953     uint cs : 16;         // code segment selector
0954     uint args : 5;       // # args, 0 for interrupt/trap gates
0955     uint rsv1 : 3;        // reserved(should be zero I guess)
0956     uint type : 4;        // type(STS_{TG,IG32,TG32})
0957     uint s : 1;          // must be 0 (system)
0958     uint dpl : 2;        // descriptor(meaning new) privilege level
0959     uint p : 1;          // Present
0960     uint off_31_16 : 16; // high bits of offset in segment
0961 };
0962
0963 // Set up a normal interrupt/trap gate descriptor.
0964 // - istrap: 1 for a trap (= exception) gate, 0 for an interrupt gate.
0965 // - sel: interrupt gate clears FL_IF, trap gate leaves FL_IF alone
0966 // - sel: Code segment selector for interrupt/trap handler
0967 // - off: Offset in code segment for interrupt/trap handler
0968 // - dpl: Descriptor Privilege Level -
0969 //       the privilege level required for software to invoke
0970 //       this interrupt/trap gate explicitly using an int instruction.
0971 #define SETGATE(gate, istrap, sel, off, d) \
0972 { \
0973     (gate).off_15_0 = (uint)(off) & 0xffff; \
0974     (gate).cs = (sel); \
0975     (gate).args = 0; \
0976     (gate).rsv1 = 0; \
0977     (gate).type = (istrap) ? STS_TG32 : STS_IG32; \
0978     (gate).s = 0; \
0979     (gate).dpl = (d); \
0980     (gate).p = 1; \
0981     (gate).off_31_16 = (uint)(off) >> 16; \
0982 }
0983
0984 #endif
0985
0986
0987
0988
0989
0990
0991
0992
0993
0994
0995
0996
0997
0998
0999

```

```

1000 // Format of an ELF executable file
1001
1002 #define ELF_MAGIC 0x464C457FU // "\x7FELF" in little endian
1003
1004 // File header
1005 struct elfhdr {
1006     uint magic; // must equal ELF_MAGIC
1007     uchar elf[12];
1008     ushort type;
1009     ushort machine;
1010     uint version;
1011     uint entry;
1012     uint phoff;
1013     uint shoff;
1014     uint flags;
1015     ushort ehsize;
1016     ushort phentsize;
1017     ushort phnum;
1018     ushort shentsize;
1019     ushort shnum;
1020     ushort shstrndx;
1021 };
1022
1023 // Program section header
1024 struct proghdr {
1025     uint type;
1026     uint off;
1027     uint vaddr;
1028     uint paddr;
1029     uint filesz;
1030     uint memsz;
1031     uint flags;
1032     uint align;
1033 };
1034
1035 // Values for Proghdr type
1036 #define ELF_PROG_LOAD 1
1037
1038 // Flag bits for Proghdr flags
1039 #define ELF_PROG_FLAG_EXEC 1
1040 #define ELF_PROG_FLAG_WRITE 2
1041 #define ELF_PROG_FLAG_READ 4
1042
1043
1044
1045
1046
1047
1048
1049

```

```

1050 # Multiboot header, for multiboot boot loaders like GNU Grub.
1051 # http://www.gnu.org/software/grub/manual/multiboot/multiboot.html
1052 #
1053 # Using GRUB 2, you can boot xv6 from a file stored in a
1054 # Linux file system by copying kernel or kernelmemfs to /boot
1055 # and then adding this menu entry:
1056 #
1057 # menuentry "xv6" {
1058 #     insmod ext2
1059 #     set root='(hd0,msdos1)'
1060 #     set kernel='/boot/kernel'
1061 #     echo "Loading ${kernel}..."
1062 #     multiboot ${kernel} ${kernel}
1063 #     boot
1064 # }
1065
1066 #include "asm.h"
1067 #include "memlayout.h"
1068 #include "mmu.h"
1069 #include "param.h"
1070
1071 # Multiboot header. Data to direct multiboot loader.
1072 .p2align 2
1073 .text
1074 .globl multiboot_header
1075 multiboot_header:
1076     #define magic 0x1badb002
1077     #define flags 0
1078     .long magic
1079     .long flags
1080     .long (-magic-flags)
1081
1082 # By convention, the _start symbol specifies the ELF entry point.
1083 # Since we haven't set up virtual memory yet, our entry point is
1084 # the physical address of 'entry'.
1085 .globl _start
1086 _start = V2P_WO(entry)
1087
1088 # Entering xv6 on boot processor, with paging off.
1089 .globl entry
1090 entry:
1091     # Turn on page size extension for 4Mbyte pages
1092     movl    %cr4, %eax
1093     orl     $(CR4_PSE), %eax
1094     movl    %eax, %cr4
1095     # Set page directory
1096     movl    $(V2P_WO(entrypgdir)), %eax
1097     movl    %eax, %cr3
1098     # Turn on paging.
1099     movl    %cr0, %eax

```

```

1100 orl    $(CR0_PG|CR0_WP), %eax
1101 movl   %eax, %cr0
1102
1103 # Set up the stack pointer.
1104 movl $(stack + KSTACKSIZE), %esp
1105
1106 # Jump to main(), and switch to executing at
1107 # high addresses. The indirect call is needed because
1108 # the assembler produces a PC-relative instruction
1109 # for a direct jump.
1110 mov $main, %eax
1111 jmp *%eax
1112
1113 .comm stack, KSTACKSIZE
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149

```

```

1150 #include "asm.h"
1151 #include "memlayout.h"
1152 #include "mmu.h"
1153
1154 # Each non-boot CPU ("AP") is started up in response to a STARTUP
1155 # IPI from the boot CPU. Section B.4.2 of the Multi-Processor
1156 # Specification says that the AP will start in real mode with CS:IP
1157 # set to XY00:0000, where XY is an 8-bit value sent with the
1158 # STARTUP. Thus this code must start at a 4096-byte boundary.
1159 #
1160 # Because this code sets DS to zero, it must sit
1161 # at an address in the low 2^16 bytes.
1162 #
1163 # Startothers (in main.c) sends the STARTUPs one at a time.
1164 # It copies this code (start) at 0x7000. It puts the address of
1165 # a newly allocated per-core stack in start-4, the address of the
1166 # place to jump to (mpenter) in start-8, and the physical address
1167 # of entrypgdir in start-12.
1168 #
1169 # This code is identical to bootasm.S except:
1170 #   - it does not need to enable A20
1171 #   - it uses the address at start-4, start-8, and start-12
1172
1173 .code16
1174 .globl start
1175 start:
1176 cli
1177
1178 xorw    %ax, %ax
1179 movw    %ax, %ds
1180 movw    %ax, %es
1181 movw    %ax, %ss
1182
1183 lgdt    gdt_desc
1184 movl    %cr0, %eax
1185 orl     $CR0_PE, %eax
1186 movl    %eax, %cr0
1187
1188 ljmpl    $(SEG_KCODE<<3), $(start32)
1189
1190 .code32
1191 start32:
1192 movw    $(SEG_KDATA<<3), %ax
1193 movw    %ax, %ds
1194 movw    %ax, %es
1195 movw    %ax, %ss
1196 movw    $0, %ax
1197 movw    %ax, %fs
1198 movw    %ax, %gs
1199

```

```

1200 # Turn on page size extension for 4Mbyte pages
1201 movl    %cr4, %eax
1202 orl     $(CR4_PSE), %eax
1203 movl    %eax, %cr4
1204 # Use enterpgdir as our initial page table
1205 movl    (start-12), %eax
1206 movl    %eax, %cr3
1207 # Turn on paging.
1208 movl    %cr0, %eax
1209 orl     $(CR0_PE|CR0_PG|CR0_WP), %eax
1210 movl    %eax, %cr0
1211
1212 # Switch to the stack allocated by startothers()
1213 movl    (start-4), %esp
1214 # Call mpenter()
1215 call    *(start-8)
1216
1217 movw    $0x8a00, %ax
1218 movw    %ax, %dx
1219 outw    %ax, %dx
1220 movw    $0x8ae0, %ax
1221 outw    %ax, %dx
1222 spin:
1223 jmp     spin
1224
1225 .p2align 2
1226 gdt:
1227 SEG_NULLASM
1228 SEG_ASM(STA_X|STA_R, 0, 0xffffffff)
1229 SEG_ASM(STA_W, 0, 0xffffffff)
1230
1231
1232 gdtdesc:
1233 .word    (gdtdesc - gdt - 1)
1234 .long    gdt
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249

```

```

1250 #include "types.h"
1251 #include "defs.h"
1252 #include "param.h"
1253 #include "memlayout.h"
1254 #include "mmu.h"
1255 #include "proc.h"
1256 #include "x86.h"
1257
1258 static void startothers(void);
1259 static void mpmain(void) __attribute__((noreturn));
1260 extern pde_t *kpgdir;
1261 extern char end[]; // first address after kernel loaded from ELF file
1262
1263 // Bootstrap processor starts running C code here.
1264 // Allocate a real stack and switch to it, first
1265 // doing some setup required for memory allocator to work.
1266 int
1267 main(void)
1268 {
1269     kinit1(end, P2V(4*1024*1024)); // phys page allocator
1270     kvmalloc(); // kernel page table
1271     mpinit(); // collect info about this machine
1272     lapicinit();
1273     seginit(); // set up segments
1274     cprintf("ncpu%d: starting xv6\n\n", cpu->id);
1275     picinit(); // interrupt controller
1276     ioapicinit(); // another interrupt controller
1277     consoleinit(); // I/O devices & their interrupts
1278     uartinit(); // serial port
1279     pinit(); // process table
1280     tvinit(); // trap vectors
1281     binit(); // buffer cache
1282     fileinit(); // file table
1283     ideinit(); // disk
1284     if(!ismp)
1285         timerinit(); // uniprocessor timer
1286     startothers(); // start other processors
1287     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
1288     userinit(); // first user process
1289     // Finish setting up this processor in mpmain.
1290     mpmain();
1291 }
1292
1293
1294
1295
1296
1297
1298
1299

```

```

1300 // Other CPUs jump here from entryother.S.
1301 static void
1302 mpenter(void)
1303 {
1304     switchkvm();
1305     seginit();
1306     lapicinit();
1307     mpmain();
1308 }
1309
1310 // Common CPU setup code.
1311 static void
1312 mpmain(void)
1313 {
1314     cprintf("cpu%d: starting\n", cpu->id);
1315     idtinit(); // load idt register
1316     xchg(&cpu->started, 1); // tell startothers() we're up
1317     scheduler(); // start running processes
1318 }
1319
1320 pde_t entrypgdir[]; // For entry.S
1321
1322 // Start the non-boot (AP) processors.
1323 static void
1324 startothers(void)
1325 {
1326     extern uchar _binary_entryother_start[], _binary_entryother_size[];
1327     uchar *code;
1328     struct cpu *c;
1329     char *stack;
1330
1331     // Write entry code to unused memory at 0x7000.
1332     // The linker has placed the image of entryother.S in
1333     // _binary_entryother_start.
1334     code = p2v(0x7000);
1335     memmove(code, _binary_entryother_start, (uint)_binary_entryother_size);
1336
1337     for(c = cpus; c < cpus+ncpu; c++){
1338         if(c == cpus+cpunum()) // We've started already.
1339             continue;
1340
1341         // Tell entryother.S what stack to use, where to enter, and what
1342         // pgdir to use. We cannot use kpgdir yet, because the AP processor
1343         // is running in low memory, so we use entrypgdir for the APs too.
1344         stack = kalloc();
1345         *(void**) (code-4) = stack + KSTACKSIZE;
1346         *(void**) (code-8) = mpenter;
1347         *(int**) (code-12) = (void *) v2p(entrypgdir);
1348
1349         lapicstartap(c->id, v2p(code));

```

```

1350     // wait for cpu to finish mpmain()
1351     while(c->started == 0)
1352         ;
1353 }
1354 }
1355
1356 // Boot page table used in entry.S and entryother.S.
1357 // Page directories (and page tables), must start on a page boundary,
1358 // hence the "__aligned__" attribute.
1359 // Use PTE_PS in page directory entry to enable 4Mbyte pages.
1360 __attribute__((__aligned__(PGSIZE)))
1361 pde_t entrypgdir[NPDENTRIES] = {
1362     // Map VA's [0, 4MB) to PA's [0, 4MB)
1363     [0] = (0) | PTE_P | PTE_W | PTE_PS,
1364     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
1365     [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
1366 };
1367
1368 // Blank page.
1369 // Blank page.
1370 // Blank page.
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399

```

```

1400 // Mutual exclusion lock.
1401 struct spinlock {
1402     uint locked;        // Is the lock held?
1403
1404     // For debugging:
1405     char *name;         // Name of lock.
1406     struct cpu *cpu;    // The cpu holding the lock.
1407     uint pcs[10];       // The call stack (an array of program counters)
1408                        // that locked the lock.
1409 };
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449

```

```

1450 // Mutual exclusion spin locks.
1451
1452 #include "types.h"
1453 #include "defs.h"
1454 #include "param.h"
1455 #include "x86.h"
1456 #include "memlayout.h"
1457 #include "mmu.h"
1458 #include "proc.h"
1459 #include "spinlock.h"
1460
1461 void
1462 initlock(struct spinlock *lk, char *name)
1463 {
1464     lk->name = name;
1465     lk->locked = 0;
1466     lk->cpu = 0;
1467 }
1468
1469 // Acquire the lock.
1470 // Loops (spins) until the lock is acquired.
1471 // Holding a lock for a long time may cause
1472 // other CPUs to waste time spinning to acquire it.
1473 void
1474 acquire(struct spinlock *lk)
1475 {
1476     pushcli(); // disable interrupts to avoid deadlock.
1477     if(holding(lk))
1478         panic("acquire");
1479
1480     // The xchg is atomic.
1481     // It also serializes, so that reads after acquire are not
1482     // reordered before it.
1483     while(xchg(&lk->locked, 1) != 0)
1484         ;
1485
1486     // Record info about lock acquisition for debugging.
1487     lk->cpu = cpu;
1488     getcallerpcs(&lk, lk->pcs);
1489 }
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499

```

```

1500 // Release the lock.
1501 void
1502 release(struct spinlock *lk)
1503 {
1504     if(!holding(lk))
1505         panic("release");
1506
1507     lk->pcs[0] = 0;
1508     lk->cpu = 0;
1509
1510     // The xchg serializes, so that reads before release are
1511     // not reordered after it. The 1996 PentiumPro manual (Volume 3,
1512     // 7.2) says reads can be carried out speculatively and in
1513     // any order, which implies we need to serialize here.
1514     // But the 2007 Intel 64 Architecture Memory Ordering White
1515     // Paper says that Intel 64 and IA-32 will not move a load
1516     // after a store. So lock->locked = 0 would work here.
1517     // The xchg being asm volatile ensures gcc emits it after
1518     // the above assignments (and after the critical section).
1519     xchg(&lk->locked, 0);
1520
1521     popcli();
1522 }
1523
1524 // Record the current call stack in pcs[] by following the %ebp chain.
1525 void
1526 getcallerpcs(void *v, uint pcs[])
1527 {
1528     uint *ebp;
1529     int i;
1530
1531     ebp = (uint*)v - 2;
1532     for(i = 0; i < 10; i++){
1533         if(ebp == 0 || ebp < (uint*)KERNBASE || ebp == (uint*)0xffffffff)
1534             break;
1535         pcs[i] = ebp[1]; // saved %eip
1536         ebp = (uint*)ebp[0]; // saved %ebp
1537     }
1538     for(; i < 10; i++)
1539         pcs[i] = 0;
1540 }
1541
1542 // Check whether this cpu is holding the lock.
1543 int
1544 holding(struct spinlock *lock)
1545 {
1546     return lock->locked && lock->cpu == cpu;
1547 }
1548
1549

```

```

1550 // Pushcli/popcli are like cli/sti except that they are matched:
1551 // it takes two popcli to undo two pushcli. Also, if interrupts
1552 // are off, then pushcli, popcli leaves them off.
1553
1554 void
1555 pushcli(void)
1556 {
1557     int eflags;
1558
1559     eflags = readeflags();
1560     cli();
1561     if(cpu->ncli++ == 0)
1562         cpu->intena = eflags & FL_IF;
1563 }
1564
1565 void
1566 popcli(void)
1567 {
1568     if(readeflags() & FL_IF)
1569         panic("popcli - interruptible");
1570     if(--cpu->ncli < 0)
1571         panic("popcli");
1572     if(cpu->ncli == 0 && cpu->intena)
1573         sti();
1574 }
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599

```



```

1600 #include "param.h"
1601 #include "types.h"
1602 #include "defs.h"
1603 #include "x86.h"
1604 #include "memlayout.h"
1605 #include "mmu.h"
1606 #include "proc.h"
1607 #include "elf.h"
1608
1609 extern char data[]; // defined by kernel.ld
1610 pde_t *kpgdir; // for use in scheduler()
1611 struct segdesc gdt[NSEGS];
1612
1613 // Set up CPU's kernel segment descriptors.
1614 // Run once on entry on each CPU.
1615 void
1616 seginit(void)
1617 {
1618     struct cpu *c;
1619
1620     // Map "logical" addresses to virtual addresses using identity map.
1621     // Cannot share a CODE descriptor for both kernel and user
1622     // because it would have to have DPL_USR, but the CPU forbids
1623     // an interrupt from CPL=0 to DPL=3.
1624     c = &cpu[cpunum()];
1625     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
1626     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
1627     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
1628     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
1629
1630     // Map cpu, and curproc
1631     c->gdt[SEG_KCPU] = SEG(STA_W, &c->cpu, 8, 0);
1632
1633     lgdt(c->gdt, sizeof(c->gdt));
1634     loadgs(SEG_KCPU << 3);
1635
1636     // Initialize cpu-local storage.
1637     cpu = c;
1638     proc = 0;
1639 }
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649

```

```

1650 // Return the address of the PTE in page table pgdir
1651 // that corresponds to virtual address va. If alloc!=0,
1652 // create any required page table pages.
1653 static pte_t *
1654 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1655 {
1656     pde_t *pde;
1657     pte_t *pgtab;
1658
1659     pde = &pgdir[PDX(va)];
1660     if(*pde & PTE_P){
1661         pgtab = (pte_t*)p2v(PTE_ADDR(*pde));
1662     } else {
1663         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1664             return 0;
1665         // Make sure all those PTE_P bits are zero.
1666         memset(pgtab, 0, PGSIZE);
1667         // The permissions here are overly generous, but they can
1668         // be further restricted by the permissions in the page table
1669         // entries, if necessary.
1670         *pde = v2p(pgtab) | PTE_P | PTE_W | PTE_U;
1671     }
1672     return &pgtab[PTX(va)];
1673 }
1674
1675 // Create PTEs for virtual addresses starting at va that refer to
1676 // physical addresses starting at pa. va and size might not
1677 // be page-aligned.
1678 static int
1679 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1680 {
1681     char *a, *last;
1682     pte_t *pte;
1683
1684     a = (char*)PGROUNDDOWN((uint)va);
1685     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
1686     for(;;){
1687         if((pte = walkpgdir(pgdir, a, 1)) == 0)
1688             return -1;
1689         if(*pte & PTE_P)
1690             panic("remap");
1691         *pte = pa | perm | PTE_P;
1692         if(a == last)
1693             break;
1694         a += PGSIZE;
1695         pa += PGSIZE;
1696     }
1697     return 0;
1698 }
1699

```

```

1700 // There is one page table per process, plus one that's used when
1701 // a CPU is not running any process (kpgdir). The kernel uses the
1702 // current process's page table during system calls and interrupts;
1703 // page protection bits prevent user code from using the kernel's
1704 // mappings.
1705 //
1706 // setupkvm() and exec() set up every page table like this:
1707 //
1708 // 0..KERNBASE: user memory (text+data+stack+heap), mapped to
1709 // phys memory allocated by the kernel
1710 // KERNBASE..KERNBASE+EXTMEM: mapped to 0..EXTMEM (for I/O space)
1711 // KERNBASE+EXTMEM..data: mapped to EXTMEM..V2P(data)
1712 // for the kernel's instructions and r/o data
1713 // data..KERNBASE+PHYSTOP: mapped to V2P(data)..PHYSTOP,
1714 // rw data + free physical memory
1715 // 0xfe000000..0: mapped direct (devices such as ioapic)
1716 //
1717 // The kernel allocates physical memory for its heap and for user memory
1718 // between V2P(end) and the end of physical memory (PHYSTOP)
1719 // (directly addressable from end..P2V(PHYSTOP)).
1720 //
1721 // This table defines the kernel's mappings, which are present in
1722 // every process's page table.
1723 static struct kmap {
1724     void *virt;
1725     uint phys_start;
1726     uint phys_end;
1727     int perm;
1728 } kmap[] = {
1729     { (void*)KERNBASE, 0, EXTMEM, PTE_W}, // I/O space
1730     { (void*)KERNLINK, V2P(KERNEL), 0}, // kern text+rodata
1731     { (void*)data, V2P(data), PHYSTOP, PTE_W}, // kern data+memory
1732     { (void*)DEVSPACE, DEVSPACE, 0, PTE_W}, // more devices
1733 };
1734
1735 // Set up kernel part of a page table.
1736 pde_t*
1737 setupkvm(void)
1738 {
1739     pde_t *pgdir;
1740     struct kmap *k;
1741
1742     if((pgdir = (pde_t*)kalloc()) == 0)
1743         return 0;
1744     memset(pgdir, 0, PGSIZE);
1745     if (p2v(PHYSTOP) > (void*)DEVSPACE)
1746         panic("PHYSTOP too high");
1747     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
1748         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
1749             (uint)k->phys_start, k->perm) < 0)

```

```

1750         return 0;
1751     return pgdir;
1752 }
1753
1754 // Allocate one page table for the machine for the kernel address
1755 // space for scheduler processes.
1756 void
1757 kvmalloc(void)
1758 {
1759     kpgdir = setupkvm();
1760     switchkvm();
1761 }
1762
1763 // Switch h/w page table register to the kernel-only page table,
1764 // for when no process is running.
1765 void
1766 switchkvm(void)
1767 {
1768     lcr3(v2p(kpgdir)); // switch to the kernel page table
1769 }
1770
1771 // Switch TSS and h/w page table to correspond to process p.
1772 void
1773 switchvm(struct proc *p)
1774 {
1775     pushcli();
1776     cpu->gdt[SEG_TSS] = SEG16(STS_T32A, &cpu->ts, sizeof(cpu->ts)-1, 0);
1777     cpu->gdt[SEG_TSS].s = 0;
1778     cpu->ts.ss0 = SEG_KDATA << 3;
1779     cpu->ts.esp0 = (uint)proc->kstack + KSTACKSIZE;
1780     ltr(SEG_TSS << 3);
1781     if(p->pgdir == 0)
1782         panic("switchvm: no pgdir");
1783     lcr3(v2p(p->pgdir)); // switch to new address space
1784     popcli();
1785 }
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799

```

```

1800 // Load the initcode into address 0 of pgdir.
1801 // sz must be less than a page.
1802 void
1803 inituvm(pde_t *pgdir, char *init, uint sz)
1804 {
1805     char *mem;
1806
1807     if(sz >= PGSIZE)
1808         panic("inituvm: more than a page");
1809     mem = kalloc();
1810     memset(mem, 0, PGSIZE);
1811     mappages(pgdir, 0, PGSIZE, v2p(mem), PTE_W|PTE_U);
1812     memmove(mem, init, sz);
1813 }
1814
1815 // Load a program segment into pgdir.  addr must be page-aligned
1816 // and the pages from addr to addr+sz must already be mapped.
1817 int
1818 loaduvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz)
1819 {
1820     uint i, pa, n;
1821     pte_t *pte;
1822
1823     if((uint) addr % PGSIZE != 0)
1824         panic("loaduvm: addr must be page aligned");
1825     for(i = 0; i < sz; i += PGSIZE){
1826         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
1827             panic("loaduvm: address should exist");
1828         pa = PTE_ADDR(*pte);
1829         if(sz - i < PGSIZE)
1830             n = sz - i;
1831         else
1832             n = PGSIZE;
1833         if(readi(ip, p2v(pa), offset+i, n) != n)
1834             return -1;
1835     }
1836     return 0;
1837 }
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849

```

```

1850 // Allocate page tables and physical memory to grow process from oldsz to
1851 // newsz, which need not be page aligned.  Returns new size or 0 on error.
1852 int
1853 allocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1854 {
1855     char *mem;
1856     uint a;
1857
1858     if(newsz >= KERNBASE)
1859         return 0;
1860     if(newsz < oldsz)
1861         return oldsz;
1862
1863     a = PGROUNDUP(oldsz);
1864     for(; a < newsz; a += PGSIZE){
1865         mem = kalloc();
1866         if(mem == 0){
1867             cprintf("allocuvm out of memory\n");
1868             deallocuvm(pgdir, newsz, oldsz);
1869             return 0;
1870         }
1871         memset(mem, 0, PGSIZE);
1872         mappages(pgdir, (char*)a, PGSIZE, v2p(mem), PTE_W|PTE_U);
1873     }
1874     return newsz;
1875 }
1876
1877 // Deallocate user pages to bring the process size from oldsz to
1878 // newsz.  oldsz and newsz need not be page-aligned, nor does newsz
1879 // need to be less than oldsz.  oldsz can be larger than the actual
1880 // process size.  Returns the new process size.
1881 int
1882 deallocuvm(pde_t *pgdir, uint oldsz, uint newsz)
1883 {
1884     pte_t *pte;
1885     uint a, pa;
1886
1887     if(newsz >= oldsz)
1888         return oldsz;
1889
1890     a = PGROUNDUP(newsz);
1891     for(; a < oldsz; a += PGSIZE){
1892         pte = walkpgdir(pgdir, (char*)a, 0);
1893         if(!pte)
1894             a += (NPENTRIES - 1) * PGSIZE;
1895         else if((*pte & PTE_P) != 0){
1896             pa = PTE_ADDR(*pte);
1897             if(pa == 0)
1898                 panic("kfree");
1899             char *v = p2v(pa);

```

```

1900     kfree(v);
1901     *pte = 0;
1902 }
1903 }
1904 return newsz;
1905 }
1906
1907 // Free a page table and all the physical memory pages
1908 // in the user part.
1909 void
1910 freevm(pde_t *pgdir)
1911 {
1912     uint i;
1913
1914     if(pgdir == 0)
1915         panic("freevm: no pgdir");
1916     deallocvm(pgdir, KERNBASE, 0);
1917     for(i = 0; i < NPENTRIES; i++){
1918         if(pgdir[i] & PTE_P){
1919             char * v = p2v(PTE_ADDR(pgdir[i]));
1920             kfree(v);
1921         }
1922     }
1923     kfree((char*)pgdir);
1924 }
1925
1926 // Clear PTE_U on a page. Used to create an inaccessible
1927 // page beneath the user stack.
1928 void
1929 clearpteu(pde_t *pgdir, char *uva)
1930 {
1931     pte_t *pte;
1932
1933     pte = walkpgdir(pgdir, uva, 0);
1934     if(pte == 0)
1935         panic("clearpteu");
1936     *pte &= ~PTE_U;
1937 }
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949

```

```

1950 // Given a parent process's page table, create a copy
1951 // of it for a child.
1952 pde_t*
1953 copyuvm(pde_t *pgdir, uint sz)
1954 {
1955     pde_t *d;
1956     pte_t *pte;
1957     uint pa, i, flags;
1958     char *mem;
1959
1960     if((d = setupkvm()) == 0)
1961         return 0;
1962     for(i = 0; i < sz; i += PGSIZE){
1963         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
1964             panic("copyuvm: pte should exist");
1965         if(!(*pte & PTE_P))
1966             panic("copyuvm: page not present");
1967         pa = PTE_ADDR(*pte);
1968         flags = PTE_FLAGS(*pte);
1969         if((mem = kalloc()) == 0)
1970             goto bad;
1971         memmove(mem, (char*)p2v(pa), PGSIZE);
1972         if(mappages(d, (void*)i, PGSIZE, v2p(mem), flags) < 0)
1973             goto bad;
1974     }
1975     return d;
1976
1977 bad:
1978     freevm(d);
1979     return 0;
1980 }
1981
1982 // Map user virtual address to kernel address.
1983 char*
1984 uva2ka(pde_t *pgdir, char *uva)
1985 {
1986     pte_t *pte;
1987
1988     pte = walkpgdir(pgdir, uva, 0);
1989     if((*pte & PTE_P) == 0)
1990         return 0;
1991     if((*pte & PTE_U) == 0)
1992         return 0;
1993     return (char*)p2v(PTE_ADDR(*pte));
1994 }
1995
1996
1997
1998
1999

```

```

2000 // Copy len bytes from p to user address va in page table pgdir.
2001 // Most useful when pgdir is not the current page table.
2002 // uva2ka ensures this only works for PTE_U pages.
2003 int
2004 copyout(pde_t *pgdir, uint va, void *p, uint len)
2005 {
2006     char *buf, *pa0;
2007     uint n, va0;
2008
2009     buf = (char*)p;
2010     while(len > 0){
2011         va0 = (uint)PGROUNDDOWN(va);
2012         pa0 = uva2ka(pgdir, (char*)va0);
2013         if(pa0 == 0)
2014             return -1;
2015         n = PGSIZE - (va - va0);
2016         if(n > len)
2017             n = len;
2018         memmove(pa0 + (va - va0), buf, n);
2019         len -= n;
2020         buf += n;
2021         va = va0 + PGSIZE;
2022     }
2023     return 0;
2024 }
2025
2026 // Blank page.
2027 // Blank page.
2028 // Blank page.
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049

```

```

2050 // Segments in proc->gdt.
2051 #define NSEGS      7
2052
2053 // Default UID and GID for init
2054 #define INITUID     0
2055 #define INITGID     0
2056
2057 // Default number of ready processes list
2058 #define NUM_READY_LISTS 7
2059 // Default starting priority number
2060 #define PRIORITY_HIGH 0
2061 // Default lowest priority number
2062 #define PRIORITY_LOW  PRIORITY_HIGH+NUM_READY_LISTS-1
2063 // Default promotion interval
2064 #define TICKS_TO_PROMOTE 200
2065 // Default process budget
2066 #define BUDGET 400
2067
2068 // Per-CPU state
2069 struct cpu {
2070     uchar id; // Local APIC ID; index into cpus[] below
2071     struct context *scheduler; // swtch() here to enter scheduler
2072     struct taskstate ts; // Used by x86 to find stack for interrupt
2073     struct segdesc gdt[NSEGS]; // x86 global descriptor table
2074     volatile uint started; // Has the CPU started?
2075     int ncli; // Depth of pushcli nesting.
2076     int intena; // Were interrupts enabled before pushcli?
2077
2078     // Cpu-local storage variables; see below
2079     struct cpu *cpu;
2080     struct proc *proc; // The currently-running process.
2081 };
2082
2083 extern struct cpu cpus[NCPU];
2084 extern int ncpu;
2085
2086 // Per-CPU variables, holding pointers to the
2087 // current cpu and to the current process.
2088 // The asm suffix tells gcc to use "%gs:0" to refer to cpu
2089 // and "%gs:4" to refer to proc. seginit sets up the
2090 // %gs segment register so that %gs refers to the memory
2091 // holding those two variables in the local cpu's struct cpu.
2092 // This is similar to how thread-local variables are implemented
2093 // in thread libraries such as Linux pthreads.
2094 extern struct cpu *cpu asm("%gs:0"); // &cpus[cpunum()]
2095 extern struct proc *proc asm("%gs:4"); // cpus[cpunum()].proc
2096
2097
2098
2099

```

```

2100 // Saved registers for kernel context switches.
2101 // Don't need to save all the segment registers (%cs, etc),
2102 // because they are constant across kernel contexts.
2103 // Don't need to save %eax, %ecx, %edx, because the
2104 // x86 convention is that the caller has saved them.
2105 // Contexts are stored at the bottom of the stack they
2106 // describe; the stack pointer is the address of the context.
2107 // The layout of the context matches the layout of the stack in swtch.S
2108 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,
2109 // but it is on the stack and allocproc() manipulates it.
2110 struct context {
2111     uint edi;
2112     uint esi;
2113     uint ebx;
2114     uint ebp;
2115     uint eip;
2116 };
2117
2118 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
2119
2120 // Per-process state
2121 struct proc {
2122     uint sz; // Size of process memory (bytes)
2123     pde_t* pgdir; // Page table
2124     char *kstack; // Bottom of kernel stack for this process
2125     enum procstate state; // Process state
2126     uint pid; // Process ID
2127     struct proc *parent; // Parent process
2128     struct trapframe *tf; // Trap frame for current syscall
2129     struct context *context; // swtch() here to run process
2130     void *chan; // If non-zero, sleeping on chan
2131     int killed; // If non-zero, have been killed
2132     struct file *ofile[NOFILE]; // Open files
2133     struct inode *cwd; // Current directory
2134     char name[16]; // Process name (debugging)
2135     uint start_ticks; // Start ticks (debugging)
2136 #ifdef CS333_P2
2137     uint cpu_ticks_total; // Total elapsed ticks in process
2138     uint cpu_ticks_in; // Ticks when scheduled
2139     uint uid; // Process owner's user id
2140     uint gid; // Process owner's group id
2141 #endif
2142
2143 #ifdef CS333_P3
2144     int priority; // Process priority
2145     int budget; // A process's time budget
2146     struct proc *next; // Next process in the process list
2147 #endif
2148 };
2149

```

```

2150 // Process memory is laid out contiguously, low addresses first:
2151 //   text
2152 //   original data and bss
2153 //   fixed-size stack
2154 //   expandable heap
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199

```

```

2200 #include "types.h"
2201 #include "defs.h"
2202 #include "param.h"
2203 #include "memlayout.h"
2204 #include "mmu.h"
2205 #include "x86.h"
2206 #include "proc.h"
2207 #include "spinlock.h"
2208 #include "uproc.h"
2209
2210 struct {
2211   struct spinlock lock;
2212   struct proc proc[NPROC];
2213 #ifdef CS333_P3
2214   struct proc *pReadyList[NUM_READY_LISTS];
2215   struct proc *pFreeList;
2216   uint PromoteAtTime;
2217 #endif
2218 } ptable;
2219
2220 static struct proc *initproc;
2221
2222 int nextpid = 1;
2223 extern void forkret(void);
2224 extern void trapret(void);
2225
2226 static void wakeup1(void *chan);
2227
2228 void
2229 pinit(void)
2230 {
2231   initlock(&ptable.lock, "ptable");
2232 }
2233
2234 #ifdef CS333_P3
2235 // Pops a process off a process queue
2236 // Return -1 if no process in the queue
2237 static struct proc*
2238 popq(struct proc **proclst)
2239 {
2240   if(!holding(&ptable.lock))
2241     panic("popq ptable.lock\n");
2242   if(proclst <= 0 || *proclst <= 0) return 0;
2243   struct proc *ret;
2244   ret = *proclst;
2245   *proclst = (*proclst)->next;
2246   ret->next = 0;
2247   return ret;
2248 }
2249
```

```

2250 // Pushs a process to the pFreeList
2251 static void
2252 pushfreeq(struct proc* input, struct proc **freelist)
2253 {
2254   if(!holding(&ptable.lock))
2255     panic("pushfreeq ptable.lock\n");
2256   else {
2257     input->next = *freelist;
2258     *freelist = input;
2259   }
2260 }
2261
2262 // Pushs a process to the pReadyList
2263 static void
2264 pushreadyq(struct proc* input, struct proc **readylist)
2265 {
2266   if(!holding(&ptable.lock))
2267     panic("pushreadyq ptable.lock\n");
2268   if(!input)
2269     return;
2270   if(!*readylist) {
2271     input->next = 0;
2272     *readylist = input;
2273   }
2274   else {
2275     struct proc* temp = *readylist;
2276     while(temp->next)
2277       temp = temp->next;
2278     temp->next = input;
2279     input->next = 0;
2280   }
2281 }
2282
2283 // Set process's priority to specified value
2284 // Return 0 if success
2285 // Assumes holding ptable lock
2286 int
2287 setpriority(int pid, int priority)
2288 {
2289   if(pid < 0)
2290     panic("pid out of bound\n");
2291   if(priority < PRIORITY_HIGH || priority > PRIORITY_LOW) {
2292     cprintf("Invalid priority value: %d, need an int between %d and %d\n",
2293             priority, PRIORITY_LOW, PRIORITY_HIGH);
2294     return -1;
2295   }
2296   acquire(&ptable.lock);
2297   struct proc *p;
2298   for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2299     if(p->pid == pid) {

```

```

2300         p->priority = priority;
2301         p->budget = BUDGET;
2302         release(&ptable.lock);
2303         return 0;
2304     }
2305     cprintf("Invalid pid: %d\n", pid);
2306     release(&ptable.lock);
2307     return -1;
2308 }
2309 #endif
2310
2311 // Look in the process table for an UNUSED proc.
2312 // If found, change state to EMBRYO and initialize
2313 // state required to run in the kernel.
2314 // Otherwise return 0.
2315 // static struct proc*
2316 // allocproc(void)
2317 {
2318     struct proc *p;
2319     char *sp;
2320
2321 #ifndef CS333_P3
2322     acquire(&ptable.lock);
2323     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2324         if(p->state == UNUSED)
2325             goto found;
2326     release(&ptable.lock);
2327 #else
2328     acquire(&ptable.lock);
2329     p = popq(&ptable.pFreeList);
2330     if(p && p->state == UNUSED)
2331         goto found;
2332     release(&ptable.lock);
2333 #endif
2334     return 0;
2335
2336 found:
2337     p->state = EMBRYO;
2338     p->pid = nextpid++;
2339     release(&ptable.lock);
2340
2341     // Allocate kernel stack.
2342     if((p->kstack = kalloc()) == 0){
2343         p->state = UNUSED;
2344         acquire(&ptable.lock);
2345         pushfreeq(p, &ptable.pFreeList);
2346         release(&ptable.lock);
2347     }
2348     return 0;

```

```

2350     }
2351     sp = p->kstack + KSTACKSIZE;
2352
2353     // Leave room for trap frame.
2354     sp -= sizeof *p->tf;
2355     p->tf = (struct trapframe*)sp;
2356
2357     // Set up new context to start executing at forkret,
2358     // which returns to trapret.
2359     sp -= 4;
2360     *(uint*)sp = (uint)trapret;
2361
2362     sp -= sizeof *p->context;
2363     p->context = (struct context*)sp;
2364     memset(p->context, 0, sizeof *p->context);
2365     p->context->eip = (uint)forkret;
2366
2367     acquire(&tickslock);
2368     p->start_ticks = ticks;
2369     release(&tickslock);
2370     p->cpu_ticks_in = 0;
2371
2372     return p;
2373 }
2374
2375 // Check if it's time to promote
2376 // Assume always hold the lock
2377 // return 1 if it's time to promote
2378 #ifdef CS333_P3
2379 static int
2380 timetopromote(void)
2381 {
2382     if(!holding(&ptable.lock))
2383         panic("timetopromote ptable.lock");
2384     acquire(&tickslock);
2385     if(ticks < ptable.PromoteAtTime) {
2386         release(&tickslock);
2387         return 0; // Not time to promote
2388     }
2389     ptable.PromoteAtTime = ticks + TICKS_TO_PROMOTE;
2390     release(&tickslock);
2391     return 1;
2392 }
2393
2394
2395
2396
2397
2398
2399

```



```

2400 #endif
2401
2402 // Set up first user process.
2403 void
2404 userinit(void)
2405 {
2406     struct proc *p;
2407     extern char _binary_initcode_start[], _binary_initcode_size[];
2408
2409 #ifdef CS333_P3
2410     acquire(&ptable.lock);
2411     ptable.pFreeList = 0;
2412     // Initialize freelist by putting UNUSED processes to the list
2413     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2414         if(p->state == UNUSED)
2415             pushfreeq(p, &ptable.pFreeList);
2416     // Initialize readylist to empty
2417     int i;
2418     for(i = PRIORITY_HIGH; i < NUM_READY_LISTS; ++i) {
2419         ptable.pReadyList[i] = 0;
2420     }
2421     release(&ptable.lock);
2422 #endif
2423
2424     p = allocproc();
2425     initproc = p;
2426     if((p->pgdir = setupkvm()) == 0)
2427         panic("userinit: out of memory?");
2428     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
2429     p->sz = PGSIZE;
2430     memset(p->tf, 0, sizeof(*p->tf));
2431     p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
2432     p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
2433     p->tf->es = p->tf->ds;
2434     p->tf->ss = p->tf->ds;
2435     p->tf->eflags = FL_IF;
2436     p->tf->esp = PGSIZE;
2437     p->tf->eip = 0; // beginning of initcode.S
2438
2439     safestrcpy(p->name, "initcode", sizeof(p->name));
2440     p->cwd = namei("/");
2441
2442     p->state = RUNNABLE;
2443 #ifdef CS333_P3
2444     acquire(&ptable.lock);
2445     ptable.pReadyList[PRIORITY_HIGH] = p;
2446     p->next = 0;
2447     release(&ptable.lock);
2448 #endif
2449

```

```

2450 #ifdef CS333_P2
2451     p->uid = INITUID;
2452     p->gid = INITGID;
2453 #endif
2454 }
2455
2456 // Grow current process's memory by n bytes.
2457 // Return 0 on success, -1 on failure.
2458 int
2459 growproc(int n)
2460 {
2461     uint sz;
2462
2463     sz = proc->sz;
2464     if(n > 0){
2465         if((sz = allocuvm(proc->pgdir, sz, sz + n)) == 0)
2466             return -1;
2467     } else if(n < 0){
2468         if((sz = deallocuvm(proc->pgdir, sz, sz + n)) == 0)
2469             return -1;
2470     }
2471     proc->sz = sz;
2472     switchuvm(proc);
2473     return 0;
2474 }
2475
2476 // Create a new process copying p as the parent.
2477 // Sets up stack to return as if from system call.
2478 // Caller must set state of returned proc to RUNNABLE.
2479 int
2480 fork(void)
2481 {
2482     int i, pid;
2483     struct proc *np;
2484
2485     // Allocate process.
2486     if((np = allocproc()) == 0)
2487         return -1;
2488
2489     // Copy process state from p.
2490     if((np->pgdir = copyuvm(proc->pgdir, proc->sz)) == 0){
2491         kfree(np->kstack);
2492         np->kstack = 0;
2493         np->state = UNUSED;
2494 #ifdef CS333_P3
2495         pushfreeq(np, &ptable.pFreeList);
2496 #endif
2497         return -1;
2498     }
2499     np->sz = proc->sz;

```

```

2500 np->parent = proc;
2501 *np->tf = *proc->tf;
2502
2503 #ifdef CS333_P2
2504 // Copy process UID, GID
2505 np->uid = proc->uid;
2506 np->gid = proc->gid;
2507 #endif
2508
2509 // Clear %eax so that fork returns 0 in the child.
2510 np->tf->eax = 0;
2511
2512 for(i = 0; i < NOFILE; i++)
2513     if(proc->ofile[i])
2514         np->ofile[i] = filedup(proc->ofile[i]);
2515 np->cwd = idup(proc->cwd);
2516
2517 safestrcpy(np->name, proc->name, sizeof(proc->name));
2518
2519 pid = np->pid;
2520
2521 // lock to force the compiler to emit the np->state write last.
2522 acquire(&ptable.lock);
2523 np->state = RUNNABLE;
2524 #ifdef CS333_P3
2525 np->priority = PRIORITY_HIGH;
2526 np->budget = BUDGET;
2527 pushreadyq(np, &ptable.pReadyList[PRIORITY_HIGH]);
2528 #endif
2529 release(&ptable.lock);
2530
2531 return pid;
2532 }
2533
2534 // Exit the current process. Does not return.
2535 // An exited process remains in the zombie state
2536 // until its parent calls wait() to find out it exited.
2537 void
2538 exit(void)
2539 {
2540     struct proc *p;
2541     int fd;
2542
2543     if(proc == initproc)
2544         panic("init exiting");
2545
2546
2547
2548
2549

```

```

2550 // Close all open files.
2551 for(fd = 0; fd < NOFILE; fd++){
2552     if(proc->ofile[fd]){
2553         fileclose(proc->ofile[fd]);
2554         proc->ofile[fd] = 0;
2555     }
2556 }
2557
2558 begin_op();
2559 iput(proc->cwd);
2560 end_op();
2561 proc->cwd = 0;
2562
2563 acquire(&ptable.lock);
2564
2565 // Parent might be sleeping in wait().
2566 wakeup1(proc->parent);
2567
2568 // Pass abandoned children to init.
2569 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2570     if(p->parent == proc){
2571         p->parent = initproc;
2572         if(p->state == ZOMBIE)
2573             wakeup1(initproc);
2574     }
2575 }
2576
2577 // Jump into the scheduler, never to return.
2578 proc->state = ZOMBIE;
2579 sched();
2580 panic("zombie exit");
2581 }
2582
2583 // Wait for a child process to exit and return its pid.
2584 // Return -1 if this process has no children.
2585 int
2586 wait(void)
2587 {
2588     struct proc *p;
2589     int havekids, pid;
2590
2591     acquire(&ptable.lock);
2592     for(;;){
2593         // Scan through table looking for zombie children.
2594         havekids = 0;
2595         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2596             if(p->parent != proc)
2597                 continue;
2598             havekids = 1;
2599             if(p->state == ZOMBIE){

```

```

2600     // Found one.
2601     pid = p->pid;
2602     kfree(p->kstack);
2603     p->kstack = 0;
2604     freevm(p->pgdir);
2605     p->state = UNUSED;
2606 #ifdef CS333_P3
2607         pushfreeq(p, &ptable.pFreeList);
2608 #endif
2609     p->pid = 0;
2610     p->parent = 0;
2611     p->name[0] = 0;
2612     p->killed = 0;
2613     release(&ptable.lock);
2614     return pid;
2615 }
2616 }
2617
2618 // No point waiting if we don't have any children.
2619 if(!havekids || proc->killed){
2620     release(&ptable.lock);
2621     return -1;
2622 }
2623
2624 // Wait for children to exit. (See wakeup1 call in proc_exit.)
2625 sleep(proc, &ptable.lock);
2626 }
2627 }
2628
2629 // Per-CPU process scheduler.
2630 // Each CPU calls scheduler() after setting itself up.
2631 // Scheduler never returns. It loops, doing:
2632 //  - choose a process to run
2633 //  - swtch to start running that process
2634 //  - eventually that process transfers control
2635 //    via swtch back to the scheduler.
2636 #ifndef CS333_P3
2637 // original xv6 scheduler. Use if CS333_P3 NOT defined.
2638 void
2639 scheduler(void)
2640 {
2641     struct proc *p;
2642
2643     for(;;){
2644         // Enable interrupts on this processor.
2645         sti();
2646
2647
2648
2649

```

```

2650     // Loop over process table looking for process to run.
2651     acquire(&ptable.lock);
2652     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2653         if(p->state != RUNNABLE)
2654             continue;
2655
2656         // Switch to chosen process. It is the process's job
2657         // to release ptable.lock and then reacquire it
2658         // before jumping back to us.
2659         proc = p;
2660         switchvm(p);
2661         p->state = RUNNING;
2662 #ifdef CS333_P2
2663         acquire(&tickslock);
2664         p->cpu_ticks_in = ticks;
2665         release(&tickslock);
2666 #endif
2667         swtch(&cpu->scheduler, proc->context);
2668         switchkvm();
2669
2670         // Process is done running for now.
2671         // It should have changed its p->state before coming back.
2672         proc = 0;
2673     }
2674     release(&ptable.lock);
2675
2676 }
2677 }
2678
2679 #else
2680 // CS333_P3 MLFQ scheduler implementation goes here
2681 void
2682 scheduler(void)
2683 {
2684     struct proc *p;
2685
2686     for(;;){
2687         // Enable interrupts on this processor.
2688         sti();
2689
2690         // If promotion timer expires promote all processes one
2691         // level up
2692         acquire(&ptable.lock);
2693         if(timetopromote()) {
2694             // Increase priority for Running, sleeping processes
2695             for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2696                 if(p->priority <= PRIORITY_HIGH)
2697                     continue;
2698                 if(p->state == RUNNING || p->state == SLEEPING){
2699                     p->budget = BUDGET;

```

```

2700         p->priority -= 1;
2701     }
2702 }
2703 // Priority queue shift up
2704 int priority;
2705 for(priority = PRIORITY_HIGH; priority < PRIORITY_LOW; ++priority)
2706     cprintf("time to promote\n");
2707     do {
2708         p = popq(&ptable.pReadyList[priority+1]);
2709         if(p) {
2710             p->priority -= 1;
2711             p->budget = BUDGET;
2712             pushreadyq(p, &ptable.pReadyList[priority]);
2713         }
2714     }while(p);
2715 }
2716 }
2717
2718 // Find the next runnable process and pop it from the ready
2719 // list
2720 int i;
2721 for(i = PRIORITY_HIGH; i < PRIORITY_LOW+1; i++) {
2722     if(!ptable.pReadyList[i])
2723         ++i;
2724     continue;
2725 }
2726 p = popq(&ptable.pReadyList[i]);
2727
2728 if(!p) {
2729     panic("popping an empty readylist");
2730 }
2731
2732 // Switch to chosen process. It is the process's job
2733 // to release ptable.lock and then reacquire it
2734 // before jumping back to us.
2735 proc = p;
2736 switchvm(p);
2737 p->state = RUNNING;
2738 #ifdef CS333_P2
2739     acquire(&tickslock);
2740     p->cpu_ticks_in = ticks;
2741     release(&tickslock);
2742 #endif
2743 swtch(&cpu->scheduler, proc->context);
2744 switchkvm();
2745
2746
2747
2748
2749

```

```

2750         // Process is done running for now.
2751         // It should have changed its p->state before coming back
2752         proc = 0;
2753     }
2754     release(&ptable.lock);
2755 }
2756 }
2757 #endif
2758
2759 // Enter scheduler. Must hold only ptable.lock
2760 // and have changed proc->state.
2761 void
2762 sched(void)
2763 {
2764     int intena;
2765
2766     if(!holding(&ptable.lock))
2767         panic("sched ptable.lock");
2768     if(cpu->ncli != 1)
2769         panic("sched locks");
2770     if(proc->state == RUNNING)
2771         panic("sched running");
2772     if(readeflags() & FL_IF)
2773         panic("sched interrible");
2774     intena = cpu->intena;
2775 #ifdef CS333_P2
2776     acquire(&tickslock);
2777     proc->cpu_ticks_total += ticks - proc->cpu_ticks_in;
2778 #ifdef CS333_P3
2779     proc->budget -= ticks - proc->cpu_ticks_in;
2780 #endif
2781     release(&tickslock);
2782 #endif
2783
2784 #ifdef CS333_P3
2785     // Check process's budget if its <= 0
2786     // demote to the next lower priority queue
2787     // else add it to the back of current queue.
2788     if(proc->budget <= 0 && proc->priority < PRIORITY_LOW){
2789         proc->priority += 1;
2790         proc->budget = BUDGET;
2791     }
2792     if(proc->state == RUNNABLE)
2793         pushreadyq(proc, &ptable.pReadyList[proc->priority]);
2794 #endif
2795     swtch(&proc->context, cpu->scheduler);
2796     cpu->intena = intena;
2797 }
2798
2799

```

```

2800 // Give up the CPU for one scheduling round.
2801 void
2802 yield(void)
2803 {
2804     acquire(&ptable.lock);
2805     proc->state = RUNNABLE;
2806     sched();
2807     release(&ptable.lock);
2808 }
2809
2810 // A fork child's very first scheduling by scheduler()
2811 // will swtch here. "Return" to user space.
2812 void
2813 forkret(void)
2814 {
2815     static int first = 1;
2816     // Still holding ptable.lock from scheduler.
2817     release(&ptable.lock);
2818
2819     if (first) {
2820         // Some initialization functions must be run in the context
2821         // of a regular process (e.g., they call sleep), and thus cannot
2822         // be run from main().
2823         first = 0;
2824         iinit(ROOTDEV);
2825         initlog(ROOTDEV);
2826     }
2827
2828     // Return to "caller", actually trapret (see allocproc).
2829 }
2830
2831 // Atomically release lock and sleep on chan.
2832 // Reacquires lock when awakened.
2833 void
2834 sleep(void *chan, struct spinlock *lk)
2835 {
2836     if (proc == 0)
2837         panic("sleep");
2838
2839     if (lk == 0)
2840         panic("sleep without lk");
2841
2842     // Must acquire ptable.lock in order to
2843     // change p->state and then call sched.
2844     // Once we hold ptable.lock, we can be
2845     // guaranteed that we won't miss any wakeup
2846     // (wakeup runs with ptable.lock locked),
2847     // so it's okay to release lk.
2848     if (lk != &ptable.lock) {
2849         acquire(&ptable.lock);

```

```

2850     release(lk);
2851 }
2852
2853 // Go to sleep.
2854 proc->chan = chan;
2855 proc->state = SLEEPING;
2856 sched();
2857
2858 // Tidy up.
2859 proc->chan = 0;
2860
2861 // Reacquire original lock.
2862 if (lk != &ptable.lock) {
2863     release(&ptable.lock);
2864     acquire(lk);
2865 }
2866 }
2867
2868 // Wake up all processes sleeping on chan.
2869 // The ptable lock must be held.
2870 static void
2871 wakeup1(void *chan)
2872 {
2873     struct proc *p;
2874
2875     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2876 #ifndef CS333_P3
2877         if (p->state == SLEEPING && p->chan == chan)
2878             p->state = RUNNABLE;
2879 #else
2880         if (p->state == SLEEPING && p->chan == chan) {
2881             p->state = RUNNABLE;
2882             pushreadyq(p, &ptable.pReadyList[p->priority]);
2883         }
2884 #endif
2885 }
2886
2887 // Wake up all processes sleeping on chan.
2888 void
2889 wakeup(void *chan)
2890 {
2891     acquire(&ptable.lock);
2892     wakeup1(chan);
2893     release(&ptable.lock);
2894 }
2895
2896
2897
2898
2899

```

```

2900 // Kill the process with the given pid.
2901 // Process won't exit until it returns
2902 // to user space (see trap in trap.c).
2903 int
2904 kill(int pid)
2905 {
2906     struct proc *p;
2907
2908     acquire(&ptable.lock);
2909     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2910         if(p->pid == pid){
2911             p->killed = 1;
2912             // Wake process from sleep if necessary.
2913 #ifndef CS333_P3
2914             if(p->state == SLEEPING)
2915                 p->state = RUNNABLE;
2916 #else
2917             if(p->state == SLEEPING) {
2918                 p->state = RUNNABLE;
2919                 pushreadyq(p, &ptable.pReadyList[p->priority]);
2920             }
2921 #endif
2922             release(&ptable.lock);
2923             return 0;
2924         }
2925     }
2926     release(&ptable.lock);
2927     return -1;
2928 }
2929
2930
2931 // Print a process listing to console. For debugging.
2932 // Runs when user types ^P on console.
2933 // No lock to avoid wedging a stuck machine further.
2934 static void
2935 print_elapsed(struct proc *p)
2936 {
2937     uint temp = p->start_ticks;
2938     temp = ticks - temp;
2939     cprintf("%d.%d", temp/100, temp%100);
2940 #ifdef CS333_P2
2941     cprintf(" %d.%d", p->cpu_ticks_total/100, p->cpu_ticks_total%100);
2942     cprintf(" %d ", p->uid);
2943     cprintf(" %d ", p->gid);
2944     if(p->parent && p->pid != 1)
2945         cprintf(" %d ", p->parent->pid);
2946     else
2947         cprintf(" %d ", p->pid);
2948 #ifdef CS333_P3
2949     cprintf(" %d ", p->priority);

```

```

2950 #endif
2951 #endif
2952 }
2953
2954 void
2955 procdump(void)
2956 {
2957     static char *states[] = {
2958         [UNUSED]    "unused",
2959         [EMBRYO]    "embryo",
2960         [SLEEPING]  "sleep ",
2961         [RUNNABLE]  "runble",
2962         [RUNNING]   "run   ",
2963         [ZOMBIE]    "zombie"
2964     };
2965     int i;
2966     struct proc *p;
2967     char *state;
2968     uint pc[10];
2969
2970 #ifdef CS333_P3
2971     cprintf("\nPID State Name Elapsed TotalCpuTime UID GID PPID\n");
2972 #else
2973 #ifdef CS333_P2
2974     cprintf("\nPID State Name Elapsed TotalCpuTime UID GID PPID\n");
2975 #else
2976     cprintf("\nPID State Name Elapsed PCs\n");
2977 #endif
2978 #endif
2979
2980     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2981         if(p->state == UNUSED)
2982             continue;
2983         if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
2984             state = states[p->state];
2985         else
2986             state = "???";
2987         cprintf("%d %s %s ", p->pid, state, p->name);
2988         print_elapsed(p);
2989         if(p->state == SLEEPING){
2990             getcallerpcs((uint*)p->context->ebp+2, pc);
2991             for(i=0; i<10 && pc[i] != 0; i++)
2992                 cprintf(" %p", pc[i]);
2993         }
2994         cprintf("\n");
2995     }
2996 }
2997
2998
2999

```

```

3000 #ifdef CS333_P2
3001 // Get process information
3002 int
3003 getprocs(uint max, struct uproc* table)
3004 {
3005     if(!table || max == 0) return -1;
3006     static char *states[] = {
3007         [UNUSED]    "unused",
3008         [EMBRYO]     "embryo",
3009         [SLEEPING]   "sleep ",
3010         [RUNNABLE]   "runble ",
3011         [RUNNING]    "run   ",
3012         [ZOMBIE]     "zombie"
3013     };
3014
3015     int procscount = 0;
3016     struct proc *p;
3017     if(max > NPROC)
3018         max = NPROC;
3019     acquire(&table.lock);
3020     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
3021         if(max <= 0) break; // break out of the loop if the max number of
3022         if(p->state == UNUSED || p->state == EMBRYO || p->state == ZOMBIE)
3023             continue;
3024         table->pid = p->pid;
3025         table->uid = p->uid;
3026         table->gid = p->gid;
3027         if(!p->parent || p->pid == 1)
3028             table->ppid = p->pid;
3029         else
3030             table->ppid = p->parent->pid;
3031 #ifdef CS333_P3
3032         table->priority = p->priority;
3033 #endif
3034         acquire(&tickslock);
3035         table->elapsed_ticks = ticks - p->start_ticks;
3036         table->CPU_total_ticks = p->cpu_ticks_total;
3037         release(&tickslock);
3038         safestrcpy(table->state, states[p->state], sizeof(table->state));
3039         table->size = p->sz;
3040         safestrcpy(table->name, p->name, sizeof(table->name));
3041         ++procscount;
3042         ++table;
3043         --max;
3044     }
3045     release(&table.lock);
3046
3047     return procscount;
3048 }
3049 #endif

```

```

3050 # Context switch
3051 #
3052 # void swtch(struct context **old, struct context *new);
3053 #
3054 # Save current register context in old
3055 # and then load register context from new.
3056
3057 .globl swtch
3058 swtch:
3059     movl 4(%esp), %eax
3060     movl 8(%esp), %edx
3061
3062     # Save old callee-save registers
3063     pushl %ebp
3064     pushl %ebx
3065     pushl %esi
3066     pushl %edi
3067
3068     # Switch stacks
3069     movl %esp, (%eax)
3070     movl %edx, %esp
3071
3072     # Load new callee-save registers
3073     popl %edi
3074     popl %esi
3075     popl %ebx
3076     popl %ebp
3077     ret
3078
3079
3080
3081
3082
3083
3084
3085
3086
3087
3088
3089
3090
3091
3092
3093
3094
3095
3096
3097
3098
3099

```

```

3100 // Physical memory allocator, intended to allocate
3101 // memory for user processes, kernel stacks, page table pages,
3102 // and pipe buffers. Allocates 4096-byte pages.
3103
3104 #include "types.h"
3105 #include "defs.h"
3106 #include "param.h"
3107 #include "memlayout.h"
3108 #include "mmu.h"
3109 #include "spinlock.h"
3110
3111 void freerange(void *vstart, void *vend);
3112 extern char end[]; // first address after kernel loaded from ELF file
3113
3114 struct run {
3115     struct run *next;
3116 };
3117
3118 struct {
3119     struct spinlock lock;
3120     int use_lock;
3121     struct run *freelist;
3122 } kmem;
3123
3124 // Initialization happens in two phases.
3125 // 1. main() calls kinit1() while still using entrypdir to place just
3126 // the pages mapped by entrypdir on free list.
3127 // 2. main() calls kinit2() with the rest of the physical pages
3128 // after installing a full page table that maps them on all cores.
3129 void
3130 kinit1(void *vstart, void *vend)
3131 {
3132     initlock(&kmem.lock, "kmem");
3133     kmem.use_lock = 0;
3134     freerange(vstart, vend);
3135 }
3136
3137 void
3138 kinit2(void *vstart, void *vend)
3139 {
3140     freerange(vstart, vend);
3141     kmem.use_lock = 1;
3142 }
3143
3144
3145
3146
3147
3148
3149

```

```

3150 void
3151 freerange(void *vstart, void *vend)
3152 {
3153     char *p;
3154     p = (char*)PGROUNDUP((uint)vstart);
3155     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
3156         kfree(p);
3157 }
3158
3159 // Free the page of physical memory pointed at by v,
3160 // which normally should have been returned by a
3161 // call to kalloc(). (The exception is when
3162 // initializing the allocator; see kinit above.)
3163 void
3164 kfree(char *v)
3165 {
3166     struct run *r;
3167
3168     if((uint)v % PGSIZE || v < end || v2p(v) >= PHYSTOP)
3169         panic("kfree");
3170
3171     // Fill with junk to catch dangling refs.
3172     memset(v, 1, PGSIZE);
3173
3174     if(kmem.use_lock)
3175         acquire(&kmem.lock);
3176     r = (struct run*)v;
3177     r->next = kmem.freelist;
3178     kmem.freelist = r;
3179     if(kmem.use_lock)
3180         release(&kmem.lock);
3181 }
3182
3183 // Allocate one 4096-byte page of physical memory.
3184 // Returns a pointer that the kernel can use.
3185 // Returns 0 if the memory cannot be allocated.
3186 char*
3187 kalloc(void)
3188 {
3189     struct run *r;
3190
3191     if(kmem.use_lock)
3192         acquire(&kmem.lock);
3193     r = kmem.freelist;
3194     if(r)
3195         kmem.freelist = r->next;
3196     if(kmem.use_lock)
3197         release(&kmem.lock);
3198     return (char*)r;
3199 }

```



```

3200 // x86 trap and interrupt constants.
3201
3202 // Processor-defined:
3203 #define T_DIVIDE      0      // divide error
3204 #define T_DEBUG      1      // debug exception
3205 #define T_NMI        2      // non-maskable interrupt
3206 #define T_BRKPT      3      // breakpoint
3207 #define T_OFLOW      4      // overflow
3208 #define T_BOUND      5      // bounds check
3209 #define T_ILLOP      6      // illegal opcode
3210 #define T_DEVICE      7      // device not available
3211 #define T_DBLFLT      8      // double fault
3212 // #define T_COPROC    9      // reserved (not used since 486)
3213 #define T_TSS        10     // invalid task switch segment
3214 #define T_SEGNP      11     // segment not present
3215 #define T_STACK      12     // stack exception
3216 #define T_GPFLT      13     // general protection fault
3217 #define T_PGFLT      14     // page fault
3218 // #define T_RES       15     // reserved
3219 #define T_FPERR      16     // floating point error
3220 #define T_ALIGN      17     // alignment check
3221 #define T_MCHK       18     // machine check
3222 #define T_SIMDERR     19     // SIMD floating point error
3223
3224 // These are arbitrarily chosen, but with care not to overlap
3225 // processor defined exceptions or interrupt vectors.
3226 #define T_SYSCALL     64     // system call
3227 #define T_DEFAULT     500    // catchall
3228
3229 #define T_IRQ0        32     // IRQ 0 corresponds to int T_IRQ
3230
3231 #define IRQ_TIMER      0
3232 #define IRQ_KBD        1
3233 #define IRQ_COM1       4
3234 #define IRQ_IDE        14
3235 #define IRQ_ERROR      19
3236 #define IRQ_SPURIOUS   31
3237
3238
3239
3240
3241
3242
3243
3244
3245
3246
3247
3248
3249

```

```

3250 #!/usr/bin/perl -w
3251
3252 # Generate vectors.S, the trap/interrupt entry points.
3253 # There has to be one entry point per interrupt number
3254 # since otherwise there's no way for trap() to discover
3255 # the interrupt number.
3256
3257 print "# generated by vectors.pl - do not edit\n";
3258 print "# handlers\n";
3259 print ".globl alltraps\n";
3260 for(my $i = 0; $i < 256; $i++){
3261     print ".globl vector$i\n";
3262     print "vector$i:\n";
3263     if(!($i == 8 || ($i >= 10 && $i <= 14) || $i == 17)){
3264         print "    pushl $0\n";
3265     }
3266     print "    pushl $$i\n";
3267     print "    jmp alltraps\n";
3268 }
3269
3270 print "\n# vector table\n";
3271 print ".data\n";
3272 print ".globl vectors\n";
3273 print "vectors:\n";
3274 for(my $i = 0; $i < 256; $i++){
3275     print "    .long vector$i\n";
3276 }
3277
3278 # sample output:
3279 # # handlers
3280 # .globl alltraps
3281 # .globl vector0
3282 # vector0:
3283 #     pushl $0
3284 #     pushl $0
3285 #     jmp alltraps
3286 # ...
3287 #
3288 # # vector table
3289 # .data
3290 # .globl vectors
3291 # vectors:
3292 #     .long vector0
3293 #     .long vector1
3294 #     .long vector2
3295 # ...
3296
3297
3298
3299

```

```

3300 #include "mmu.h"
3301
3302 # vectors.S sends all traps here.
3303 .globl alltraps
3304 alltraps:
3305 # Build trap frame.
3306 pushl %ds
3307 pushl %es
3308 pushl %fs
3309 pushl %gs
3310 pushal
3311
3312 # Set up data and per-cpu segments.
3313 movw $(SEG_KDATA<<3), %ax
3314 movw %ax, %ds
3315 movw %ax, %es
3316 movw $(SEG_KCPU<<3), %ax
3317 movw %ax, %fs
3318 movw %ax, %gs
3319
3320 # Call trap(tf), where tf=%esp
3321 pushl %esp
3322 call trap
3323 addl $4, %esp
3324
3325 # Return falls through to trapret...
3326 .globl trapret
3327 trapret:
3328 popal
3329 popl %gs
3330 popl %fs
3331 popl %es
3332 popl %ds
3333 addl $0x8, %esp # trapno and errcode
3334 iret
3335
3336
3337
3338
3339
3340
3341
3342
3343
3344
3345
3346
3347
3348
3349

```

```

3350 #include "types.h"
3351 #include "defs.h"
3352 #include "param.h"
3353 #include "memlayout.h"
3354 #include "mmu.h"
3355 #include "proc.h"
3356 #include "x86.h"
3357 #include "traps.h"
3358 #include "spinlock.h"
3359
3360 // Interrupt descriptor table (shared by all CPUs).
3361 struct gatedesc idt[256];
3362 extern uint vectors[]; // in vectors.S: array of 256 entry pointers
3363 struct spinlock tickslock;
3364 uint ticks;
3365
3366 void
3367 tvinit(void)
3368 {
3369     int i;
3370
3371     for(i = 0; i < 256; i++)
3372         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
3373     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
3374
3375     initlock(&tickslock, "time");
3376 }
3377
3378 void
3379 idtinit(void)
3380 {
3381     lidt(idt, sizeof(idt));
3382 }
3383
3384 void
3385 trap(struct trapframe *tf)
3386 {
3387     if(tf->trapno == T_SYSCALL){
3388         if(proc->killed)
3389             exit();
3390         proc->tf = tf;
3391         syscall();
3392         if(proc->killed)
3393             exit();
3394         return;
3395     }
3396
3397     switch(tf->trapno){
3398     case T_IRQ0 + IRQ_TIMER:
3399         if(cpu->id == 0){

```

```

3400     acquire(&tickslock);
3401     ticks++;
3402     release(&tickslock);    // NOTE: MarkM has reversed these two lines.
3403     wakeup(&ticks);        // wakeup() should not require the tickslock t
3404 }
3405 lapiceoi();
3406 break;
3407 case T_IRQ0 + IRQ_IDE:
3408     ideintr();
3409     lapiceoi();
3410     break;
3411 case T_IRQ0 + IRQ_IDE+1:
3412     // Bochs generates spurious IDE1 interrupts.
3413     break;
3414 case T_IRQ0 + IRQ_KBD:
3415     kbdintr();
3416     lapiceoi();
3417     break;
3418 case T_IRQ0 + IRQ_COM1:
3419     uartintr();
3420     lapiceoi();
3421     break;
3422 case T_IRQ0 + 7:
3423 case T_IRQ0 + IRQ_SPURIOUS:
3424     cprintf("cpu%d: spurious interrupt at %x:%x\n",
3425             cpu->id, tf->cs, tf->eip);
3426     lapiceoi();
3427     break;
3428
3429 default:
3430     if(proc == 0 || (tf->cs&3) == 0){
3431         // In kernel, it must be our mistake.
3432         cprintf("unexpected trap %d from cpu %d eip %x (cr2=0x%x)\n",
3433                 tf->trapno, cpu->id, tf->eip, rcr2());
3434         panic("trap");
3435     }
3436     // In user space, assume process misbehaved.
3437     cprintf("pid %d %s: trap %d err %d on cpu %d "
3438             "eip 0x%x addr 0x%x--kill proc\n",
3439             proc->pid, proc->name, tf->trapno, tf->err, cpu->id, tf->eip,
3440             rcr2());
3441     proc->killed = 1;
3442 }
3443
3444 // Force process exit if it has been killed and is in user space.
3445 // (If it is still executing in the kernel, let it keep running
3446 // until it gets to the regular system call return.)
3447 if(proc && proc->killed && (tf->cs&3) == DPL_USER)
3448     exit();
3449

```

```

3450 // Force process to give up CPU on clock tick.
3451 // If interrupts were on while locks held, would need to check nlock.
3452 if(proc && proc->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER)
3453     yield();
3454
3455 // Check if the process has been killed since we yielded
3456 if(proc && proc->killed && (tf->cs&3) == DPL_USER)
3457     exit();
3458 }
3459
3460
3461
3462
3463
3464
3465
3466
3467
3468
3469
3470
3471
3472
3473
3474
3475
3476
3477
3478
3479
3480
3481
3482
3483
3484
3485
3486
3487
3488
3489
3490
3491
3492
3493
3494
3495
3496
3497
3498
3499

```

```

3500 // System call numbers
3501 #define SYS_fork          1
3502 #define SYS_exit          SYS_fork+1
3503 #define SYS_wait          SYS_exit+1
3504 #define SYS_pipe          SYS_wait+1
3505 #define SYS_read          SYS_pipe+1
3506 #define SYS_kill          SYS_read+1
3507 #define SYS_exec          SYS_kill+1
3508 #define SYS_fstat         SYS_exec+1
3509 #define SYS_chdir         SYS_fstat+1
3510 #define SYS_dup           SYS_chdir+1
3511 #define SYS_getpid        SYS_dup+1
3512 #define SYS_sbrk          SYS_getpid+1
3513 #define SYS_sleep         SYS_sbrk+1
3514 #define SYS_uptime        SYS_sleep+1
3515 #define SYS_open          SYS_uptime+1
3516 #define SYS_write         SYS_open+1
3517 #define SYS_mknod         SYS_write+1
3518 #define SYS_unlink        SYS_mknod+1
3519 #define SYS_link          SYS_unlink+1
3520 #define SYS_mkdir         SYS_link+1
3521 #define SYS_close         SYS_mkdir+1
3522 #define SYS_halt          SYS_close+1
3523 // student system calls begin here. Follow the existing pattern.
3524 #define SYS_date           SYS_halt+1
3525 #define SYS_getuid         SYS_date+1
3526 #define SYS_getgid         SYS_getuid+1
3527 #define SYS_getppid        SYS_getgid+1
3528 #define SYS_setuid         SYS_getppid+1
3529 #define SYS_setgid         SYS_setuid+1
3530 #define SYS_getprocs       SYS_setgid+1
3531 #define SYS_setpriority    SYS_getprocs+1
3532
3533
3534
3535
3536
3537
3538
3539
3540
3541
3542
3543
3544
3545
3546
3547
3548
3549

```

```

3550 #include "types.h"
3551 #include "defs.h"
3552 #include "param.h"
3553 #include "memlayout.h"
3554 #include "mmu.h"
3555 #include "proc.h"
3556 #include "x86.h"
3557 #include "syscall.h"
3558
3559 // User code makes a system call with INT T_SYSCALL.
3560 // System call number in %eax.
3561 // Arguments on the stack, from the user call to the C
3562 // library system call function. The saved user %esp points
3563 // to a saved program counter, and then the first argument.
3564
3565 // Fetch the int at addr from the current process.
3566 int
3567 fetchint(uint addr, int *ip)
3568 {
3569     if(addr >= proc->sz || addr+4 > proc->sz)
3570         return -1;
3571     *ip = *(int*)(addr);
3572     return 0;
3573 }
3574
3575 // Fetch the nul-terminated string at addr from the current process.
3576 // Doesn't actually copy the string - just sets *pp to point at it.
3577 // Returns length of string, not including nul.
3578 int
3579 fetchstr(uint addr, char **pp)
3580 {
3581     char *s, *ep;
3582
3583     if(addr >= proc->sz)
3584         return -1;
3585     *pp = (char*)addr;
3586     ep = (char*)proc->sz;
3587     for(s = *pp; s < ep; s++)
3588         if(*s == 0)
3589             return s - *pp;
3590     return -1;
3591 }
3592
3593 // Fetch the nth 32-bit system call argument.
3594 int
3595 argint(int n, int *ip)
3596 {
3597     return fetchint(proc->tf->esp + 4 + 4*n, ip);
3598 }
3599

```

```

3600 // Fetch the nth word-sized system call argument as a pointer
3601 // to a block of memory of size n bytes. Check that the pointer
3602 // lies within the process address space.
3603 int
3604 argptr(int n, char **pp, int size)
3605 {
3606     int i;
3607
3608     if(argint(n, &i) < 0)
3609         return -1;
3610     if((uint)i >= proc->sz || (uint)i+size > proc->sz)
3611         return -1;
3612     *pp = (char*)i;
3613     return 0;
3614 }
3615
3616 // Fetch the nth word-sized system call argument as a string pointer.
3617 // Check that the pointer is valid and the string is nul-terminated.
3618 // (There is no shared writable memory, so the string can't change
3619 // between this check and being used by the kernel.)
3620 int
3621 argstr(int n, char **pp)
3622 {
3623     int addr;
3624     if(argint(n, &addr) < 0)
3625         return -1;
3626     return fetchstr(addr, pp);
3627 }
3628
3629 extern int sys_chdir(void);
3630 extern int sys_close(void);
3631 extern int sys_dup(void);
3632 extern int sys_exec(void);
3633 extern int sys_exit(void);
3634 extern int sys_fork(void);
3635 extern int sys_fstat(void);
3636 extern int sys_getpid(void);
3637 extern int sys_kill(void);
3638 extern int sys_link(void);
3639 extern int sys_mkdir(void);
3640 extern int sys_mknod(void);
3641 extern int sys_open(void);
3642 extern int sys_pipe(void);
3643 extern int sys_read(void);
3644 extern int sys_sbrk(void);
3645 extern int sys_sleep(void);
3646 extern int sys_unlink(void);
3647 extern int sys_wait(void);
3648 extern int sys_write(void);
3649 extern int sys_uptime(void);

```

```

3650 extern int sys_halt(void);
3651 extern int sys_date(void);
3652 #ifdef CS333_P2
3653 extern int sys_getuid(void);
3654 extern int sys_getgid(void);
3655 extern int sys_getppid(void);
3656 extern int sys_setuid(void);
3657 extern int sys_setgid(void);
3658 extern int sys_getprocs(void);
3659 #endif
3660 #ifdef CS333_P3
3661 extern int sys_setpriority(void);
3662 #endif
3663
3664 static int (*syscalls[])(void) = {
3665     [SYS_fork]    sys_fork,
3666     [SYS_exit]    sys_exit,
3667     [SYS_wait]    sys_wait,
3668     [SYS_pipe]    sys_pipe,
3669     [SYS_read]    sys_read,
3670     [SYS_kill]    sys_kill,
3671     [SYS_exec]    sys_exec,
3672     [SYS_fstat]   sys_fstat,
3673     [SYS_chdir]   sys_chdir,
3674     [SYS_dup]     sys_dup,
3675     [SYS_getpid]  sys_getpid,
3676     [SYS_sbrk]    sys_sbrk,
3677     [SYS_sleep]   sys_sleep,
3678     [SYS_uptime]  sys_uptime,
3679     [SYS_open]    sys_open,
3680     [SYS_write]   sys_write,
3681     [SYS_mknod]   sys_mknod,
3682     [SYS_unlink]  sys_unlink,
3683     [SYS_link]    sys_link,
3684     [SYS_mkdir]   sys_mkdir,
3685     [SYS_close]   sys_close,
3686     [SYS_halt]    sys_halt,
3687     [SYS_date]    sys_date,
3688     #ifdef CS333_P2
3689     [SYS_getuid]  sys_getuid,
3690     [SYS_getgid]  sys_getgid,
3691     [SYS_getppid] sys_getppid,
3692     [SYS_setuid]  sys_setuid,
3693     [SYS_setgid]  sys_setgid,
3694     [SYS_getprocs] sys_getprocs,
3695     #endif
3696     #ifdef CS333_P3
3697     [SYS_setpriority] sys_setpriority,
3698     #endif
3699 };

```

```

3700 // put data structure for printing out system call invocation information he:
3701 #ifdef PRINT_SYSCALLS
3702 static const char * (print_syscalls[]) = {
3703 [SYS_fork] = "fork",
3704 [SYS_exit] = "exit",
3705 [SYS_wait] = "wait",
3706 [SYS_pipe] = "pipe",
3707 [SYS_read] = "read",
3708 [SYS_kill] = "kill",
3709 [SYS_exec] = "exec",
3710 [SYS_fstat] = "fstat",
3711 [SYS_chdir] = "chdir",
3712 [SYS_dup] = "dup",
3713 [SYS_getpid] = "getpid",
3714 [SYS_sbrk] = "sbrk",
3715 [SYS_sleep] = "sleep",
3716 [SYS_uptime] = "uptime",
3717 [SYS_open] = "open",
3718 [SYS_write] = "write",
3719 [SYS_mknod] = "mknod",
3720 [SYS_unlink] = "unlink",
3721 [SYS_link] = "link",
3722 [SYS_mkdir] = "mkdir",
3723 [SYS_close] = "close",
3724 [SYS_halt] = "halt",
3725 [SYS_date] = "date",
3726 #ifdef CS333_P2
3727 [SYS_getgid] = "getuid",
3728 [SYS_getuid] = "getgid",
3729 [SYS_getppid] = "getppid",
3730 [SYS_setgid] = "setuid",
3731 [SYS_setuid] = "setgid",
3732 [SYS_getprocs] = "getprocs",
3733 #endif
3734 #ifdef CS333_P3
3735 [SYS_setpriority] = "setpriority",
3736 #endif
3737 };
3738
3739
3740
3741
3742
3743
3744
3745
3746
3747
3748
3749

```

```

3750 #endif
3751
3752 void
3753 syscall(void)
3754 {
3755     int num;
3756
3757     num = proc->tf->eax;
3758     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
3759         proc->tf->eax = syscalls[num]();
3760         // some code goes here
3761         #ifdef PRINT_SYSCALLS
3762             cprintf("%s -> %d\n", print_syscalls[num], proc->tf->eax);
3763         #endif
3764     } else {
3765         cprintf("%d %s: unknown sys call %d\n",
3766             proc->pid, proc->name, num);
3767         proc->tf->eax = -1;
3768     }
3769 }
3770
3771
3772
3773
3774
3775
3776
3777
3778
3779
3780
3781
3782
3783
3784
3785
3786
3787
3788
3789
3790
3791
3792
3793
3794
3795
3796
3797
3798
3799

```

```

3800 #include "types.h"
3801 #include "x86.h"
3802 #include "defs.h"
3803 #include "date.h"
3804 #include "param.h"
3805 #include "memlayout.h"
3806 #include "mmu.h"
3807 #include "proc.h"
3808 #include "uproc.h"
3809
3810 int
3811 sys_fork(void)
3812 {
3813     return fork();
3814 }
3815
3816 int
3817 sys_exit(void)
3818 {
3819     exit();
3820     return 0; // not reached
3821 }
3822
3823 int
3824 sys_wait(void)
3825 {
3826     return wait();
3827 }
3828
3829 int
3830 sys_kill(void)
3831 {
3832     int pid;
3833
3834     if(argint(0, &pid) < 0)
3835         return -1;
3836     return kill(pid);
3837 }
3838
3839 int
3840 sys_getpid(void)
3841 {
3842     return proc->pid;
3843 }
3844
3845
3846
3847
3848
3849

```

```

3850 int
3851 sys_sbrk(void)
3852 {
3853     int addr;
3854     int n;
3855
3856     if(argint(0, &n) < 0)
3857         return -1;
3858     addr = proc->sz;
3859     if(growproc(n) < 0)
3860         return -1;
3861     return addr;
3862 }
3863
3864 int
3865 sys_sleep(void)
3866 {
3867     int n;
3868     uint ticks0;
3869
3870     if(argint(0, &n) < 0)
3871         return -1;
3872     acquire(&tickslock);
3873     ticks0 = ticks;
3874     while(ticks - ticks0 < n){
3875         if(proc->killed){
3876             release(&tickslock);
3877             return -1;
3878         }
3879         sleep(&ticks, &tickslock);
3880     }
3881     release(&tickslock);
3882     return 0;
3883 }
3884
3885 // return how many clock tick interrupts have occurred
3886 // since start.
3887 int
3888 sys_uptime(void)
3889 {
3890     uint xticks;
3891
3892     acquire(&tickslock);
3893     xticks = ticks;
3894     release(&tickslock);
3895     return xticks;
3896 }
3897
3898
3899

```

```

3900 //Turn of the computer
3901 int sys_halt(void){
3902     cprintf("Shutting down ...\n");
3903     //outw(0xB004, 0x0 | 0x2000);
3904     outw(0x604, 0x0 | 0x2000);
3905     return 0;
3906 }
3907 }
3908
3909 //Get current UTC date of the system
3910 int
3911 sys_date(void)
3912 {
3913     struct rtcdate *d;
3914     if(argptr(0, (void*)&d, sizeof(*d)) < 0)
3915         return -1;
3916     cmostime(d);
3917     return 0;
3918 }
3919
3920 #ifdef CS333_P2
3921 // Set UID
3922 int
3923 sys_setuid(void)
3924 {
3925     uint new_uid;
3926     if(argint(0, (int*) &new_uid) < 0)
3927         return -1;
3928     if(new_uid < 0 || new_uid > 32767)
3929         return -1;
3930     proc->uid = new_uid;
3931     return 0;
3932 }
3933
3934 // Set GID
3935 int
3936 sys_setgid(void)
3937 {
3938     uint new_gid;
3939     if(argint(0, (int*) &new_gid) < 0)
3940         return -1;
3941     if(new_gid < 0 || new_gid > 32767)
3942         return -1;
3943     proc->gid = new_gid;
3944     return 0;
3945 }
3946
3947
3948
3949

```

```

3950 // Get UID of current process
3951 int
3952 sys_getuid(void)
3953 {
3954     return proc->uid;
3955 }
3956
3957 // Get GID of current process
3958 int
3959 sys_getgid(void)
3960 {
3961     return proc->gid;
3962 }
3963
3964 // Get PPID of current process
3965 int
3966 sys_getppid(void)
3967 {
3968     if(proc->pid == 1)
3969         return proc->pid;
3970     if(!proc->parent)
3971         return proc->pid;
3972     return proc->parent->pid;
3973 }
3974
3975 // Get process info
3976 int
3977 sys_getprocs(void)
3978 {
3979     uint arg1;
3980     struct uproc* table;
3981     if(argint(0, (int*) &arg1) < 0)
3982         return -1;
3983     if(argptr(1, (void*)&table, sizeof(*table)) < 0)
3984         return -1;;
3985     return getprocs(arg1, table);
3986 }
3987
3988
3989
3990
3991
3992
3993
3994
3995
3996
3997
3998
3999

```



```
4000 #endif
4001
4002 #ifdef CS333_P3
4003 int
4004 sys_setpriority(void)
4005 {
4006     int value;
4007     int pid;
4008     if(argint(0, (int*) &pid) < 0)
4009         return -1;
4010     if(argint(1, (int*) &value) < 0)
4011         return -1;
4012     return setpriority(pid, value);
4013 }
4014 #endif
4015
4016
4017
4018
4019
4020
4021
4022
4023
4024
4025
4026
4027
4028
4029
4030
4031
4032
4033
4034
4035
4036
4037
4038
4039
4040
4041
4042
4043
4044
4045
4046
4047
4048
4049
```

```
4050 // halt the system.
4051 #include "types.h"
4052 #include "user.h"
4053
4054 int
4055 main(void) {
4056     halt();
4057     return 0;
4058 }
4059
4060
4061
4062
4063
4064
4065
4066
4067
4068
4069
4070
4071
4072
4073
4074
4075
4076
4077
4078
4079
4080
4081
4082
4083
4084
4085
4086
4087
4088
4089
4090
4091
4092
4093
4094
4095
4096
4097
4098
4099
```

```
4100 struct buf {
4101     int flags;
4102     uint dev;
4103     uint blockno;
4104     struct buf *prev; // LRU cache list
4105     struct buf *next;
4106     struct buf *qnext; // disk queue
4107     uchar data[BSIZE];
4108 };
4109 #define B_BUSY 0x1 // buffer is locked by some process
4110 #define B_VALID 0x2 // buffer has been read from disk
4111 #define B_DIRTY 0x4 // buffer needs to be written to disk
4112
4113
4114
4115
4116
4117
4118
4119
4120
4121
4122
4123
4124
4125
4126
4127
4128
4129
4130
4131
4132
4133
4134
4135
4136
4137
4138
4139
4140
4141
4142
4143
4144
4145
4146
4147
4148
4149
```

```
4150 #define O_RDONLY 0x000
4151 #define O_WRONLY 0x001
4152 #define O_RDWR 0x002
4153 #define O_CREATE 0x200
4154
4155
4156
4157
4158
4159
4160
4161
4162
4163
4164
4165
4166
4167
4168
4169
4170
4171
4172
4173
4174
4175
4176
4177
4178
4179
4180
4181
4182
4183
4184
4185
4186
4187
4188
4189
4190
4191
4192
4193
4194
4195
4196
4197
4198
4199
```

```

4200 #define T_DIR 1 // Directory
4201 #define T_FILE 2 // File
4202 #define T_DEV 3 // Device
4203
4204 struct stat {
4205     short type; // Type of file
4206     int dev; // File system's disk device
4207     uint ino; // Inode number
4208     short nlink; // Number of links to file
4209     uint size; // Size of file in bytes
4210 };
4211
4212
4213
4214
4215
4216
4217
4218
4219
4220
4221
4222
4223
4224
4225
4226
4227
4228
4229
4230
4231
4232
4233
4234
4235
4236
4237
4238
4239
4240
4241
4242
4243
4244
4245
4246
4247
4248
4249

```

```

4250 // On-disk file system format.
4251 // Both the kernel and user programs use this header file.
4252
4253
4254 #define ROOTINO 1 // root i-number
4255 #define BSIZE 512 // block size
4256
4257 // Disk layout:
4258 // [ boot block | super block | log | inode blocks | free bit map | data blocks ]
4259 //
4260 // mkfs computes the super block and builds an initial file system. The super block
4261 // the disk layout:
4262 struct superblock {
4263     uint size; // Size of file system image (blocks)
4264     uint nblocks; // Number of data blocks
4265     uint ninodes; // Number of inodes.
4266     uint nlog; // Number of log blocks
4267     uint logstart; // Block number of first log block
4268     uint inodestart; // Block number of first inode block
4269     uint bmapstart; // Block number of first free map block
4270 };
4271
4272 #define NDIRECT 12
4273 #define NINDIRECT (BSIZE / sizeof(uint))
4274 #define MAXFILE (NDIRECT + NINDIRECT)
4275
4276 // On-disk inode structure
4277 struct dinode {
4278     short type; // File type
4279     short major; // Major device number (T_DEV only)
4280     short minor; // Minor device number (T_DEV only)
4281     short nlink; // Number of links to inode in file system
4282     uint size; // Size of file (bytes)
4283     uint addrs[NDIRECT+1]; // Data block addresses
4284 };
4285
4286
4287
4288
4289
4290
4291
4292
4293
4294
4295
4296
4297
4298
4299

```

```

4300 // Inodes per block.
4301 #define IPB          (BSIZE / sizeof(struct dinode))
4302
4303 // Block containing inode i
4304 #define IBLOCK(i, sb) ((i) / IPB + sb.inodestart)
4305
4306 // Bitmap bits per block
4307 #define BPB          (BSIZE*8)
4308
4309 // Block of free map containing bit for block b
4310 #define BBLOCK(b, sb) (b/BPB + sb.bmapstart)
4311
4312 // Directory is a file containing a sequence of dirent structures.
4313 #define DIRSIZ 14
4314
4315 struct dirent {
4316     ushort inum;
4317     char name[DIRSIZ];
4318 };
4319
4320
4321
4322
4323
4324
4325
4326
4327
4328
4329
4330
4331
4332
4333
4334
4335
4336
4337
4338
4339
4340
4341
4342
4343
4344
4345
4346
4347
4348
4349

```

```

4350 struct file {
4351     enum { FD_NONE, FD_PIPE, FD_INODE } type;
4352     int ref; // reference count
4353     char readable;
4354     char writable;
4355     struct pipe *pipe;
4356     struct inode *ip;
4357     uint off;
4358 };
4359
4360
4361 // in-memory copy of an inode
4362 struct inode {
4363     uint dev;           // Device number
4364     uint inum;          // Inode number
4365     int ref;            // Reference count
4366     int flags;          // I_BUSY, I_VALID
4367
4368     short type;         // copy of disk inode
4369     short major;
4370     short minor;
4371     short nlink;
4372     uint size;
4373     uint addrs[NDIRECT+1];
4374 };
4375 #define I_BUSY 0x1
4376 #define I_VALID 0x2
4377
4378 // table mapping major device number to
4379 // device functions
4380 struct devsw {
4381     int (*read)(struct inode*, char*, int);
4382     int (*write)(struct inode*, char*, int);
4383 };
4384
4385 extern struct devsw devsw[];
4386
4387 #define CONSOLE 1
4388
4389 // Blank page.
4390
4391
4392
4393
4394
4395
4396
4397
4398
4399

```

```

4400 // Simple PIO-based (non-DMA) IDE driver code.
4401
4402 #include "types.h"
4403 #include "defs.h"
4404 #include "param.h"
4405 #include "memlayout.h"
4406 #include "mmu.h"
4407 #include "proc.h"
4408 #include "x86.h"
4409 #include "traps.h"
4410 #include "spinlock.h"
4411 #include "fs.h"
4412 #include "buf.h"
4413
4414 #define SECTOR_SIZE 512
4415 #define IDE_BSY 0x80
4416 #define IDE_DRDY 0x40
4417 #define IDE_DF 0x20
4418 #define IDE_ERR 0x01
4419
4420 #define IDE_CMD_READ 0x20
4421 #define IDE_CMD_WRITE 0x30
4422
4423 // idequeue points to the buf now being read/written to the disk.
4424 // idequeue->qnext points to the next buf to be processed.
4425 // You must hold idelock while manipulating queue.
4426
4427 static struct spinlock idelock;
4428 static struct buf *idequeue;
4429
4430 static int havedisk1;
4431 static void idestart(struct buf*);
4432
4433 // Wait for IDE disk to become ready.
4434 static int
4435 idewait(int checkerr)
4436 {
4437     int r;
4438
4439     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)
4440         ;
4441     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)
4442         return -1;
4443     return 0;
4444 }
4445
4446
4447
4448
4449

```

```

4450 void
4451 ideinit(void)
4452 {
4453     int i;
4454
4455     initlock(&idelock, "ide");
4456     picenable(IRQ_IDE);
4457     ioapicenable(IRQ_IDE, ncpu - 1);
4458     idewait(0);
4459
4460     // Check if disk 1 is present
4461     outb(0x1f6, 0xe0 | (1<<4));
4462     for(i=0; i<1000; i++){
4463         if(inb(0x1f7) != 0){
4464             havedisk1 = 1;
4465             break;
4466         }
4467     }
4468
4469     // Switch back to disk 0.
4470     outb(0x1f6, 0xe0 | (0<<4));
4471 }
4472
4473 // Start the request for b. Caller must hold idelock.
4474 static void
4475 idestart(struct buf *b)
4476 {
4477     if(b == 0)
4478         panic("idestart");
4479     if(b->blockno >= FSSIZE)
4480         panic("incorrect blockno");
4481     int sector_per_block = BSIZE/SECTOR_SIZE;
4482     int sector = b->blockno * sector_per_block;
4483
4484     if (sector_per_block > 7) panic("idestart");
4485
4486     idewait(0);
4487     outb(0x3f6, 0); // generate interrupt
4488     outb(0x1f2, sector_per_block); // number of sectors
4489     outb(0x1f3, sector & 0xff);
4490     outb(0x1f4, (sector >> 8) & 0xff);
4491     outb(0x1f5, (sector >> 16) & 0xff);
4492     outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((sector>>24)&0x0f));
4493     if(b->flags & B_DIRTY){
4494         outb(0x1f7, IDE_CMD_WRITE);
4495         outsl(0x1f0, b->data, BSIZE/4);
4496     } else {
4497         outb(0x1f7, IDE_CMD_READ);
4498     }
4499 }

```

```

4500 // Interrupt handler.
4501 void
4502 ideintr(void)
4503 {
4504     struct buf *b;
4505
4506     // First queued buffer is the active request.
4507     acquire(&idelock);
4508     if((b = idequeue) == 0){
4509         release(&idelock);
4510         // cprintf("spurious IDE interrupt\n");
4511         return;
4512     }
4513     idequeue = b->qnext;
4514
4515     // Read data if needed.
4516     if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
4517         insl(0x1f0, b->data, BSIZE/4);
4518
4519     // Wake process waiting for this buf.
4520     b->flags |= B_VALID;
4521     b->flags &= ~B_DIRTY;
4522     wakeup(b);
4523
4524     // Start disk on next buf in queue.
4525     if(idequeue != 0)
4526         idestart(idequeue);
4527
4528     release(&idelock);
4529 }
4530
4531 // Sync buf with disk.
4532 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
4533 // Else if B_VALID is not set, read buf from disk, set B_VALID.
4534 void
4535 iderw(struct buf *b)
4536 {
4537     struct buf **pp;
4538
4539     if(!(b->flags & B_BUSY))
4540         panic("iderw: buf not busy");
4541     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
4542         panic("iderw: nothing to do");
4543     if(b->dev != 0 && !havedisk1)
4544         panic("iderw: ide disk 1 not present");
4545
4546     acquire(&idelock);
4547
4548
4549

```

```

4550 // Append b to idequeue.
4551 b->qnext = 0;
4552 for(pp=&idequeue; *pp; pp=&(*pp)->qnext)
4553     ;
4554 *pp = b;
4555
4556 // Start disk if necessary.
4557 if(idequeue == b)
4558     idestart(b);
4559
4560 // Wait for request to finish.
4561 while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
4562     sleep(b, &idelock);
4563 }
4564
4565 release(&idelock);
4566 }
4567
4568
4569
4570
4571
4572
4573
4574
4575
4576
4577
4578
4579
4580
4581
4582
4583
4584
4585
4586
4587
4588
4589
4590
4591
4592
4593
4594
4595
4596
4597
4598
4599

```

```

4600 // Buffer cache.
4601 //
4602 // The buffer cache is a linked list of buf structures holding
4603 // cached copies of disk block contents. Caching disk blocks
4604 // in memory reduces the number of disk reads and also provides
4605 // a synchronization point for disk blocks used by multiple processes.
4606 //
4607 // Interface:
4608 // * To get a buffer for a particular disk block, call bread.
4609 // * After changing buffer data, call bwrite to write it to disk.
4610 // * When done with the buffer, call brelse.
4611 // * Do not use the buffer after calling brelse.
4612 // * Only one process at a time can use a buffer,
4613 //   so do not keep them longer than necessary.
4614 //
4615 // The implementation uses three state flags internally:
4616 // * B_BUSY: the block has been returned from bread
4617 //   and has not been passed back to brelse.
4618 // * B_VALID: the buffer data has been read from the disk.
4619 // * B_DIRTY: the buffer data has been modified
4620 //   and needs to be written to disk.
4621
4622 #include "types.h"
4623 #include "defs.h"
4624 #include "param.h"
4625 #include "spinlock.h"
4626 #include "fs.h"
4627 #include "buf.h"
4628
4629 struct {
4630   struct spinlock lock;
4631   struct buf buf[NBUF];
4632
4633   // Linked list of all buffers, through prev/next.
4634   // head.next is most recently used.
4635   struct buf head;
4636 } bcache;
4637
4638 void
4639 binit(void)
4640 {
4641   struct buf *b;
4642
4643   initlock(&bcache.lock, "bcache");
4644
4645   // Create linked list of buffers
4646   bcache.head.prev = &bcache.head;
4647   bcache.head.next = &bcache.head;
4648   for(b = bcache.buf; b < bcache.buf+NBUF; b++){
4649     b->next = bcache.head.next;

```

```

4650     b->prev = &bcache.head;
4651     b->dev = -1;
4652     bcache.head.next->prev = b;
4653     bcache.head.next = b;
4654   }
4655 }
4656
4657 // Look through buffer cache for block on device dev.
4658 // If not found, allocate a buffer.
4659 // In either case, return B_BUSY buffer.
4660 static struct buf*
4661 bget(uint dev, uint blockno)
4662 {
4663   struct buf *b;
4664
4665   acquire(&bcache.lock);
4666
4667   loop:
4668   // Is the block already cached?
4669   for(b = bcache.head.next; b != &bcache.head; b = b->next){
4670     if(b->dev == dev && b->blockno == blockno){
4671       if(!(b->flags & B_BUSY)){
4672         b->flags |= B_BUSY;
4673         release(&bcache.lock);
4674         return b;
4675       }
4676       sleep(b, &bcache.lock);
4677       goto loop;
4678     }
4679   }
4680
4681   // Not cached; recycle some non-busy and clean buffer.
4682   // "clean" because B_DIRTY and !B_BUSY means log.c
4683   // hasn't yet committed the changes to the buffer.
4684   for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
4685     if((b->flags & B_BUSY) == 0 && (b->flags & B_DIRTY) == 0){
4686       b->dev = dev;
4687       b->blockno = blockno;
4688       b->flags = B_BUSY;
4689       release(&bcache.lock);
4690       return b;
4691     }
4692   }
4693   panic("bget: no buffers");
4694 }
4695
4696
4697
4698
4699

```

```

4700 // Return a B_BUSY buf with the contents of the indicated block.
4701 struct buf*
4702 bread(uint dev, uint blockno)
4703 {
4704     struct buf *b;
4705
4706     b = bget(dev, blockno);
4707     if(!(b->flags & B_VALID)) {
4708         iderw(b);
4709     }
4710     return b;
4711 }
4712
4713 // Write b's contents to disk. Must be B_BUSY.
4714 void
4715 bwrite(struct buf *b)
4716 {
4717     if((b->flags & B_BUSY) == 0)
4718         panic("bwrite");
4719     b->flags |= B_DIRTY;
4720     iderw(b);
4721 }
4722
4723 // Release a B_BUSY buffer.
4724 // Move to the head of the MRU list.
4725 void
4726 brelse(struct buf *b)
4727 {
4728     if((b->flags & B_BUSY) == 0)
4729         panic("brelse");
4730
4731     acquire(&bcache.lock);
4732
4733     b->next->prev = b->prev;
4734     b->prev->next = b->next;
4735     b->next = bcache.head.next;
4736     b->prev = &bcache.head;
4737     bcache.head.next->prev = b;
4738     bcache.head.next = b;
4739
4740     b->flags &= ~B_BUSY;
4741     wakeup(b);
4742
4743     release(&bcache.lock);
4744 }
4745 // Blank page.
4746
4747
4748
4749

```

```

4750 #include "types.h"
4751 #include "defs.h"
4752 #include "param.h"
4753 #include "spinlock.h"
4754 #include "fs.h"
4755 #include "buf.h"
4756
4757 // Simple logging that allows concurrent FS system calls.
4758 //
4759 // A log transaction contains the updates of multiple FS system
4760 // calls. The logging system only commits when there are
4761 // no FS system calls active. Thus there is never
4762 // any reasoning required about whether a commit might
4763 // write an uncommitted system call's updates to disk.
4764 //
4765 // A system call should call begin_op()/end_op() to mark
4766 // its start and end. Usually begin_op() just increments
4767 // the count of in-progress FS system calls and returns.
4768 // But if it thinks the log is close to running out, it
4769 // sleeps until the last outstanding end_op() commits.
4770 //
4771 // The log is a physical re-do log containing disk blocks.
4772 // The on-disk log format:
4773 //   header block, containing block #s for block A, B, C, ...
4774 //   block A
4775 //   block B
4776 //   block C
4777 //   ...
4778 // Log appends are synchronous.
4779
4780 // Contents of the header block, used for both the on-disk header block
4781 // and to keep track in memory of logged block# before commit.
4782 struct logheader {
4783     int n;
4784     int block[LOGSIZE];
4785 };
4786
4787 struct log {
4788     struct spinlock lock;
4789     int start;
4790     int size;
4791     int outstanding; // how many FS sys calls are executing.
4792     int committing;  // in commit(), please wait.
4793     int dev;
4794     struct logheader lh;
4795 };
4796
4797
4798
4799

```



```

4800 struct log log;
4801
4802 static void recover_from_log(void);
4803 static void commit();
4804
4805 void
4806 initlog(int dev)
4807 {
4808     if (sizeof(struct logheader) >= BSIZE)
4809         panic("initlog: too big logheader");
4810
4811     struct superblock sb;
4812     initlock(&log.lock, "log");
4813     readsb(dev, &sb);
4814     log.start = sb.logstart;
4815     log.size = sb.nlog;
4816     log.dev = dev;
4817     recover_from_log();
4818 }
4819
4820 // Copy committed blocks from log to their home location
4821 static void
4822 install_trans(void)
4823 {
4824     int tail;
4825
4826     for (tail = 0; tail < log.lh.n; tail++) {
4827         struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block
4828         struct buf *dbuf = bread(log.dev, log.lh.block[tail]); // read dst
4829         memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst
4830         bwrite(dbuf); // write dst to disk
4831         brelse(lbuf);
4832         brelse(dbuf);
4833     }
4834 }
4835
4836 // Read the log header from disk into the in-memory log header
4837 static void
4838 read_head(void)
4839 {
4840     struct buf *buf = bread(log.dev, log.start);
4841     struct logheader *lh = (struct logheader *) (buf->data);
4842     int i;
4843     log.lh.n = lh->n;
4844     for (i = 0; i < log.lh.n; i++) {
4845         log.lh.block[i] = lh->block[i];
4846     }
4847     brelse(buf);
4848 }
4849

```

```

4850 // Write in-memory log header to disk.
4851 // This is the true point at which the
4852 // current transaction commits.
4853 static void
4854 write_head(void)
4855 {
4856     struct buf *buf = bread(log.dev, log.start);
4857     struct logheader *hb = (struct logheader *) (buf->data);
4858     int i;
4859     hb->n = log.lh.n;
4860     for (i = 0; i < log.lh.n; i++) {
4861         hb->block[i] = log.lh.block[i];
4862     }
4863     bwrite(buf);
4864     brelse(buf);
4865 }
4866
4867 static void
4868 recover_from_log(void)
4869 {
4870     read_head();
4871     install_trans(); // if committed, copy from log to disk
4872     log.lh.n = 0;
4873     write_head(); // clear the log
4874 }
4875
4876 // called at the start of each FS system call.
4877 void
4878 begin_op(void)
4879 {
4880     acquire(&log.lock);
4881     while(1){
4882         if(log.committing){
4883             sleep(&log, &log.lock);
4884         } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){
4885             // this op might exhaust log space; wait for commit.
4886             sleep(&log, &log.lock);
4887         } else {
4888             log.outstanding += 1;
4889             release(&log.lock);
4890             break;
4891         }
4892     }
4893 }
4894
4895
4896
4897
4898
4899

```

```

4900 // called at the end of each FS system call.
4901 // commits if this was the last outstanding operation.
4902 void
4903 end_op(void)
4904 {
4905     int do_commit = 0;
4906
4907     acquire(&log.lock);
4908     log.outstanding -= 1;
4909     if(log.committing)
4910         panic("log.committing");
4911     if(log.outstanding == 0){
4912         do_commit = 1;
4913         log.committing = 1;
4914     } else {
4915         // begin_op() may be waiting for log space.
4916         wakeup(&log);
4917     }
4918     release(&log.lock);
4919
4920     if(do_commit){
4921         // call commit w/o holding locks, since not allowed
4922         // to sleep with locks.
4923         commit();
4924         acquire(&log.lock);
4925         log.committing = 0;
4926         wakeup(&log);
4927         release(&log.lock);
4928     }
4929 }
4930
4931 // Copy modified blocks from cache to log.
4932 static void
4933 write_log(void)
4934 {
4935     int tail;
4936
4937     for (tail = 0; tail < log.lh.n; tail++) {
4938         struct buf *to = bread(log.dev, log.start+tail+1); // log block
4939         struct buf *from = bread(log.dev, log.lh.block[tail]); // cache block
4940         memmove(to->data, from->data, BSIZE);
4941         bwrite(to); // write the log
4942         brelse(from);
4943         brelse(to);
4944     }
4945 }
4946
4947
4948
4949

```

```

4950 static void
4951 commit()
4952 {
4953     if (log.lh.n > 0) {
4954         write_log(); // Write modified blocks from cache to log
4955         write_head(); // Write header to disk -- the real commit
4956         install_trans(); // Now install writes to home locations
4957         log.lh.n = 0;
4958         write_head(); // Erase the transaction from the log
4959     }
4960 }
4961
4962 // Caller has modified b->data and is done with the buffer.
4963 // Record the block number and pin in the cache with B_DIRTY.
4964 // commit()/write_log() will do the disk write.
4965 //
4966 // log_write() replaces bwrite(); a typical use is:
4967 //   bp = bread(...)
4968 //   modify bp->data[]
4969 //   log_write(bp)
4970 //   brelse(bp)
4971 void
4972 log_write(struct buf *b)
4973 {
4974     int i;
4975
4976     if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
4977         panic("too big a transaction");
4978     if (log.outstanding < 1)
4979         panic("log_write outside of trans");
4980
4981     acquire(&log.lock);
4982     for (i = 0; i < log.lh.n; i++) {
4983         if (log.lh.block[i] == b->blockno) // log absorption
4984             break;
4985     }
4986     log.lh.block[i] = b->blockno;
4987     if (i == log.lh.n)
4988         log.lh.n++;
4989     b->flags |= B_DIRTY; // prevent eviction
4990     release(&log.lock);
4991 }
4992
4993
4994
4995
4996
4997
4998
4999

```

```

5000 // File system implementation. Five layers:
5001 //   + Blocks: allocator for raw disk blocks.
5002 //   + Log: crash recovery for multi-step updates.
5003 //   + Files: inode allocator, reading, writing, metadata.
5004 //   + Directories: inode with special contents (list of other inodes!)
5005 //   + Names: paths like /usr/rtn/xv6/fs.c for convenient naming.
5006 //
5007 // This file contains the low-level file system manipulation
5008 // routines. The (higher-level) system call implementations
5009 // are in sysfile.c.
5010
5011 #include "types.h"
5012 #include "defs.h"
5013 #include "param.h"
5014 #include "stat.h"
5015 #include "mmu.h"
5016 #include "proc.h"
5017 #include "spinlock.h"
5018 #include "fs.h"
5019 #include "buf.h"
5020 #include "file.h"
5021
5022 #define min(a, b) ((a) < (b) ? (a) : (b))
5023 static void itrunc(struct inode*);
5024 struct superblock sb; // there should be one per dev, but we run with one
5025
5026 // Read the super block.
5027 void
5028 readsb(int dev, struct superblock *sb)
5029 {
5030     struct buf *bp;
5031
5032     bp = bread(dev, 1);
5033     memmove(sb, bp->data, sizeof(*sb));
5034     brelse(bp);
5035 }
5036
5037 // Zero a block.
5038 static void
5039 bzero(int dev, int bno)
5040 {
5041     struct buf *bp;
5042
5043     bp = bread(dev, bno);
5044     memset(bp->data, 0, BSIZE);
5045     log_write(bp);
5046     brelse(bp);
5047 }
5048
5049

```

```

5050 // Blocks.
5051
5052 // Allocate a zeroed disk block.
5053 static uint
5054 balloc(uint dev)
5055 {
5056     int b, bi, m;
5057     struct buf *bp;
5058
5059     bp = 0;
5060     for(b = 0; b < sb.size; b += BPB){
5061         bp = bread(dev, BBLOCK(b, sb));
5062         for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
5063             m = 1 << (bi % 8);
5064             if((bp->data[bi/8] & m) == 0){ // Is block free?
5065                 bp->data[bi/8] |= m; // Mark block in use.
5066                 log_write(bp);
5067                 brelse(bp);
5068                 bzero(dev, b + bi);
5069                 return b + bi;
5070             }
5071         }
5072         brelse(bp);
5073     }
5074     panic("balloc: out of blocks");
5075 }
5076
5077 // Free a disk block.
5078 static void
5079 bfree(int dev, uint b)
5080 {
5081     struct buf *bp;
5082     int bi, m;
5083
5084     readsb(dev, &sb);
5085     bp = bread(dev, BBLOCK(b, sb));
5086     bi = b % BPB;
5087     m = 1 << (bi % 8);
5088     if((bp->data[bi/8] & m) == 0)
5089         panic("freeing free block");
5090     bp->data[bi/8] &= ~m;
5091     log_write(bp);
5092     brelse(bp);
5093 }
5094
5095
5096
5097
5098
5099

```

```

5100 // Inodes.
5101 //
5102 // An inode describes a single unnamed file.
5103 // The inode disk structure holds metadata: the file's type,
5104 // its size, the number of links referring to it, and the
5105 // list of blocks holding the file's content.
5106 //
5107 // The inodes are laid out sequentially on disk at
5108 // sb.startinode. Each inode has a number, indicating its
5109 // position on the disk.
5110 //
5111 // The kernel keeps a cache of in-use inodes in memory
5112 // to provide a place for synchronizing access
5113 // to inodes used by multiple processes. The cached
5114 // inodes include book-keeping information that is
5115 // not stored on disk: ip->ref and ip->flags.
5116 //
5117 // An inode and its in-memory representative go through a
5118 // sequence of states before they can be used by the
5119 // rest of the file system code.
5120 //
5121 // * Allocation: an inode is allocated if its type (on disk)
5122 //   is non-zero. ialloc() allocates, iput() frees if
5123 //   the link count has fallen to zero.
5124 //
5125 // * Referencing in cache: an entry in the inode cache
5126 //   is free if ip->ref is zero. Otherwise ip->ref tracks
5127 //   the number of in-memory pointers to the entry (open
5128 //   files and current directories). iget() to find or
5129 //   create a cache entry and increment its ref, iput()
5130 //   to decrement ref.
5131 //
5132 // * Valid: the information (type, size, &c) in an inode
5133 //   cache entry is only correct when the I_VALID bit
5134 //   is set in ip->flags. ilock() reads the inode from
5135 //   the disk and sets I_VALID, while iput() clears
5136 //   I_VALID if ip->ref has fallen to zero.
5137 //
5138 // * Locked: file system code may only examine and modify
5139 //   the information in an inode and its content if it
5140 //   has first locked the inode. The I_BUSY flag indicates
5141 //   that the inode is locked. ilock() sets I_BUSY,
5142 //   while iunlock clears it.
5143 //
5144 // Thus a typical sequence is:
5145 //   ip = iget(dev, inum)
5146 //   ilock(ip)
5147 //   ... examine and modify ip->xxx ...
5148 //   iunlock(ip)
5149 //   iput(ip)

```

```

5150 //
5151 // ilock() is separate from iget() so that system calls can
5152 // get a long-term reference to an inode (as for an open file)
5153 // and only lock it for short periods (e.g., in read()).
5154 // The separation also helps avoid deadlock and races during
5155 // pathname lookup. iget() increments ip->ref so that the inode
5156 // stays cached and pointers to it remain valid.
5157 //
5158 // Many internal file system functions expect the caller to
5159 // have locked the inodes involved; this lets callers create
5160 // multi-step atomic operations.
5161 //
5162 struct {
5163   struct spinlock lock;
5164   struct inode inode[NINODE];
5165 } icache;
5166
5167 void
5168 iinit(int dev)
5169 {
5170   initlock(&icache.lock, "icache");
5171   readsb(dev, &sb);
5172   cprintf("sb: size %d nblocks %d ninodes %d nlog %d logstart %d inodestart %d\n",
5173           sb.nblocks, sb.ninodes, sb.nlog, sb.logstart, sb.inodestart, sb.bmap);
5174 }
5175
5176 static struct inode* iget(uint dev, uint inum);
5177
5178 // Allocate a new inode with the given type on device dev.
5179 // A free inode has a type of zero.
5180 struct inode*
5181 ialloc(uint dev, short type)
5182 {
5183   int inum;
5184   struct buf *bp;
5185   struct dinode *dip;
5186
5187   for(inum = 1; inum < sb.ninodes; inum++){
5188     bp = bread(dev, IBLOCK(inum, sb));
5189     dip = (struct dinode*)bp->data + inum%IPB;
5190     if(dip->type == 0){ // a free inode
5191       memset(dip, 0, sizeof(*dip));
5192       dip->type = type;
5193       log_write(bp); // mark it allocated on the disk
5194       brelse(bp);
5195       return iget(dev, inum);
5196     }
5197     brelse(bp);
5198   }
5199   panic("ialloc: no inodes");

```

```

5200 }
5201
5202 // Copy a modified in-memory inode to disk.
5203 void
5204 iupdate(struct inode *ip)
5205 {
5206     struct buf *bp;
5207     struct dinode *dip;
5208
5209     bp = bread(ip->dev, IBLOCK(ip->inum, sb));
5210     dip = (struct dinode*)bp->data + ip->inum%IPB;
5211     dip->type = ip->type;
5212     dip->major = ip->major;
5213     dip->minor = ip->minor;
5214     dip->nlink = ip->nlink;
5215     dip->size = ip->size;
5216     memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
5217     log_write(bp);
5218     brelse(bp);
5219 }
5220
5221 // Find the inode with number inum on device dev
5222 // and return the in-memory copy. Does not lock
5223 // the inode and does not read it from disk.
5224 static struct inode*
5225 iget(uint dev, uint inum)
5226 {
5227     struct inode *ip, *empty;
5228
5229     acquire(&icache.lock);
5230
5231     // Is the inode already cached?
5232     empty = 0;
5233     for(ip = &icache.inode[0]; ip < &icache.inode[NINODE]; ip++){
5234         if(ip->ref > 0 && ip->dev == dev && ip->inum == inum){
5235             ip->ref++;
5236             release(&icache.lock);
5237             return ip;
5238         }
5239         if(empty == 0 && ip->ref == 0)    // Remember empty slot.
5240             empty = ip;
5241     }
5242
5243     // Recycle an inode cache entry.
5244     if(empty == 0)
5245         panic("iget: no inodes");
5246
5247
5248
5249

```

```

5250     ip = empty;
5251     ip->dev = dev;
5252     ip->inum = inum;
5253     ip->ref = 1;
5254     ip->flags = 0;
5255     release(&icache.lock);
5256
5257     return ip;
5258 }
5259
5260 // Increment reference count for ip.
5261 // Returns ip to enable ip = idup(ip1) idiom.
5262 struct inode*
5263 idup(struct inode *ip)
5264 {
5265     acquire(&icache.lock);
5266     ip->ref++;
5267     release(&icache.lock);
5268     return ip;
5269 }
5270
5271 // Lock the given inode.
5272 // Reads the inode from disk if necessary.
5273 void
5274 ilock(struct inode *ip)
5275 {
5276     struct buf *bp;
5277     struct dinode *dip;
5278
5279     if(ip == 0 || ip->ref < 1)
5280         panic("ilock");
5281
5282     acquire(&icache.lock);
5283     while(ip->flags & I_BUSY)
5284         sleep(ip, &icache.lock);
5285     ip->flags |= I_BUSY;
5286     release(&icache.lock);
5287
5288     if(!(ip->flags & I_VALID)){
5289         bp = bread(ip->dev, IBLOCK(ip->inum, sb));
5290         dip = (struct dinode*)bp->data + ip->inum%IPB;
5291         ip->type = dip->type;
5292         ip->major = dip->major;
5293         ip->minor = dip->minor;
5294         ip->nlink = dip->nlink;
5295         ip->size = dip->size;
5296         memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
5297         brelse(bp);
5298         ip->flags |= I_VALID;
5299         if(ip->type == 0)

```

```

5300     panic("ilock: no type");
5301 }
5302 }
5303
5304 // Unlock the given inode.
5305 void
5306 iunlock(struct inode *ip)
5307 {
5308     if(ip == 0 || !(ip->flags & I_BUSY) || ip->ref < 1)
5309         panic("iunlock");
5310
5311     acquire(&icache.lock);
5312     ip->flags &= ~I_BUSY;
5313     wakeup(ip);
5314     release(&icache.lock);
5315 }
5316
5317 // Drop a reference to an in-memory inode.
5318 // If that was the last reference, the inode cache entry can
5319 // be recycled.
5320 // If that was the last reference and the inode has no links
5321 // to it, free the inode (and its content) on disk.
5322 // All calls to iput() must be inside a transaction in
5323 // case it has to free the inode.
5324 void
5325 iput(struct inode *ip)
5326 {
5327     acquire(&icache.lock);
5328     if(ip->ref == 1 && (ip->flags & I_VALID) && ip->nlink == 0){
5329         // inode has no links and no other references: truncate and free.
5330         if(ip->flags & I_BUSY)
5331             panic("iput busy");
5332         ip->flags |= I_BUSY;
5333         release(&icache.lock);
5334         itrunc(ip);
5335         ip->type = 0;
5336         iupdate(ip);
5337         acquire(&icache.lock);
5338         ip->flags = 0;
5339         wakeup(ip);
5340     }
5341     ip->ref--;
5342     release(&icache.lock);
5343 }
5344
5345
5346
5347
5348
5349

```

```

5350 // Common idiom: unlock, then put.
5351 void
5352 iunlockput(struct inode *ip)
5353 {
5354     iunlock(ip);
5355     iput(ip);
5356 }
5357
5358 // Inode content
5359 //
5360 // The content (data) associated with each inode is stored
5361 // in blocks on the disk. The first NDIRECT block numbers
5362 // are listed in ip->addrs[]. The next NINDIRECT blocks are
5363 // listed in block ip->addrs[NDIRECT].
5364
5365 // Return the disk block address of the nth block in inode ip.
5366 // If there is no such block, bmap allocates one.
5367 static uint
5368 bmap(struct inode *ip, uint bn)
5369 {
5370     uint addr, *a;
5371     struct buf *bp;
5372
5373     if(bn < NDIRECT){
5374         if((addr = ip->addrs[bn]) == 0)
5375             ip->addrs[bn] = addr = balloc(ip->dev);
5376         return addr;
5377     }
5378     bn -= NDIRECT;
5379
5380     if(bn < NINDIRECT){
5381         // Load indirect block, allocating if necessary.
5382         if((addr = ip->addrs[NDIRECT]) == 0)
5383             ip->addrs[NDIRECT] = addr = balloc(ip->dev);
5384         bp = bread(ip->dev, addr);
5385         a = (uint*)bp->data;
5386         if((addr = a[bn]) == 0){
5387             a[bn] = addr = balloc(ip->dev);
5388             log_write(bp);
5389         }
5390         brelse(bp);
5391         return addr;
5392     }
5393
5394     panic("bmap: out of range");
5395 }
5396
5397
5398
5399

```

```

5400 // Truncate inode (discard contents).
5401 // Only called when the inode has no links
5402 // to it (no directory entries referring to it)
5403 // and has no in-memory reference to it (is
5404 // not an open file or current directory).
5405 static void
5406 itrunc(struct inode *ip)
5407 {
5408     int i, j;
5409     struct buf *bp;
5410     uint *a;
5411
5412     for(i = 0; i < NDIRECT; i++){
5413         if(ip->addrs[i]){
5414             bfree(ip->dev, ip->addrs[i]);
5415             ip->addrs[i] = 0;
5416         }
5417     }
5418
5419     if(ip->addrs[NDIRECT]){
5420         bp = bread(ip->dev, ip->addrs[NDIRECT]);
5421         a = (uint*)bp->data;
5422         for(j = 0; j < NINDIRECT; j++){
5423             if(a[j])
5424                 bfree(ip->dev, a[j]);
5425         }
5426         brelse(bp);
5427         bfree(ip->dev, ip->addrs[NDIRECT]);
5428         ip->addrs[NDIRECT] = 0;
5429     }
5430
5431     ip->size = 0;
5432     iupdate(ip);
5433 }
5434
5435 // Copy stat information from inode.
5436 void
5437 stati(struct inode *ip, struct stat *st)
5438 {
5439     st->dev = ip->dev;
5440     st->ino = ip->inum;
5441     st->type = ip->type;
5442     st->nlink = ip->nlink;
5443     st->size = ip->size;
5444 }
5445
5446
5447
5448
5449

```

```

5450 // Read data from inode.
5451 int
5452 readi(struct inode *ip, char *dst, uint off, uint n)
5453 {
5454     uint tot, m;
5455     struct buf *bp;
5456
5457     if(ip->type == T_DEV){
5458         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
5459             return -1;
5460         return devsw[ip->major].read(ip, dst, n);
5461     }
5462
5463     if(off > ip->size || off + n < off)
5464         return -1;
5465     if(off + n > ip->size)
5466         n = ip->size - off;
5467
5468     for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
5469         bp = bread(ip->dev, bmap(ip, off/BSIZE));
5470         m = min(n - tot, BSIZE - off%BSIZE);
5471         memmove(dst, bp->data + off%BSIZE, m);
5472         brelse(bp);
5473     }
5474     return n;
5475 }
5476
5477 // Write data to inode.
5478 int
5479 writei(struct inode *ip, char *src, uint off, uint n)
5480 {
5481     uint tot, m;
5482     struct buf *bp;
5483
5484     if(ip->type == T_DEV){
5485         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
5486             return -1;
5487         return devsw[ip->major].write(ip, src, n);
5488     }
5489
5490     if(off > ip->size || off + n < off)
5491         return -1;
5492     if(off + n > MAXFILE*BSIZE)
5493         return -1;
5494
5495     for(tot=0; tot<n; tot+=m, off+=m, src+=m){
5496         bp = bread(ip->dev, bmap(ip, off/BSIZE));
5497         m = min(n - tot, BSIZE - off%BSIZE);
5498         memmove(bp->data + off%BSIZE, src, m);
5499         log_write(bp);

```

```

5500     brelse(bp);
5501 }
5502
5503 if(n > 0 && off > ip->size){
5504     ip->size = off;
5505     iupdate(ip);
5506 }
5507 return n;
5508 }
5509
5510 // Directories
5511
5512 int
5513 namecmp(const char *s, const char *t)
5514 {
5515     return strncmp(s, t, DIRSIZ);
5516 }
5517
5518 // Look for a directory entry in a directory.
5519 // If found, set *poff to byte offset of entry.
5520 struct inode*
5521 dirlookup(struct inode *dp, char *name, uint *poff)
5522 {
5523     uint off, inum;
5524     struct dirent de;
5525
5526     if(dp->type != T_DIR)
5527         panic("dirlookup not DIR");
5528
5529     for(off = 0; off < dp->size; off += sizeof(de)){
5530         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5531             panic("dirlink read");
5532         if(de.inum == 0)
5533             continue;
5534         if(namecmp(name, de.name) == 0){
5535             // entry matches path element
5536             if(poff)
5537                 *poff = off;
5538             inum = de.inum;
5539             return iget(dp->dev, inum);
5540         }
5541     }
5542
5543     return 0;
5544 }
5545
5546
5547
5548
5549

```

```

5550 // Write a new directory entry (name, inum) into the directory dp.
5551 int
5552 dirlink(struct inode *dp, char *name, uint inum)
5553 {
5554     int off;
5555     struct dirent de;
5556     struct inode *ip;
5557
5558     // Check that name is not present.
5559     if((ip = dirlookup(dp, name, 0)) != 0){
5560         iput(ip);
5561         return -1;
5562     }
5563
5564     // Look for an empty dirent.
5565     for(off = 0; off < dp->size; off += sizeof(de)){
5566         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5567             panic("dirlink read");
5568         if(de.inum == 0)
5569             break;
5570     }
5571
5572     strncpy(de.name, name, DIRSIZ);
5573     de.inum = inum;
5574     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
5575         panic("dirlink");
5576
5577     return 0;
5578 }
5579
5580 // Paths
5581
5582 // Copy the next path element from path into name.
5583 // Return a pointer to the element following the copied one.
5584 // The returned path has no leading slashes,
5585 // so the caller can check *path=='\0' to see if the name is the last one.
5586 // If no name to remove, return 0.
5587 //
5588 // Examples:
5589 //   skipelem("a/bb/c", name) = "bb/c", setting name = "a"
5590 //   skipelem("///a//bb", name) = "bb", setting name = "a"
5591 //   skipelem("a", name) = "", setting name = "a"
5592 //   skipelem("", name) = skipelem("///", name) = 0
5593 //
5594 static char*
5595 skipelem(char *path, char *name)
5596 {
5597     char *s;
5598     int len;
5599

```



```

5600 while(*path == '/')
5601     path++;
5602 if(*path == 0)
5603     return 0;
5604 s = path;
5605 while(*path != '/' && *path != 0)
5606     path++;
5607 len = path - s;
5608 if(len >= DIRSIZ)
5609     memmove(name, s, DIRSIZ);
5610 else {
5611     memmove(name, s, len);
5612     name[len] = 0;
5613 }
5614 while(*path == '/')
5615     path++;
5616 return path;
5617 }
5618
5619 // Look up and return the inode for a path name.
5620 // If parent != 0, return the inode for the parent and copy the final
5621 // path element into name, which must have room for DIRSIZ bytes.
5622 // Must be called inside a transaction since it calls iput().
5623 static struct inode*
5624 namex(char *path, int nameparent, char *name)
5625 {
5626     struct inode *ip, *next;
5627
5628     if(*path == '/')
5629         ip = iget(ROOTDEV, ROOTINO);
5630     else
5631         ip = idup(proc->cwd);
5632
5633     while((path = skipelem(path, name)) != 0){
5634         ilock(ip);
5635         if(ip->type != T_DIR){
5636             iunlockput(ip);
5637             return 0;
5638         }
5639         if(nameparent && *path == '\0'){
5640             // Stop one level early.
5641             iunlock(ip);
5642             return ip;
5643         }
5644         if((next = dirlookup(ip, name, 0)) == 0){
5645             iunlockput(ip);
5646             return 0;
5647         }
5648         iunlockput(ip);
5649         ip = next;

```

```

5650     }
5651     if(nameparent){
5652         iput(ip);
5653         return 0;
5654     }
5655     return ip;
5656 }
5657
5658 struct inode*
5659 namei(char *path)
5660 {
5661     char name[DIRSIZ];
5662     return namex(path, 0, name);
5663 }
5664
5665 struct inode*
5666 nameiparent(char *path, char *name)
5667 {
5668     return namex(path, 1, name);
5669 }
5670
5671
5672
5673
5674
5675
5676
5677
5678
5679
5680
5681
5682
5683
5684
5685
5686
5687
5688
5689
5690
5691
5692
5693
5694
5695
5696
5697
5698
5699

```

```

5700 //
5701 // File descriptors
5702 //
5703
5704 #include "types.h"
5705 #include "defs.h"
5706 #include "param.h"
5707 #include "fs.h"
5708 #include "file.h"
5709 #include "spinlock.h"
5710
5711 struct devsw devsw[NDEV];
5712 struct {
5713     struct spinlock lock;
5714     struct file file[NFILE];
5715 } ftable;
5716
5717 void
5718 fileinit(void)
5719 {
5720     initlock(&ftable.lock, "ftable");
5721 }
5722
5723 // Allocate a file structure.
5724 struct file*
5725 filealloc(void)
5726 {
5727     struct file *f;
5728
5729     acquire(&ftable.lock);
5730     for(f = ftable.file; f < ftable.file + NFILE; f++){
5731         if(f->ref == 0){
5732             f->ref = 1;
5733             release(&ftable.lock);
5734             return f;
5735         }
5736     }
5737     release(&ftable.lock);
5738     return 0;
5739 }
5740
5741
5742
5743
5744
5745
5746
5747
5748
5749

```

```

5750 // Increment ref count for file f.
5751 struct file*
5752 filedup(struct file *f)
5753 {
5754     acquire(&ftable.lock);
5755     if(f->ref < 1)
5756         panic("filedup");
5757     f->ref++;
5758     release(&ftable.lock);
5759     return f;
5760 }
5761
5762 // Close file f. (Decrement ref count, close when reaches 0.)
5763 void
5764 fileclose(struct file *f)
5765 {
5766     struct file ff;
5767
5768     acquire(&ftable.lock);
5769     if(f->ref < 1)
5770         panic("fileclose");
5771     if(--f->ref > 0){
5772         release(&ftable.lock);
5773         return;
5774     }
5775     ff = *f;
5776     f->ref = 0;
5777     f->type = FD_NONE;
5778     release(&ftable.lock);
5779
5780     if(ff.type == FD_PIPE)
5781         pipeclose(ff.pipe, ff.writable);
5782     else if(ff.type == FD_INODE){
5783         begin_op();
5784         iput(ff.ip);
5785         end_op();
5786     }
5787 }
5788
5789
5790
5791
5792
5793
5794
5795
5796
5797
5798
5799

```

```

5800 // Get metadata about file f.
5801 int
5802 filestat(struct file *f, struct stat *st)
5803 {
5804     if(f->type == FD_INODE){
5805         ilock(f->ip);
5806         stati(f->ip, st);
5807         iunlock(f->ip);
5808         return 0;
5809     }
5810     return -1;
5811 }
5812
5813 // Read from file f.
5814 int
5815 fileread(struct file *f, char *addr, int n)
5816 {
5817     int r;
5818
5819     if(f->readable == 0)
5820         return -1;
5821     if(f->type == FD_PIPE)
5822         return piperead(f->pipe, addr, n);
5823     if(f->type == FD_INODE){
5824         ilock(f->ip);
5825         if((r = readi(f->ip, addr, f->off, n)) > 0)
5826             f->off += r;
5827         iunlock(f->ip);
5828         return r;
5829     }
5830     panic("fileread");
5831 }
5832
5833 // Write to file f.
5834 int
5835 filewrite(struct file *f, char *addr, int n)
5836 {
5837     int r;
5838
5839     if(f->writable == 0)
5840         return -1;
5841     if(f->type == FD_PIPE)
5842         return pipewrite(f->pipe, addr, n);
5843     if(f->type == FD_INODE){
5844         // write a few blocks at a time to avoid exceeding
5845         // the maximum log transaction size, including
5846         // i-node, indirect block, allocation blocks,
5847         // and 2 blocks of slop for non-aligned writes.
5848         // this really belongs lower down, since writei()
5849         // might be writing a device like the console.

```

```

5850     int max = ((LOGSIZE-1-1-2) / 2) * 512;
5851     int i = 0;
5852     while(i < n){
5853         int n1 = n - i;
5854         if(n1 > max)
5855             n1 = max;
5856
5857         begin_op();
5858         ilock(f->ip);
5859         if ((r = writei(f->ip, addr + i, f->off, n1)) > 0)
5860             f->off += r;
5861         iunlock(f->ip);
5862         end_op();
5863
5864         if(r < 0)
5865             break;
5866         if(r != n1)
5867             panic("short filewrite");
5868         i += r;
5869     }
5870     return i == n ? n : -1;
5871 }
5872 panic("filewrite");
5873 }
5874
5875
5876
5877
5878
5879
5880
5881
5882
5883
5884
5885
5886
5887
5888
5889
5890
5891
5892
5893
5894
5895
5896
5897
5898
5899

```

```

5900 //
5901 // File-system system calls.
5902 // Mostly argument checking, since we don't trust
5903 // user code, and calls into file.c and fs.c.
5904 //
5905
5906 #include "types.h"
5907 #include "defs.h"
5908 #include "param.h"
5909 #include "stat.h"
5910 #include "mmu.h"
5911 #include "proc.h"
5912 #include "fs.h"
5913 #include "file.h"
5914 #include "fcntl.h"
5915
5916 // Fetch the nth word-sized system call argument as a file descriptor
5917 // and return both the descriptor and the corresponding struct file.
5918 static int
5919 argfd(int n, int *pfd, struct file **pf)
5920 {
5921     int fd;
5922     struct file *f;
5923
5924     if(argint(n, &fd) < 0)
5925         return -1;
5926     if(fd < 0 || fd >= NOFILE || (f=proc->ofile[fd]) == 0)
5927         return -1;
5928     if(pfd)
5929         *pfd = fd;
5930     if(pf)
5931         *pf = f;
5932     return 0;
5933 }
5934
5935 // Allocate a file descriptor for the given file.
5936 // Takes over file reference from caller on success.
5937 static int
5938 fdalloc(struct file *f)
5939 {
5940     int fd;
5941
5942     for(fd = 0; fd < NOFILE; fd++){
5943         if(proc->ofile[fd] == 0){
5944             proc->ofile[fd] = f;
5945             return fd;
5946         }
5947     }
5948     return -1;
5949 }

```

```

5950 int
5951 sys_dup(void)
5952 {
5953     struct file *f;
5954     int fd;
5955
5956     if(argfd(0, 0, &f) < 0)
5957         return -1;
5958     if((fd=fdalloc(f)) < 0)
5959         return -1;
5960     filedup(f);
5961     return fd;
5962 }
5963
5964 int
5965 sys_read(void)
5966 {
5967     struct file *f;
5968     int n;
5969     char *p;
5970
5971     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
5972         return -1;
5973     return fileread(f, p, n);
5974 }
5975
5976 int
5977 sys_write(void)
5978 {
5979     struct file *f;
5980     int n;
5981     char *p;
5982
5983     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
5984         return -1;
5985     return filewrite(f, p, n);
5986 }
5987
5988 int
5989 sys_close(void)
5990 {
5991     int fd;
5992     struct file *f;
5993
5994     if(argfd(0, &fd, &f) < 0)
5995         return -1;
5996     proc->ofile[fd] = 0;
5997     fileclose(f);
5998     return 0;
5999 }

```

```

6000 int
6001 sys_fstat(void)
6002 {
6003     struct file *f;
6004     struct stat *st;
6005
6006     if(argfd(0, 0, &f) < 0 || argptr(1, (void*)&st, sizeof(*st)) < 0)
6007         return -1;
6008     return filestat(f, st);
6009 }
6010
6011 // Create the path new as a link to the same inode as old.
6012 int
6013 sys_link(void)
6014 {
6015     char name[DIRSIZ], *new, *old;
6016     struct inode *dp, *ip;
6017
6018     if(argstr(0, &old) < 0 || argstr(1, &new) < 0)
6019         return -1;
6020
6021     begin_op();
6022     if((ip = namei(old)) == 0){
6023         end_op();
6024         return -1;
6025     }
6026
6027     ilock(ip);
6028     if(ip->type == T_DIR){
6029         iunlockput(ip);
6030         end_op();
6031         return -1;
6032     }
6033
6034     ip->nlink++;
6035     iupdate(ip);
6036     iunlock(ip);
6037
6038     if((dp = nameiparent(new, name)) == 0)
6039         goto bad;
6040     ilock(dp);
6041     if(dp->dev != ip->dev || dirlink(dp, name, ip->inum) < 0){
6042         iunlockput(dp);
6043         goto bad;
6044     }
6045     iunlockput(dp);
6046     iput(ip);
6047
6048     end_op();
6049

```

```

6050     return 0;
6051
6052 bad:
6053     ilock(ip);
6054     ip->nlink--;
6055     iupdate(ip);
6056     iunlockput(ip);
6057     end_op();
6058     return -1;
6059 }
6060
6061 // Is the directory dp empty except for "." and ".." ?
6062 static int
6063 isdirempty(struct inode *dp)
6064 {
6065     int off;
6066     struct dirent de;
6067
6068     for(off=2*sizeof(de); off<dp->size; off+=sizeof(de)){
6069         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
6070             panic("isdirempty: readi");
6071         if(de.inum != 0)
6072             return 0;
6073     }
6074     return 1;
6075 }
6076
6077 int
6078 sys_unlink(void)
6079 {
6080     struct inode *ip, *dp;
6081     struct dirent de;
6082     char name[DIRSIZ], *path;
6083     uint off;
6084
6085     if(argstr(0, &path) < 0)
6086         return -1;
6087
6088     begin_op();
6089     if((dp = nameiparent(path, name)) == 0){
6090         end_op();
6091         return -1;
6092     }
6093
6094     ilock(dp);
6095
6096     // Cannot unlink "." or "..".
6097     if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
6098         goto bad;
6099

```

```

6100 if((ip = dirlookup(dp, name, &off)) == 0)
6101     goto bad;
6102 ilock(ip);
6103
6104 if(ip->nlink < 1)
6105     panic("unlink: nlink < 1");
6106 if(ip->type == T_DIR && !isdirempty(ip)){
6107     iunlockput(ip);
6108     goto bad;
6109 }
6110
6111 memset(&de, 0, sizeof(de));
6112 if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
6113     panic("unlink: writei");
6114 if(ip->type == T_DIR){
6115     dp->nlink--;
6116     iupdate(dp);
6117 }
6118 iunlockput(dp);
6119
6120 ip->nlink--;
6121 iupdate(ip);
6122 iunlockput(ip);
6123
6124 end_op();
6125
6126 return 0;
6127
6128 bad:
6129 iunlockput(dp);
6130 end_op();
6131 return -1;
6132 }
6133
6134 static struct inode*
6135 create(char *path, short type, short major, short minor)
6136 {
6137     uint off;
6138     struct inode *ip, *dp;
6139     char name[DIRSIZ];
6140
6141     if((dp = nameiparent(path, name)) == 0)
6142         return 0;
6143     ilock(dp);
6144
6145     if((ip = dirlookup(dp, name, &off)) != 0){
6146         iunlockput(dp);
6147         ilock(ip);
6148         if(type == T_FILE && ip->type == T_FILE)
6149             return ip;

```

```

6150     iunlockput(ip);
6151     return 0;
6152 }
6153
6154 if((ip = ialloc(dp->dev, type)) == 0)
6155     panic("create: ialloc");
6156
6157 ilock(ip);
6158 ip->major = major;
6159 ip->minor = minor;
6160 ip->nlink = 1;
6161 iupdate(ip);
6162
6163 if(type == T_DIR){ // Create . and .. entries.
6164     dp->nlink++; // for "."
6165     iupdate(dp);
6166     // No ip->nlink++ for ".": avoid cyclic ref count.
6167     if(dirlink(ip, ".", ip->inum) < 0 || dirlink(ip, "..", dp->inum) < 0)
6168         panic("create dots");
6169 }
6170
6171 if(dirlink(dp, name, ip->inum) < 0)
6172     panic("create: dirlink");
6173
6174 iunlockput(dp);
6175
6176 return ip;
6177 }
6178
6179 int
6180 sys_open(void)
6181 {
6182     char *path;
6183     int fd, omode;
6184     struct file *f;
6185     struct inode *ip;
6186
6187     if(argstr(0, &path) < 0 || argint(1, &omode) < 0)
6188         return -1;
6189
6190     begin_op();
6191
6192     if(omode & O_CREATE){
6193         ip = create(path, T_FILE, 0, 0);
6194         if(ip == 0){
6195             end_op();
6196             return -1;
6197         }
6198     } else {
6199         if((ip = namei(path)) == 0){

```

```

6200     end_op();
6201     return -1;
6202 }
6203 ilock(ip);
6204 if(ip->type == T_DIR && omode != O_RDONLY){
6205     iunlockput(ip);
6206     end_op();
6207     return -1;
6208 }
6209 }
6210
6211 if((f = filealloc()) == 0 || (fd = fdalloc(f)) < 0){
6212     if(f)
6213         fileclose(f);
6214     iunlockput(ip);
6215     end_op();
6216     return -1;
6217 }
6218 iunlock(ip);
6219 end_op();
6220
6221 f->type = FD_INODE;
6222 f->ip = ip;
6223 f->off = 0;
6224 f->readable = !(omode & O_WRONLY);
6225 f->writable = (omode & O_WRONLY) || (omode & O_RDWR);
6226 return fd;
6227 }
6228
6229 int
6230 sys_mkdir(void)
6231 {
6232     char *path;
6233     struct inode *ip;
6234
6235     begin_op();
6236     if(argstr(0, &path) < 0 || (ip = create(path, T_DIR, 0, 0)) == 0){
6237         end_op();
6238         return -1;
6239     }
6240     iunlockput(ip);
6241     end_op();
6242     return 0;
6243 }
6244
6245
6246
6247
6248
6249

```

```

6250 int
6251 sys_mknod(void)
6252 {
6253     struct inode *ip;
6254     char *path;
6255     int len;
6256     int major, minor;
6257
6258     begin_op();
6259     if((len=argstr(0, &path)) < 0 ||
6260         argint(1, &major) < 0 ||
6261         argint(2, &minor) < 0 ||
6262         (ip = create(path, T_DEV, major, minor)) == 0){
6263         end_op();
6264         return -1;
6265     }
6266     iunlockput(ip);
6267     end_op();
6268     return 0;
6269 }
6270
6271 int
6272 sys_chdir(void)
6273 {
6274     char *path;
6275     struct inode *ip;
6276
6277     begin_op();
6278     if(argstr(0, &path) < 0 || (ip = namei(path)) == 0){
6279         end_op();
6280         return -1;
6281     }
6282     ilock(ip);
6283     if(ip->type != T_DIR){
6284         iunlockput(ip);
6285         end_op();
6286         return -1;
6287     }
6288     iunlock(ip);
6289     iput(proc->cwd);
6290     end_op();
6291     proc->cwd = ip;
6292     return 0;
6293 }
6294
6295
6296
6297
6298
6299

```

```

6300 int
6301 sys_exec(void)
6302 {
6303     char *path, *argv[MAXARG];
6304     int i;
6305     uint uargv, uarg;
6306
6307     if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
6308         return -1;
6309     }
6310     memset(argv, 0, sizeof(argv));
6311     for(i=0; i++){
6312         if(i >= NELEM(argv))
6313             return -1;
6314         if(fetchint(uargv+4*i, (int*)&uarg) < 0)
6315             return -1;
6316         if(uarg == 0){
6317             argv[i] = 0;
6318             break;
6319         }
6320         if(fetchstr(uarg, &argv[i]) < 0)
6321             return -1;
6322     }
6323     return exec(path, argv);
6324 }
6325
6326 int
6327 sys_pipe(void)
6328 {
6329     int *fd;
6330     struct file *rf, *wf;
6331     int fd0, fd1;
6332
6333     if(argptr(0, (void*)&fd, 2*sizeof(fd[0])) < 0)
6334         return -1;
6335     if(pipealloc(&rf, &wf) < 0)
6336         return -1;
6337     fd0 = -1;
6338     if((fd0 = fdalloc(rf)) < 0 || (fd1 = fdalloc(wf)) < 0){
6339         if(fd0 >= 0)
6340             proc->ofile[fd0] = 0;
6341         fileclose(rf);
6342         fileclose(wf);
6343         return -1;
6344     }
6345     fd[0] = fd0;
6346     fd[1] = fd1;
6347     return 0;
6348 }
6349

```

```

6350 #include "types.h"
6351 #include "param.h"
6352 #include "memlayout.h"
6353 #include "mmu.h"
6354 #include "proc.h"
6355 #include "defs.h"
6356 #include "x86.h"
6357 #include "elf.h"
6358
6359 int
6360 exec(char *path, char **argv)
6361 {
6362     char *s, *last;
6363     int i, off;
6364     uint argc, sz, sp, ustack[3+MAXARG+1];
6365     struct elfhdr elf;
6366     struct inode *ip;
6367     struct proghdr ph;
6368     pde_t *pgdir, *oldpgdir;
6369
6370     begin_op();
6371     if((ip = namei(path)) == 0){
6372         end_op();
6373         return -1;
6374     }
6375     ilock(ip);
6376     pgdir = 0;
6377
6378     // Check ELF header
6379     if(readi(ip, (char*)&elf, 0, sizeof(elf)) < sizeof(elf))
6380         goto bad;
6381     if(elf.magic != ELF_MAGIC)
6382         goto bad;
6383
6384     if((pgdir = setupkvm()) == 0)
6385         goto bad;
6386
6387     // Load program into memory.
6388     sz = 0;
6389     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6390         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6391             goto bad;
6392         if(ph.type != ELF_PROG_LOAD)
6393             continue;
6394         if(ph.memsz < ph.filesz)
6395             goto bad;
6396         if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
6397             goto bad;
6398         if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
6399             goto bad;

```



```

6400 }
6401 iunlockput(ip);
6402 end_op();
6403 ip = 0;
6404
6405 // Allocate two pages at the next page boundary.
6406 // Make the first inaccessible. Use the second as the user stack.
6407 sz = PGROUNDUP(sz);
6408 if((sz = allocvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
6409     goto bad;
6410 clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
6411 sp = sz;
6412
6413 // Push argument strings, prepare rest of stack in ustack.
6414 for(argc = 0; argv[argc]; argc++) {
6415     if(argc >= MAXARG)
6416         goto bad;
6417     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
6418     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
6419         goto bad;
6420     ustack[3+argc] = sp;
6421 }
6422 ustack[3+argc] = 0;
6423
6424 ustack[0] = 0xffffffff; // fake return PC
6425 ustack[1] = argc;
6426 ustack[2] = sp - (argc+1)*4; // argv pointer
6427
6428 sp -= (3+argc+1) * 4;
6429 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
6430     goto bad;
6431
6432 // Save program name for debugging.
6433 for(last=s=path; *s; s++)
6434     if(*s == '/')
6435         last = s+1;
6436 safestrcpy(proc->name, last, sizeof(proc->name));
6437
6438 // Commit to the user image.
6439 oldpgdir = proc->pgdir;
6440 proc->pgdir = pgdir;
6441 proc->sz = sz;
6442 proc->tf->eip = elf.entry; // main
6443 proc->tf->esp = sp;
6444 switchvm(proc);
6445 freevm(oldpgdir);
6446 return 0;
6447
6448
6449

```

```

6450 bad:
6451     if(pgdir)
6452         freevm(pgdir);
6453     if(ip){
6454         iunlockput(ip);
6455         end_op();
6456     }
6457     return -1;
6458 }
6459
6460
6461
6462
6463
6464
6465
6466
6467
6468
6469
6470
6471
6472
6473
6474
6475
6476
6477
6478
6479
6480
6481
6482
6483
6484
6485
6486
6487
6488
6489
6490
6491
6492
6493
6494
6495
6496
6497
6498
6499

```

```

6500 #include "types.h"
6501 #include "defs.h"
6502 #include "param.h"
6503 #include "mmu.h"
6504 #include "proc.h"
6505 #include "fs.h"
6506 #include "file.h"
6507 #include "spinlock.h"
6508
6509 #define PIPESIZE 512
6510
6511 struct pipe {
6512     struct spinlock lock;
6513     char data[PIPESIZE];
6514     uint nread; // number of bytes read
6515     uint nwrite; // number of bytes written
6516     int readopen; // read fd is still open
6517     int writeopen; // write fd is still open
6518 };
6519
6520 int
6521 pipealloc(struct file **f0, struct file **f1)
6522 {
6523     struct pipe *p;
6524
6525     p = 0;
6526     *f0 = *f1 = 0;
6527     if((*f0 = filealloc()) == 0 || (*f1 = filealloc()) == 0)
6528         goto bad;
6529     if((p = (struct pipe*)kalloc()) == 0)
6530         goto bad;
6531     p->readopen = 1;
6532     p->writeopen = 1;
6533     p->nwrite = 0;
6534     p->nread = 0;
6535     initlock(&p->lock, "pipe");
6536     (*f0)->type = FD_PIPE;
6537     (*f0)->readable = 1;
6538     (*f0)->writable = 0;
6539     (*f0)->pipe = p;
6540     (*f1)->type = FD_PIPE;
6541     (*f1)->readable = 0;
6542     (*f1)->writable = 1;
6543     (*f1)->pipe = p;
6544     return 0;
6545
6546
6547
6548
6549

```

```

6550 bad:
6551     if(p)
6552         kfree((char*)p);
6553     if(*f0)
6554         fileclose(*f0);
6555     if(*f1)
6556         fileclose(*f1);
6557     return -1;
6558 }
6559
6560 void
6561 pipeclose(struct pipe *p, int writable)
6562 {
6563     acquire(&p->lock);
6564     if(writable){
6565         p->writeopen = 0;
6566         wakeup(&p->nread);
6567     } else {
6568         p->readopen = 0;
6569         wakeup(&p->nwrite);
6570     }
6571     if(p->readopen == 0 && p->writeopen == 0){
6572         release(&p->lock);
6573         kfree((char*)p);
6574     } else
6575         release(&p->lock);
6576 }
6577
6578 int
6579 pipewrite(struct pipe *p, char *addr, int n)
6580 {
6581     int i;
6582
6583     acquire(&p->lock);
6584     for(i = 0; i < n; i++){
6585         while(p->nwrite == p->nread + PIPESIZE){
6586             if(p->readopen == 0 || proc->killed){
6587                 release(&p->lock);
6588                 return -1;
6589             }
6590             wakeup(&p->nread);
6591             sleep(&p->nwrite, &p->lock);
6592         }
6593         p->data[p->nwrite++ % PIPESIZE] = addr[i];
6594     }
6595     wakeup(&p->nread);
6596     release(&p->lock);
6597     return n;
6598 }
6599

```

```

6600 int
6601 piperead(struct pipe *p, char *addr, int n)
6602 {
6603     int i;
6604
6605     acquire(&p->lock);
6606     while(p->nread == p->nwrite && p->writeopen){
6607         if(proc->killed){
6608             release(&p->lock);
6609             return -1;
6610         }
6611         sleep(&p->nread, &p->lock);
6612     }
6613     for(i = 0; i < n; i++){
6614         if(p->nread == p->nwrite)
6615             break;
6616         addr[i] = p->data[p->nread++ % PIPESIZE];
6617     }
6618     wakeup(&p->nwrite);
6619     release(&p->lock);
6620     return i;
6621 }
6622
6623
6624
6625
6626
6627
6628
6629
6630
6631
6632
6633
6634
6635
6636
6637
6638
6639
6640
6641
6642
6643
6644
6645
6646
6647
6648
6649

```

```

6650 #include "types.h"
6651 #include "x86.h"
6652
6653 void*
6654 memset(void *dst, int c, uint n)
6655 {
6656     if ((int)dst%4 == 0 && n%4 == 0){
6657         c &= 0xFF;
6658         stosl(dst, (c<<24)|(c<<16)|(c<<8)|c, n/4);
6659     } else
6660         stosb(dst, c, n);
6661     return dst;
6662 }
6663
6664 int
6665 memcmp(const void *v1, const void *v2, uint n)
6666 {
6667     const uchar *s1, *s2;
6668
6669     s1 = v1;
6670     s2 = v2;
6671     while(n-- > 0){
6672         if(*s1 != *s2)
6673             return *s1 - *s2;
6674         s1++, s2++;
6675     }
6676
6677     return 0;
6678 }
6679
6680 void*
6681 memmove(void *dst, const void *src, uint n)
6682 {
6683     const char *s;
6684     char *d;
6685
6686     s = src;
6687     d = dst;
6688     if(s < d && s + n > d){
6689         s += n;
6690         d += n;
6691         while(n-- > 0)
6692             *--d = *--s;
6693     } else
6694         while(n-- > 0)
6695             *d++ = *s++;
6696
6697     return dst;
6698 }
6699

```

```

6700 // memcpy exists to placate GCC. Use memmove.
6701 void*
6702 memcpy(void *dst, const void *src, uint n)
6703 {
6704     return memmove(dst, src, n);
6705 }
6706
6707 int
6708 strncmp(const char *p, const char *q, uint n)
6709 {
6710     while(n > 0 && *p && *p == *q)
6711         n--, p++, q++;
6712     if(n == 0)
6713         return 0;
6714     return (uchar)*p - (uchar)*q;
6715 }
6716
6717 char*
6718 strncpy(char *s, const char *t, int n)
6719 {
6720     char *os;
6721
6722     os = s;
6723     while(n-- > 0 && (*s++ = *t++) != 0)
6724         ;
6725     while(n-- > 0)
6726         *s++ = 0;
6727     return os;
6728 }
6729
6730 // Like strncpy but guaranteed to NUL-terminate.
6731 char*
6732 safestrcpy(char *s, const char *t, int n)
6733 {
6734     char *os;
6735
6736     os = s;
6737     if(n <= 0)
6738         return os;
6739     while(--n > 0 && (*s++ = *t++) != 0)
6740         ;
6741     *s = 0;
6742     return os;
6743 }
6744
6745
6746
6747
6748
6749

```

```

6750 int
6751 strlen(const char *s)
6752 {
6753     int n;
6754
6755     for(n = 0; s[n]; n++)
6756         ;
6757     return n;
6758 }
6759
6760
6761
6762
6763
6764
6765
6766
6767
6768
6769
6770
6771
6772
6773
6774
6775
6776
6777
6778
6779
6780
6781
6782
6783
6784
6785
6786
6787
6788
6789
6790
6791
6792
6793
6794
6795
6796
6797
6798
6799

```

```

6800 // See MultiProcessor Specification Version 1.[14]
6801
6802 struct mp {           // floating pointer
6803     uchar signature[4]; // "_MP_"
6804     void *physaddr;      // phys addr of MP config table
6805     uchar length;        // 1
6806     uchar specrev;       // [14]
6807     uchar checksum;      // all bytes must add up to 0
6808     uchar type;          // MP system config type
6809     uchar imcrp;
6810     uchar reserved[3];
6811 };
6812
6813 struct mpconf {        // configuration table header
6814     uchar signature[4]; // "PCMP"
6815     ushort length;      // total table length
6816     uchar version;      // [14]
6817     uchar checksum;     // all bytes must add up to 0
6818     uchar product[20];  // product id
6819     uint *oemtable;     // OEM table pointer
6820     ushort oemlength;   // OEM table length
6821     ushort entry;       // entry count
6822     uint *lapicaddr;    // address of local APIC
6823     ushort xlength;     // extended table length
6824     uchar xchecksum;    // extended table checksum
6825     uchar reserved;
6826 };
6827
6828 struct mpproc {        // processor table entry
6829     uchar type;         // entry type (0)
6830     uchar apicid;       // local APIC id
6831     uchar version;      // local APIC verison
6832     uchar flags;        // CPU flags
6833     #define MPBOOT 0x02 // This proc is the bootstrap processor.
6834     uchar signature[4]; // CPU signature
6835     uint feature;       // feature flags from CPUID instruction
6836     uchar reserved[8];
6837 };
6838
6839 struct mpioapic {      // I/O APIC table entry
6840     uchar type;         // entry type (2)
6841     uchar apicno;       // I/O APIC id
6842     uchar version;      // I/O APIC version
6843     uchar flags;        // I/O APIC flags
6844     uint *addr;         // I/O APIC address
6845 };
6846
6847
6848
6849

```

```

6850 // Table entry types
6851 #define MPPROC 0x00 // One per processor
6852 #define MPBUS 0x01 // One per bus
6853 #define MPIOAPIC 0x02 // One per I/O APIC
6854 #define MPIOINTR 0x03 // One per bus interrupt source
6855 #define MPLINTR 0x04 // One per system interrupt source
6856
6857 // Blank page.
6858
6859
6860
6861
6862
6863
6864
6865
6866
6867
6868
6869
6870
6871
6872
6873
6874
6875
6876
6877
6878
6879
6880
6881
6882
6883
6884
6885
6886
6887
6888
6889
6890
6891
6892
6893
6894
6895
6896
6897
6898
6899

```

```

6900 // Multiprocessor support
6901 // Search memory for MP description structures.
6902 // http://developer.intel.com/design/pentium/datashts/24201606.pdf
6903
6904 #include "types.h"
6905 #include "defs.h"
6906 #include "param.h"
6907 #include "memlayout.h"
6908 #include "mp.h"
6909 #include "x86.h"
6910 #include "mmu.h"
6911 #include "proc.h"
6912
6913 struct cpu cpus[NCPU];
6914 static struct cpu *bcpu;
6915 int ismp;
6916 int ncpu;
6917 uchar ioapicid;
6918
6919 int
6920 mpbcpu(void)
6921 {
6922     return bcpu-cpus;
6923 }
6924
6925 static uchar
6926 sum(uchar *addr, int len)
6927 {
6928     int i, sum;
6929
6930     sum = 0;
6931     for(i=0; i<len; i++)
6932         sum += addr[i];
6933     return sum;
6934 }
6935
6936 // Look for an MP structure in the len bytes at addr.
6937 static struct mp*
6938 mpsearch1(uint a, int len)
6939 {
6940     uchar *e, *p, *addr;
6941
6942     addr = p2v(a);
6943     e = addr+len;
6944     for(p = addr; p < e; p += sizeof(struct mp))
6945         if(memcmp(p, "_MP_", 4) == 0 && sum(p, sizeof(struct mp)) == 0)
6946             return (struct mp*)p;
6947     return 0;
6948 }
6949
```

```

6950 // Search for the MP Floating Pointer Structure, which according to the
6951 // spec is in one of the following three locations:
6952 // 1) in the first KB of the EBDA;
6953 // 2) in the last KB of system base memory;
6954 // 3) in the BIOS ROM between 0xE0000 and 0xFFFFF.
6955 static struct mp*
6956 mpsearch(void)
6957 {
6958     uchar *bda;
6959     uint p;
6960     struct mp *mp;
6961
6962     bda = (uchar *) P2V(0x400);
6963     if((p = ((bda[0x0F]<<8) | bda[0x0E]) << 4)){
6964         if((mp = mpsearch1(p, 1024)))
6965             return mp;
6966     } else {
6967         p = ((bda[0x14]<<8) | bda[0x13])*1024;
6968         if((mp = mpsearch1(p-1024, 1024)))
6969             return mp;
6970     }
6971     return mpsearch1(0xF0000, 0x10000);
6972 }
6973
6974 // Search for an MP configuration table. For now,
6975 // don't accept the default configurations (physaddr == 0).
6976 // Check for correct signature, calculate the checksum and,
6977 // if correct, check the version.
6978 // To do: check extended table checksum.
6979 static struct mpconf*
6980 mpconfig(struct mp **pmp)
6981 {
6982     struct mpconf *conf;
6983     struct mp *mp;
6984
6985     if((mp = mpsearch()) == 0 || mp->physaddr == 0)
6986         return 0;
6987     conf = (struct mpconf*) p2v((uint) mp->physaddr);
6988     if(memcmp(conf, "PCMP", 4) != 0)
6989         return 0;
6990     if(conf->version != 1 && conf->version != 4)
6991         return 0;
6992     if(sum((uchar*)conf, conf->length) != 0)
6993         return 0;
6994     *pmp = mp;
6995     return conf;
6996 }
6997
6998
6999

```

```

7000 void
7001 mpinit(void)
7002 {
7003     uchar *p, *e;
7004     struct mp *mp;
7005     struct mpconf *conf;
7006     struct mpproc *proc;
7007     struct mpioapic *ioapic;
7008
7009     bcpu = &cpus[0];
7010     if((conf = mpconfig(&mp)) == 0)
7011         return;
7012     ismp = 1;
7013     lapic = (uint*)conf->lapicaddr;
7014     for(p=(uchar*)(conf+1), e=(uchar*)conf+conf->length; p<e; ){
7015         switch(*p){
7016             case MPPROC:
7017                 proc = (struct mpproc*)p;
7018                 if(ncpu != proc->apicid){
7019                     cprintf("mpinit: ncpu=%d apicid=%d\n", ncpu, proc->apicid);
7020                     ismp = 0;
7021                 }
7022                 if(proc->flags & MPBOOT)
7023                     bcpu = &cpus[ncpu];
7024                 cpus[ncpu].id = ncpu;
7025                 ncpu++;
7026                 p += sizeof(struct mpproc);
7027                 continue;
7028             case MPIOAPIC:
7029                 ioapic = (struct mpioapic*)p;
7030                 ioapicid = ioapic->apicno;
7031                 p += sizeof(struct mpioapic);
7032                 continue;
7033             case MPBUS:
7034             case MPIOINTR:
7035             case MPLINTR:
7036                 p += 8;
7037                 continue;
7038             default:
7039                 cprintf("mpinit: unknown config type %x\n", *p);
7040                 ismp = 0;
7041         }
7042     }
7043     if(!ismp){
7044         // Didn't like what we found; fall back to no MP.
7045         ncpu = 1;
7046         lapic = 0;
7047         ioapicid = 0;
7048         return;
7049     }

```

```

7050     if(mp->imcrp){
7051         // Bochs doesn't support IMCR, so this doesn't run on Bochs.
7052         // But it would on real hardware.
7053         outb(0x22, 0x70); // Select IMCR
7054         outb(0x23, inb(0x23) | 1); // Mask external interrupts.
7055     }
7056 }
7057
7058
7059
7060
7061
7062
7063
7064
7065
7066
7067
7068
7069
7070
7071
7072
7073
7074
7075
7076
7077
7078
7079
7080
7081
7082
7083
7084
7085
7086
7087
7088
7089
7090
7091
7092
7093
7094
7095
7096
7097
7098
7099

```

```

7100 // The local APIC manages internal (non-I/O) interrupts.
7101 // See Chapter 8 & Appendix C of Intel processor manual volume 3.
7102 // As of 7/26/2016, Intel processor manual Chapter 10 of Volume 3
7103
7104 #include "types.h"
7105 #include "defs.h"
7106 #include "date.h"
7107 #include "memlayout.h"
7108 #include "traps.h"
7109 #include "mmu.h"
7110 #include "x86.h"
7111
7112 // Local APIC registers, divided by 4 for use as uint[] indices.
7113 #define ID      (0x0020/4) // ID
7114 #define VER      (0x0030/4) // Version
7115 #define TPR      (0x0080/4) // Task Priority
7116 #define EOI      (0x00B0/4) // EOI
7117 #define SVR      (0x00F0/4) // Spurious Interrupt Vector
7118 #define ENABLE    0x00000100 // Unit Enable
7119 #define ESR      (0x0280/4) // Error Status
7120 #define ICRLO    (0x0300/4) // Interrupt Command
7121 #define INIT      0x00000500 // INIT/RESET
7122 #define STARTUP  0x00000600 // Startup IPI
7123 #define DELIVS    0x00001000 // Delivery status
7124 #define ASSERT    0x00004000 // Assert interrupt (vs deassert)
7125 #define DEASSERT  0x00000000
7126 #define LEVEL     0x00008000 // Level triggered
7127 #define BCAST     0x00008000 // Send to all APICs, including self.
7128 #define BUSY      0x00001000
7129 #define FIXED      0x00000000
7130 #define ICRHI     (0x0310/4) // Interrupt Command [63:32]
7131 #define TIMER     (0x0320/4) // Local Vector Table 0 (TIMER)
7132 #define X1         0x0000000B // divide counts by 1
7133 #define PERIODIC   0x00020000 // Periodic
7134 #define PCINT     (0x0340/4) // Performance Counter LVT
7135 #define LINT0     (0x0350/4) // Local Vector Table 1 (LINT0)
7136 #define LINT1     (0x0360/4) // Local Vector Table 2 (LINT1)
7137 #define ERROR     (0x0370/4) // Local Vector Table 3 (ERROR)
7138 #define MASKED    0x00010000 // Interrupt masked
7139 #define TICR      (0x0380/4) // Timer Initial Count
7140 #define TCCR      (0x0390/4) // Timer Current Count
7141 #define TDCR      (0x03E0/4) // Timer Divide Configuration
7142
7143 volatile uint *lapic; // Initialized in mp.c
7144
7145 static void
7146 lapicw(int index, int value)
7147 {
7148     lapic[index] = value;
7149     lapic[ID]; // wait for write to finish, by reading

```

```

7150 }
7151
7152 void
7153 lapicinit(void)
7154 {
7155     if(!lapic)
7156         return;
7157
7158     // Enable local APIC; set spurious interrupt vector.
7159     lapicw(SVR, ENABLE | (T_IRQ0 + IRQ_SPURIOUS));
7160
7161     // The timer repeatedly counts down at bus frequency
7162     // from lapic[TICR] and then issues an interrupt.
7163     // If xv6 cared more about precise timekeeping,
7164     // TICR would be calibrated using an external time source.
7165     lapicw(TDCR, X1);
7166     lapicw(TIMER, PERIODIC | (T_IRQ0 + IRQ_TIMER));
7167     lapicw(TICR, 10000000);
7168
7169     // Disable logical interrupt lines.
7170     lapicw(LINT0, MASKED);
7171     lapicw(LINT1, MASKED);
7172
7173     // Disable performance counter overflow interrupts
7174     // on machines that provide that interrupt entry.
7175     if(((lapic[VER]>>16) & 0xFF) >= 4)
7176         lapicw(PCINT, MASKED);
7177
7178     // Map error interrupt to IRQ_ERROR.
7179     lapicw(ERROR, T_IRQ0 + IRQ_ERROR);
7180
7181     // Clear error status register (requires back-to-back writes).
7182     lapicw(ESR, 0);
7183     lapicw(ESR, 0);
7184
7185     // Ack any outstanding interrupts.
7186     lapicw(EOI, 0);
7187
7188     // Send an Init Level De-Assert to synchronise arbitration ID's.
7189     lapicw(ICRHI, 0);
7190     lapicw(ICRLO, BCAST | INIT | LEVEL);
7191     while(lapic[ICRLO] & DELIVS)
7192         ;
7193
7194     // Enable interrupts on the APIC (but not on the processor).
7195     lapicw(TPR, 0);
7196 }
7197
7198
7199

```



```

7200 int
7201 cpunum(void)
7202 {
7203     // Cannot call cpu when interrupts are enabled:
7204     // result not guaranteed to last long enough to be used!
7205     // Would prefer to panic but even printing is chancy here:
7206     // almost everything, including cprintf and panic, calls cpu,
7207     // often indirectly through acquire and release.
7208     if(readeflags() & FL_IF){
7209         static int n;
7210         if(n++ == 0)
7211             cprintf("cpu called from %x with interrupts enabled\n",
7212                 __builtin_return_address(0));
7213     }
7214
7215     if(lapic)
7216         return lapic[ID]>>24;
7217     return 0;
7218 }
7219
7220 // Acknowledge interrupt.
7221 void
7222 lapiceoi(void)
7223 {
7224     if(lapic)
7225         lapicw(EOI, 0);
7226 }
7227
7228 // Spin for a given number of microseconds.
7229 // On real hardware would want to tune this dynamically.
7230 void
7231 microdelay(int us)
7232 {
7233 }
7234
7235 #define CMOS_PORT    0x70
7236 #define CMOS_RETURN  0x71
7237
7238 // Start additional processor running entry code at addr.
7239 // See Appendix B of MultiProcessor Specification.
7240 void
7241 lapicstartap(uchar apicid, uint addr)
7242 {
7243     int i;
7244     ushort *wrv;
7245
7246     // "The BSP must initialize CMOS shutdown code to 0AH
7247     // and the warm reset vector (DWORD based at 40:67) to point at
7248     // the AP startup code prior to the [universal startup algorithm]."
7249     outb(CMOS_PORT, 0xF); // offset 0xF is shutdown code

```

```

7250     outb(CMOS_PORT+1, 0x0A);
7251     wrv = (ushort*)P2V((0x40<<4 | 0x67)); // Warm reset vector
7252     wrv[0] = 0;
7253     wrv[1] = addr >> 4;
7254
7255     // "Universal startup algorithm."
7256     // Send INIT (level-triggered) interrupt to reset other CPU.
7257     lapicw(ICRHI, apicid<<24);
7258     lapicw(ICRLO, INIT | LEVEL | ASSERT);
7259     microdelay(200);
7260     lapicw(ICRLO, INIT | LEVEL);
7261     microdelay(100); // should be 10ms, but too slow in Bochs!
7262
7263     // Send startup IPI (twice!) to enter code.
7264     // Regular hardware is supposed to only accept a STARTUP
7265     // when it is in the halted state due to an INIT. So the second
7266     // should be ignored, but it is part of the official Intel algorithm.
7267     // Bochs complains about the second one. Too bad for Bochs.
7268     for(i = 0; i < 2; i++){
7269         lapicw(ICRHI, apicid<<24);
7270         lapicw(ICRLO, STARTUP | (addr>>12));
7271         microdelay(200);
7272     }
7273 }
7274
7275 #define CMOS_STATA    0x0a
7276 #define CMOS_STATB    0x0b
7277 #define CMOS_UIP      (1 << 7) // RTC update in progress
7278
7279 #define SECS          0x00
7280 #define MINS          0x02
7281 #define HOURS         0x04
7282 #define DAY           0x07
7283 #define MONTH         0x08
7284 #define YEAR          0x09
7285
7286 static uint cmos_read(uint reg)
7287 {
7288     outb(CMOS_PORT, reg);
7289     microdelay(200);
7290
7291     return inb(CMOS_RETURN);
7292 }
7293
7294
7295
7296
7297
7298
7299

```

```

7300 static void fill_rtcddate(struct rtcdate *r)
7301 {
7302     r->second = cmos_read(SECS);
7303     r->minute = cmos_read(MINS);
7304     r->hour   = cmos_read(HOURS);
7305     r->day     = cmos_read(DAY);
7306     r->month   = cmos_read(MONTH);
7307     r->year    = cmos_read(YEAR);
7308 }
7309
7310 // qemu seems to use 24-hour GWT and the values are BCD encoded
7311 void cmostime(struct rtcdate *r)
7312 {
7313     struct rtcdate t1, t2;
7314     int sb, bcd;
7315
7316     sb = cmos_read(CMOS_STATB);
7317
7318     bcd = (sb & (1 << 2)) == 0;
7319
7320     // make sure CMOS doesn't modify time while we read it
7321     for (;;) {
7322         fill_rtcddate(&t1);
7323         if (cmos_read(CMOS_STATA) & CMOS_UIP)
7324             continue;
7325         fill_rtcddate(&t2);
7326         if (memcmp(&t1, &t2, sizeof(t1)) == 0)
7327             break;
7328     }
7329
7330     // convert
7331     if (bcd) {
7332 #define CONV(x)      (t1.x = ((t1.x >> 4) * 10) + (t1.x & 0xf))
7333         CONV(second);
7334         CONV(minute);
7335         CONV(hour);
7336         CONV(day);
7337         CONV(month);
7338         CONV(year);
7339 #undef CONV
7340     }
7341
7342     *r = t1;
7343     r->year += 2000;
7344 }
7345
7346
7347
7348
7349

```

```

7350 // The I/O APIC manages hardware interrupts for an SMP system.
7351 // http://www.intel.com/design/chipsets/datashts/29056601.pdf
7352 // See also picirq.c.
7353
7354 #include "types.h"
7355 #include "defs.h"
7356 #include "traps.h"
7357
7358 #define IOAPIC    0xFEC00000    // Default physical address of IO APIC
7359
7360 #define REG_ID     0x00    // Register index: ID
7361 #define REG_VER    0x01    // Register index: version
7362 #define REG_TABLE  0x10    // Redirection table base
7363
7364 // The redirection table starts at REG_TABLE and uses
7365 // two registers to configure each interrupt.
7366 // The first (low) register in a pair contains configuration bits.
7367 // The second (high) register contains a bitmask telling which
7368 // CPUs can serve that interrupt.
7369 #define INT_DISABLED 0x00010000 // Interrupt disabled
7370 #define INT_LEVEL    0x00008000 // Level-triggered (vs edge-)
7371 #define INT_ACTIVELOW 0x00002000 // Active low (vs high)
7372 #define INT_LOGICAL  0x00000800 // Destination is CPU id (vs APIC ID)
7373
7374 volatile struct ioapic *ioapic;
7375
7376 // IO APIC MMIO structure: write reg, then read or write data.
7377 struct ioapic {
7378     uint reg;
7379     uint pad[3];
7380     uint data;
7381 };
7382
7383 static uint
7384 ioapicread(int reg)
7385 {
7386     ioapic->reg = reg;
7387     return ioapic->data;
7388 }
7389
7390 static void
7391 ioapicwrite(int reg, uint data)
7392 {
7393     ioapic->reg = reg;
7394     ioapic->data = data;
7395 }
7396
7397
7398
7399

```

```

7400 void
7401 ioapicinit(void)
7402 {
7403     int i, id, maxintr;
7404
7405     if(!ismp)
7406         return;
7407
7408     ioapic = (volatile struct ioapic*)IOAPIC;
7409     maxintr = (ioapicread(REG_VER) >> 16) & 0xFF;
7410     id = ioapicread(REG_ID) >> 24;
7411     if(id != ioapicid)
7412         cprintf("ioapicinit: id isn't equal to ioapicid; not a MP\n");
7413
7414     // Mark all interrupts edge-triggered, active high, disabled,
7415     // and not routed to any CPUs.
7416     for(i = 0; i <= maxintr; i++){
7417         ioapicwrite(REG_TABLE+2*i, INT_DISABLED | (T_IRQ0 + i));
7418         ioapicwrite(REG_TABLE+2*i+1, 0);
7419     }
7420 }
7421
7422 void
7423 ioapicenable(int irq, int cpunum)
7424 {
7425     if(!ismp)
7426         return;
7427
7428     // Mark interrupt edge-triggered, active high,
7429     // enabled, and routed to the given cpunum,
7430     // which happens to be that cpu's APIC ID.
7431     ioapicwrite(REG_TABLE+2*irq, T_IRQ0 + irq);
7432     ioapicwrite(REG_TABLE+2*irq+1, cpunum << 24);
7433 }
7434
7435
7436
7437
7438
7439
7440
7441
7442
7443
7444
7445
7446
7447
7448
7449

```

```

7450 // Intel 8259A programmable interrupt controllers.
7451
7452 #include "types.h"
7453 #include "x86.h"
7454 #include "traps.h"
7455
7456 // I/O Addresses of the two programmable interrupt controllers
7457 #define IO_PIC1      0x20    // Master (IRQs 0-7)
7458 #define IO_PIC2      0xA0    // Slave (IRQs 8-15)
7459
7460 #define IRQ_SLAVE     2      // IRQ at which slave connects to master
7461
7462 // Current IRQ mask.
7463 // Initial IRQ mask has interrupt 2 enabled (for slave 8259A).
7464 static ushort irqmask = 0xFFFF & ~(1<<IRQ_SLAVE);
7465
7466 static void
7467 picsetmask(ushort mask)
7468 {
7469     irqmask = mask;
7470     outb(IO_PIC1+1, mask);
7471     outb(IO_PIC2+1, mask >> 8);
7472 }
7473
7474 void
7475 picenable(int irq)
7476 {
7477     picsetmask(irqmask & ~(1<<irq));
7478 }
7479
7480 // Initialize the 8259A interrupt controllers.
7481 void
7482 picinit(void)
7483 {
7484     // mask all interrupts
7485     outb(IO_PIC1+1, 0xFF);
7486     outb(IO_PIC2+1, 0xFF);
7487
7488     // Set up master (8259A-1)
7489
7490     // ICW1: 0001g0hi
7491     //   g: 0 = edge triggering, 1 = level triggering
7492     //   h: 0 = cascaded PICs, 1 = master only
7493     //   i: 0 = no ICW4, 1 = ICW4 required
7494     outb(IO_PIC1, 0x11);
7495
7496     // ICW2: Vector offset
7497     outb(IO_PIC1+1, T_IRQ0);
7498
7499

```

```

7500 // ICW3: (master PIC) bit mask of IR lines connected to slaves
7501 //      (slave PIC) 3-bit # of slave's connection to master
7502 outb(IO_PIC1+1, 1<<IRQ_SLAVE);
7503
7504 // ICW4: 000nbmap
7505 //      n: 1 = special fully nested mode
7506 //      b: 1 = buffered mode
7507 //      m: 0 = slave PIC, 1 = master PIC
7508 //      (ignored when b is 0, as the master/slave role
7509 //      can be hardwired).
7510 //      a: 1 = Automatic EOI mode
7511 //      p: 0 = MCS-80/85 mode, 1 = intel x86 mode
7512 outb(IO_PIC1+1, 0x3);
7513
7514 // Set up slave (8259A-2)
7515 outb(IO_PIC2, 0x11); // ICW1
7516 outb(IO_PIC2+1, T_IRQ0 + 8); // ICW2
7517 outb(IO_PIC2+1, IRQ_SLAVE); // ICW3
7518 // NB Automatic EOI mode doesn't tend to work on the slave.
7519 // Linux source code says it's "to be investigated".
7520 outb(IO_PIC2+1, 0x3); // ICW4
7521
7522 // OCW3: 0ef01prs
7523 //      ef: 0x = NOP, 10 = clear specific mask, 11 = set specific mask
7524 //      p: 0 = no polling, 1 = polling mode
7525 //      rs: 0x = NOP, 10 = read IRR, 11 = read ISR
7526 outb(IO_PIC1, 0x68); // clear specific mask
7527 outb(IO_PIC1, 0x0a); // read IRR by default
7528
7529 outb(IO_PIC2, 0x68); // OCW3
7530 outb(IO_PIC2, 0x0a); // OCW3
7531
7532 if(irqmask != 0xFFFF)
7533     picsetmask(irqmask);
7534 }
7535
7536
7537
7538
7539
7540
7541
7542
7543
7544
7545
7546
7547
7548
7549

```

```

7550 // PC keyboard interface constants
7551
7552 #define KBSTATP      0x64    // kbd controller status port(I)
7553 #define KBS_DIB      0x01    // kbd data in buffer
7554 #define KBDATAP      0x60    // kbd data port(I)
7555
7556 #define NO            0
7557
7558 #define SHIFT         (1<<0)
7559 #define CTL           (1<<1)
7560 #define ALT           (1<<2)
7561
7562 #define CAPSLOCK      (1<<3)
7563 #define NUMLOCK       (1<<4)
7564 #define SCROLLLOCK    (1<<5)
7565
7566 #define E0ESC         (1<<6)
7567
7568 // Special keycodes
7569 #define KEY_HOME      0xE0
7570 #define KEY_END       0xE1
7571 #define KEY_UP        0xE2
7572 #define KEY_DN        0xE3
7573 #define KEY_LF        0xE4
7574 #define KEY_RT        0xE5
7575 #define KEY_PGUP      0xE6
7576 #define KEY_PGDN      0xE7
7577 #define KEY_INS       0xE8
7578 #define KEY_DEL       0xE9
7579
7580 // C('A') == Control-A
7581 #define C(x) (x - '@')
7582
7583 static uchar shiftcode[256] =
7584 {
7585     [0x1D] CTL,
7586     [0x2A] SHIFT,
7587     [0x36] SHIFT,
7588     [0x38] ALT,
7589     [0x9D] CTL,
7590     [0xB8] ALT
7591 };
7592
7593 static uchar togglecode[256] =
7594 {
7595     [0x3A] CAPSLOCK,
7596     [0x45] NUMLOCK,
7597     [0x46] SCROLLLOCK
7598 };
7599

```

```

7600 static uchar normalmap[256] =
7601 {
7602     NO,    0x1B, '1', '2', '3', '4', '5', '6', // 0x00
7603     '7', '8', '9', '0', '-', '=', '\b', '\t',
7604     'q', 'w', 'e', 'r', 't', 'y', 'u', 'i', // 0x10
7605     'o', 'p', '[', ']', '\n', NO, 'a', 's',
7606     'd', 'f', 'g', 'h', 'j', 'k', 'l', ';' // 0x20
7607     '\'', '\'', NO, '\\', 'z', 'x', 'c', 'v',
7608     'b', 'n', 'm', ',', '.', '/', NO, '*', // 0x30
7609     NO, ' ', NO, NO, NO, NO, NO, NO,
7610     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
7611     '8', '9', '-', '4', '5', '6', '+', '1',
7612     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
7613     [0x9C] '\n', // KP_Enter
7614     [0xB5] '/', // KP_Div
7615     [0xC8] KEY_UP, [0xD0] KEY_DN,
7616     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
7617     [0xCB] KEY_LF, [0xCD] KEY_RT,
7618     [0x97] KEY_HOME, [0xCF] KEY_END,
7619     [0xD2] KEY_INS, [0xD3] KEY_DEL
7620 };
7621
7622 static uchar shiftmap[256] =
7623 {
7624     NO,    033, '!', '@', '#', '$', '%', '^', // 0x00
7625     '&', '*', '(', ')', '-', '+', '\b', '\t',
7626     'Q', 'W', 'E', 'R', 'T', 'Y', 'U', 'I', // 0x10
7627     'O', 'P', '{', '}', '\n', NO, 'A', 'S',
7628     'D', 'F', 'G', 'H', 'J', 'K', 'L', ':' // 0x20
7629     '"', '~', NO, '|', 'Z', 'X', 'C', 'V',
7630     'B', 'N', 'M', '<', '>', '?', NO, '*', // 0x30
7631     NO, ' ', NO, NO, NO, NO, NO, NO,
7632     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
7633     '8', '9', '-', '4', '5', '6', '+', '1',
7634     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
7635     [0x9C] '\n', // KP_Enter
7636     [0xB5] '/', // KP_Div
7637     [0xC8] KEY_UP, [0xD0] KEY_DN,
7638     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
7639     [0xCB] KEY_LF, [0xCD] KEY_RT,
7640     [0x97] KEY_HOME, [0xCF] KEY_END,
7641     [0xD2] KEY_INS, [0xD3] KEY_DEL
7642 };
7643
7644
7645
7646
7647
7648
7649

```

```

7650 static uchar ctlmap[256] =
7651 {
7652     NO,    NO,    NO,    NO,    NO,    NO,    NO,    NO,
7653     NO,    NO,    NO,    NO,    NO,    NO,    NO,    NO,
7654     C('Q'), C('W'), C('E'), C('R'), C('T'), C('Y'), C('U'), C('I'),
7655     C('O'), C('P'), NO,    NO,    '\r', NO,    C('A'), C('S'),
7656     C('D'), C('F'), C('G'), C('H'), C('J'), C('K'), C('L'), NO,
7657     NO,    NO,    NO,    C('\'), C('Z'), C('X'), C('C'), C('V'),
7658     C('B'), C('N'), C('M'), NO,    NO,    C('/'), NO,    NO,
7659     [0x9C] '\r', // KP_Enter
7660     [0xB5] C('/'), // KP_Div
7661     [0xC8] KEY_UP, [0xD0] KEY_DN,
7662     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
7663     [0xCB] KEY_LF, [0xCD] KEY_RT,
7664     [0x97] KEY_HOME, [0xCF] KEY_END,
7665     [0xD2] KEY_INS, [0xD3] KEY_DEL
7666 };
7667
7668
7669
7670
7671
7672
7673
7674
7675
7676
7677
7678
7679
7680
7681
7682
7683
7684
7685
7686
7687
7688
7689
7690
7691
7692
7693
7694
7695
7696
7697
7698
7699

```

```

7700 #include "types.h"
7701 #include "x86.h"
7702 #include "defs.h"
7703 #include "kbd.h"
7704
7705 int
7706 kbdgetc(void)
7707 {
7708     static uint shift;
7709     static uchar *charcode[4] = {
7710         normalmap, shiftmap, ctlmap, ctlmap
7711     };
7712     uint st, data, c;
7713
7714     st = inb(KBSTATP);
7715     if((st & KBS_DIB) == 0)
7716         return -1;
7717     data = inb(KBDATAP);
7718
7719     if(data == 0xE0){
7720         shift |= E0ESC;
7721         return 0;
7722     } else if(data & 0x80){
7723         // Key released
7724         data = (shift & E0ESC ? data : data & 0x7F);
7725         shift &= ~(shiftcode[data] | E0ESC);
7726         return 0;
7727     } else if(shift & E0ESC){
7728         // Last character was an E0 escape; or with 0x80
7729         data |= 0x80;
7730         shift &= ~E0ESC;
7731     }
7732
7733     shift |= shiftcode[data];
7734     shift ^= togglecode[data];
7735     c = charcode[shift & (CTL | SHIFT)][data];
7736     if(shift & CAPSLOCK){
7737         if('a' <= c && c <= 'z')
7738             c += 'A' - 'a';
7739         else if('A' <= c && c <= 'Z')
7740             c += 'a' - 'A';
7741     }
7742     return c;
7743 }
7744
7745 void
7746 kbdintr(void)
7747 {
7748     consoleintr(kbdgetc);
7749 }

```

```

7750 // Console input and output.
7751 // Input is from the keyboard or serial port.
7752 // Output is written to the screen and serial port.
7753
7754 #include "types.h"
7755 #include "defs.h"
7756 #include "param.h"
7757 #include "traps.h"
7758 #include "spinlock.h"
7759 #include "fs.h"
7760 #include "file.h"
7761 #include "memlayout.h"
7762 #include "mmu.h"
7763 #include "proc.h"
7764 #include "x86.h"
7765
7766 static void consputc(int);
7767
7768 static int panicked = 0;
7769
7770 static struct {
7771     struct spinlock lock;
7772     int locking;
7773 } cons;
7774
7775 static void
7776 printint(int xx, int base, int sign)
7777 {
7778     static char digits[] = "0123456789abcdef";
7779     char buf[16];
7780     int i;
7781     uint x;
7782
7783     if(sign && (sign = xx < 0))
7784         x = -xx;
7785     else
7786         x = xx;
7787
7788     i = 0;
7789     do{
7790         buf[i++] = digits[x % base];
7791     }while((x /= base) != 0);
7792
7793     if(sign)
7794         buf[i++] = '-';
7795
7796     while(--i >= 0)
7797         consputc(buf[i]);
7798 }
7799

```

```

7800 // Print to the console. only understands %d, %x, %p, %s.
7801 void
7802 cprintf(char *fmt, ...)
7803 {
7804     int i, c, locking;
7805     uint *argp;
7806     char *s;
7807
7808     locking = cons.locking;
7809     if(locking)
7810         acquire(&cons.lock);
7811
7812     if (fmt == 0)
7813         panic("null fmt");
7814
7815     argp = (uint*)(void*)&fmt + 1;
7816     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
7817         if(c != '%'){
7818             consputc(c);
7819             continue;
7820         }
7821         c = fmt[++i] & 0xff;
7822         if(c == 0)
7823             break;
7824         switch(c){
7825             case 'd':
7826                 printint(*argp++, 10, 1);
7827                 break;
7828             case 'x':
7829             case 'p':
7830                 printint(*argp++, 16, 0);
7831                 break;
7832             case 's':
7833                 if((s = (char*)*argp++) == 0)
7834                     s = "(null)";
7835                 for(; *s; s++)
7836                     consputc(*s);
7837                 break;
7838             case '%':
7839                 consputc('%');
7840                 break;
7841             default:
7842                 // Print unknown % sequence to draw attention.
7843                 consputc('%');
7844                 consputc(c);
7845                 break;
7846         }
7847     }
7848 }
7849

```

```

7850     if(locking)
7851         release(&cons.lock);
7852 }
7853
7854 void
7855 panic(char *s)
7856 {
7857     int i;
7858     uint pcs[10];
7859
7860     cli();
7861     cons.locking = 0;
7862     cprintf("cpu%d: panic: ", cpu->id);
7863     cprintf(s);
7864     cprintf("\n");
7865     getcallerpcs(&s, pcs);
7866     for(i=0; i<10; i++)
7867         cprintf(" %p", pcs[i]);
7868     panicked = 1; // freeze other CPU
7869     for(;;)
7870         ;
7871 }
7872
7873 #define BACKSPACE 0x100
7874 #define CRTPORT 0x3d4
7875 static ushort *crt = (ushort*)P2V(0xb8000); // CGA memory
7876
7877 static void
7878 cgaputc(int c)
7879 {
7880     int pos;
7881
7882     // Cursor position: col + 80*row.
7883     outb(CRTPORT, 14);
7884     pos = inb(CRTPORT+1) << 8;
7885     outb(CRTPORT, 15);
7886     pos |= inb(CRTPORT+1);
7887
7888     if(c == '\n')
7889         pos += 80 - pos%80;
7890     else if(c == BACKSPACE){
7891         if(pos > 0) --pos;
7892     } else
7893         crt[pos++] = (c&0xff) | 0x0700; // black on white
7894
7895     if(pos < 0 || pos > 25*80)
7896         panic("pos under/overflow");
7897 }
7898
7899

```

```

7900 if((pos/80) >= 24){ // Scroll up.
7901     memmove(crt, crt+80, sizeof(crt[0])*23*80);
7902     pos -= 80;
7903     memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
7904 }
7905
7906 outb(CRTPORT, 14);
7907 outb(CRTPORT+1, pos>>8);
7908 outb(CRTPORT, 15);
7909 outb(CRTPORT+1, pos);
7910 crt[pos] = ' ' | 0x0700;
7911 }
7912
7913 void
7914 consputc(int c)
7915 {
7916     if(panicked){
7917         cli();
7918         for(;;)
7919             ;
7920     }
7921
7922     if(c == BACKSPACE){
7923         uartputc('\b'); uartputc(' '); uartputc('\b');
7924     } else
7925         uartputc(c);
7926     cgaputc(c);
7927 }
7928
7929 #define INPUT_BUF 128
7930 struct {
7931     char buf[INPUT_BUF];
7932     uint r; // Read index
7933     uint w; // Write index
7934     uint e; // Edit index
7935 } input;
7936
7937 #define C(x) ((x)-'@') // Control-x
7938
7939 void
7940 consoleintr(int (*getc)(void))
7941 {
7942     int c, doprocdump = 0;
7943
7944     acquire(&cons.lock);
7945     while((c = getc()) >= 0){
7946         switch(c){
7947             case C('P'): // Process listing.
7948                 doprocdump = 1; // procdump() locks cons.lock indirectly; invoke lat
7949                 break;

```

```

7950     case C('U'): // Kill line.
7951         while(input.e != input.w &&
7952             input.buf[(input.e-1) % INPUT_BUF] != '\n'){
7953             input.e--;
7954             consputc(BACKSPACE);
7955         }
7956         break;
7957     case C('H'): case '\x7f': // Backspace
7958         if(input.e != input.w){
7959             input.e--;
7960             consputc(BACKSPACE);
7961         }
7962         break;
7963     default:
7964         if(c != 0 && input.e-input.r < INPUT_BUF){
7965             c = (c == '\r') ? '\n' : c;
7966             input.buf[input.e++ % INPUT_BUF] = c;
7967             consputc(c);
7968             if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
7969                 input.w = input.e;
7970                 wakeup(&input.r);
7971             }
7972         }
7973         break;
7974     }
7975 }
7976 release(&cons.lock);
7977 if(doprocdump) {
7978     procdump(); // now call procdump() wo. cons.lock held
7979 }
7980 }
7981
7982 int
7983 consoleread(struct inode *ip, char *dst, int n)
7984 {
7985     uint target;
7986     int c;
7987
7988     iunlock(ip);
7989     target = n;
7990     acquire(&cons.lock);
7991     while(n > 0){
7992         while(input.r == input.w){
7993             if(proc->killed){
7994                 release(&cons.lock);
7995                 ilock(ip);
7996                 return -1;
7997             }
7998             sleep(&input.r, &cons.lock);
7999         }

```



```

8000     c = input.buf[input.r++ % INPUT_BUF];
8001     if(c == C('D')){ // EOF
8002         if(n < target){
8003             // Save ^D for next time, to make sure
8004             // caller gets a 0-byte result.
8005             input.r--;
8006         }
8007         break;
8008     }
8009     *dst++ = c;
8010     --n;
8011     if(c == '\n')
8012         break;
8013 }
8014 release(&cons.lock);
8015 ilock(ip);
8016
8017 return target - n;
8018 }
8019
8020 int
8021 consolewrite(struct inode *ip, char *buf, int n)
8022 {
8023     int i;
8024
8025     iunlock(ip);
8026     acquire(&cons.lock);
8027     for(i = 0; i < n; i++){
8028         consputc(buf[i] & 0xff);
8029     }
8030     release(&cons.lock);
8031     ilock(ip);
8032
8033     return n;
8034 }
8035
8036 void
8037 consoleinit(void)
8038 {
8039     initlock(&cons.lock, "console");
8040
8041     devsw[CONSOLE].write = consolewrite;
8042     devsw[CONSOLE].read = consoleread;
8043     cons.locking = 1;
8044
8045     picenable IRQ_KBD;
8046     ioapicenable(IRQ_KBD, 0);
8047 }
8048
8049

```

```

8050 // Intel 8253/8254/82C54 Programmable Interval Timer (PIT).
8051 // Only used on uniprocessors;
8052 // SMP machines use the local APIC timer.
8053
8054 #include "types.h"
8055 #include "defs.h"
8056 #include "traps.h"
8057 #include "x86.h"
8058
8059 #define IO_TIMER1      0x040          // 8253 Timer #1
8060
8061 // Frequency of all three count-down timers;
8062 // (TIMER_FREQ/freq) is the appropriate count
8063 // to generate a frequency of freq Hz.
8064
8065 #define TIMER_FREQ      1193182
8066 #define TIMER_DIV(x)    ((TIMER_FREQ+(x)/2)/(x))
8067
8068 #define TIMER_MODE      (IO_TIMER1 + 3) // timer mode port
8069 #define TIMER_SELO      0x00          // select counter 0
8070 #define TIMER_RATEGEN    0x04          // mode 2, rate generator
8071 #define TIMER_16BIT      0x30          // r/w counter 16 bits, LSB first
8072
8073 void
8074 timerinit(void)
8075 {
8076     // Interrupt 100 times/sec.
8077     outb(TIMER_MODE, TIMER_SELO | TIMER_RATEGEN | TIMER_16BIT);
8078     outb(IO_TIMER1, TIMER_DIV(100) % 256);
8079     outb(IO_TIMER1, TIMER_DIV(100) / 256);
8080     picenable(IRQ_TIMER);
8081 }
8082
8083
8084
8085
8086
8087
8088
8089
8090
8091
8092
8093
8094
8095
8096
8097
8098
8099

```

```

8100 // Intel 8250 serial port (UART).
8101
8102 #include "types.h"
8103 #include "defs.h"
8104 #include "param.h"
8105 #include "traps.h"
8106 #include "spinlock.h"
8107 #include "fs.h"
8108 #include "file.h"
8109 #include "mmu.h"
8110 #include "proc.h"
8111 #include "x86.h"
8112
8113 #define COM1      0x3f8
8114
8115 static int uart;    // is there a uart?
8116
8117 void
8118 uartinit(void)
8119 {
8120     char *p;
8121
8122     // Turn off the FIFO
8123     outb(COM1+2, 0);
8124
8125     // 9600 baud, 8 data bits, 1 stop bit, parity off.
8126     outb(COM1+3, 0x80);    // Unlock divisor
8127     outb(COM1+0, 115200/9600);
8128     outb(COM1+1, 0);
8129     outb(COM1+3, 0x03);    // Lock divisor, 8 data bits.
8130     outb(COM1+4, 0);
8131     outb(COM1+1, 0x01);    // Enable receive interrupts.
8132
8133     // If status is 0xFF, no serial port.
8134     if(inb(COM1+5) == 0xFF)
8135         return;
8136     uart = 1;
8137
8138     // Acknowledge pre-existing interrupt conditions;
8139     // enable interrupts.
8140     inb(COM1+2);
8141     inb(COM1+0);
8142     picenable(IRQ_COM1);
8143     ioapicenable(IRQ_COM1, 0);
8144
8145     // Announce that we're here.
8146     for(p="xv6...\n"; *p; p++)
8147         uartputc(*p);
8148 }
8149

```

```

8150 void
8151 uartputc(int c)
8152 {
8153     int i;
8154
8155     if(!uart)
8156         return;
8157     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)
8158         microdelay(10);
8159     outb(COM1+0, c);
8160 }
8161
8162 static int
8163 uartgetc(void)
8164 {
8165     if(!uart)
8166         return -1;
8167     if(!(inb(COM1+5) & 0x01))
8168         return -1;
8169     return inb(COM1+0);
8170 }
8171
8172 void
8173 uartintr(void)
8174 {
8175     consoleintr(uartgetc);
8176 }
8177
8178
8179
8180
8181
8182
8183
8184
8185
8186
8187
8188
8189
8190
8191
8192
8193
8194
8195
8196
8197
8198
8199

```

```
8200 # Initial process execs /init.
8201
8202 #include "syscall.h"
8203 #include "traps.h"
8204
8205
8206 # exec(init, argv)
8207 .globl start
8208 start:
8209     pushl $argv
8210     pushl $init
8211     pushl $0 // where caller pc would be
8212     movl $SYS_exec, %eax
8213     int $T_SYSCALL
8214
8215 # for(;;) exit();
8216 exit:
8217     movl $SYS_exit, %eax
8218     int $T_SYSCALL
8219     jmp exit
8220
8221 # char init[] = "/init\0";
8222 init:
8223     .string "/init\0"
8224
8225 # char *argv[] = { init, 0 };
8226 .p2align 2
8227 argv:
8228     .long init
8229     .long 0
8230
8231
8232
8233
8234
8235
8236
8237
8238
8239
8240
8241
8242
8243
8244
8245
8246
8247
8248
8249
```

```
8250 #include "syscall.h"
8251 #include "traps.h"
8252
8253 #define SYSCALL(name) \
8254     .globl name; \
8255     name: \
8256     movl $SYS_ ## name, %eax; \
8257     int $T_SYSCALL; \
8258     ret
8259
8260 SYSCALL(fork)
8261 SYSCALL(exit)
8262 SYSCALL(wait)
8263 SYSCALL(pipe)
8264 SYSCALL(read)
8265 SYSCALL(write)
8266 SYSCALL(close)
8267 SYSCALL(kill)
8268 SYSCALL(exec)
8269 SYSCALL(open)
8270 SYSCALL(mknod)
8271 SYSCALL(unlink)
8272 SYSCALL(fstat)
8273 SYSCALL(link)
8274 SYSCALL(mkdir)
8275 SYSCALL(chdir)
8276 SYSCALL(dup)
8277 SYSCALL(getpid)
8278 SYSCALL(sbrk)
8279 SYSCALL(sleep)
8280 SYSCALL(uptime)
8281 SYSCALL(halt)
8282 SYSCALL(date)
8283 SYSCALL(getuid)
8284 SYSCALL(getgid)
8285 SYSCALL(getppid)
8286 SYSCALL(setuid)
8287 SYSCALL(setgid)
8288 SYSCALL(getprocs)
8289 SYSCALL(setpriority)
8290
8291
8292
8293
8294
8295
8296
8297
8298
8299
```

```

8300 // init: The initial user-level program
8301
8302 #include "types.h"
8303 #include "stat.h"
8304 #include "user.h"
8305 #include "fcntl.h"
8306
8307 char *argv[] = { "sh", 0 };
8308
8309 int
8310 main(void)
8311 {
8312     int pid, wpid;
8313
8314     if(open("console", O_RDWR) < 0){
8315         mknod("console", 1, 1);
8316         open("console", O_RDWR);
8317     }
8318     dup(0); // stdout
8319     dup(0); // stderr
8320
8321     for(;;){
8322         printf(1, "init: starting sh\n");
8323         pid = fork();
8324         if(pid < 0){
8325             printf(1, "init: fork failed\n");
8326             exit();
8327         }
8328         if(pid == 0){
8329             exec("sh", argv);
8330             printf(1, "init: exec sh failed\n");
8331             exit();
8332         }
8333         while((wpid=wait()) >= 0 && wpid != pid)
8334             printf(1, "zombie!\n");
8335     }
8336 }
8337
8338
8339
8340
8341
8342
8343
8344
8345
8346
8347
8348
8349

```

```

8350 // Shell.
8351 // 2015-12-21. Added very simple processing for builtin commands
8352
8353 #include "types.h"
8354 #include "user.h"
8355 #include "fcntl.h"
8356
8357 // Parsed command representation
8358 #define EXEC 1
8359 #define REDIR 2
8360 #define PIPE 3
8361 #define LIST 4
8362 #define BACK 5
8363
8364 #define MAXARGS 10
8365
8366 struct cmd {
8367     int type;
8368 };
8369
8370 struct execcmd {
8371     int type;
8372     char *argv[MAXARGS];
8373     char *eargv[MAXARGS];
8374 };
8375
8376 struct redircmd {
8377     int type;
8378     struct cmd *cmd;
8379     char *file;
8380     char *efile;
8381     int mode;
8382     int fd;
8383 };
8384
8385 struct pipecmd {
8386     int type;
8387     struct cmd *left;
8388     struct cmd *right;
8389 };
8390
8391 struct listcmd {
8392     int type;
8393     struct cmd *left;
8394     struct cmd *right;
8395 };
8396
8397
8398
8399

```

```

8400 struct backcmd {
8401     int type;
8402     struct cmd *cmd;
8403 };
8404
8405 int fork1(void); // Fork but panics on failure.
8406 void panic(char*);
8407 struct cmd *parsecmd(char*);
8408
8409 // Execute cmd. Never returns.
8410 void
8411 runcmd(struct cmd *cmd)
8412 {
8413     int p[2];
8414     struct backcmd *bcmd;
8415     struct execcmd *ecmd;
8416     struct listcmd *lcmd;
8417     struct pipecmd *pcmd;
8418     struct redircmd *rcmd;
8419
8420     if(cmd == 0)
8421         exit();
8422
8423     switch(cmd->type){
8424     default:
8425         panic("runcmd");
8426
8427     case EXEC:
8428         ecmd = (struct execcmd*)cmd;
8429         if(ecmd->argv[0] == 0)
8430             exit();
8431         exec(ecmd->argv[0], ecmd->argv);
8432         printf(2, "exec %s failed\n", ecmd->argv[0]);
8433         break;
8434
8435     case REDIR:
8436         rcmd = (struct redircmd*)cmd;
8437         close(rcmd->fd);
8438         if(open(rcmd->file, rcmd->mode) < 0){
8439             printf(2, "open %s failed\n", rcmd->file);
8440             exit();
8441         }
8442         runcmd(rcmd->cmd);
8443         break;
8444
8445     case LIST:
8446         lcmd = (struct listcmd*)cmd;
8447         if(fork1() == 0)
8448             runcmd(lcmd->left);
8449         wait();

```

```

8450     runcmd(lcmd->right);
8451     break;
8452
8453     case PIPE:
8454         pcmd = (struct pipecmd*)cmd;
8455         if(pipe(p) < 0)
8456             panic("pipe");
8457         if(fork1() == 0){
8458             close(1);
8459             dup(p[1]);
8460             close(p[0]);
8461             close(p[1]);
8462             runcmd(pcmd->left);
8463         }
8464         if(fork1() == 0){
8465             close(0);
8466             dup(p[0]);
8467             close(p[0]);
8468             close(p[1]);
8469             runcmd(pcmd->right);
8470         }
8471         close(p[0]);
8472         close(p[1]);
8473         wait();
8474         wait();
8475         break;
8476
8477     case BACK:
8478         bcmd = (struct backcmd*)cmd;
8479         if(fork1() == 0)
8480             runcmd(bcmd->cmd);
8481         break;
8482     }
8483     exit();
8484 }
8485
8486 int
8487 getcmd(char *buf, int nbuf)
8488 {
8489     printf(2, "$ ");
8490     memset(buf, 0, nbuf);
8491     gets(buf, nbuf);
8492     if(buf[0] == 0) // EOF
8493         return -1;
8494     return 0;
8495 }
8496
8497
8498
8499

```

```

8500 #ifdef USE_BUILTINS
8501 // ***** processing for shell builtins begins here *****
8502
8503 int
8504 strncmp(const char *p, const char *q, uint n)
8505 {
8506     while(n > 0 && *p && *p == *q)
8507         n--, p++, q++;
8508     if(n == 0)
8509         return 0;
8510     return (uchar)*p - (uchar)*q;
8511 }
8512
8513 int
8514 makeint(char *p)
8515 {
8516     int val = 0;
8517
8518     while ((*p >= '0') && (*p <= '9')) {
8519         val = 10*val + (*p-'0');
8520         ++p;
8521     }
8522     return val;
8523 }
8524
8525 int
8526 setbuiltin(char *p)
8527 {
8528     int i;
8529
8530     p += strlen("_set");
8531     while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
8532     if (strncmp("uid", p, 3) == 0) {
8533         p += strlen("uid");
8534         while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
8535         i = makeint(p); // ugly
8536         return (setuid(i));
8537     } else
8538     if (strncmp("gid", p, 3) == 0) {
8539         p += strlen("gid");
8540         while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
8541         i = makeint(p); // ugly
8542         return (setgid(i));
8543     }
8544     printf(2, "Invalid _set parameter\n");
8545     return -1;
8546 }
8547
8548
8549

```

```

8550 int
8551 getbuiltin(char *p)
8552 {
8553     p += strlen("_get");
8554     while (strncmp(p, " ", 1) == 0) p++; // chomp spaces
8555     if (strncmp("uid", p, 3) == 0) {
8556         printf(2, "%d\n", getuid());
8557         return 0;
8558     }
8559     if (strncmp("gid", p, 3) == 0) {
8560         printf(2, "%d\n", getgid());
8561         return 0;
8562     }
8563     printf(2, "Invalid _get parameter\n");
8564     return -1;
8565 }
8566
8567 typedef int funcPtr_t(char *);
8568 typedef struct {
8569     char *cmd;
8570     funcPtr_t *name;
8571 } dispatchTableEntry_t;
8572
8573 // Use a simple function dispatch table (FDT) to process builtin commands
8574 dispatchTableEntry_t fdt[] = {
8575     {"_set", setbuiltin},
8576     {"_get", getbuiltin}
8577 };
8578 int FDTcount = sizeof(fdt) / sizeof(fdt[0]); // # entris in FDT
8579
8580 void
8581 dobuiltin(char *cmd) {
8582     int i;
8583
8584     for (i=0; i<FDTcount; i++)
8585         if (strncmp(cmd, fdt[i].cmd, strlen(fdt[i].cmd)) == 0)
8586             (*fdt[i].name)(cmd);
8587 }
8588
8589
8590
8591
8592
8593
8594
8595
8596
8597
8598
8599

```

```

8600 // ***** processing for shell builtins ends here *****
8601 #endif
8602
8603 int
8604 main(void)
8605 {
8606     static char buf[100];
8607     int fd;
8608
8609     // Assumes three file descriptors open.
8610     while((fd = open("console", O_RDWR)) >= 0){
8611         if(fd >= 3){
8612             close(fd);
8613             break;
8614         }
8615     }
8616
8617     // Read and run input commands.
8618     while(getcmd(buf, sizeof(buf)) >= 0){
8619         // add support for built-ins here. cd is a built-in
8620         if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
8621             // Clumsy but will have to do for now.
8622             // Chdir has no effect on the parent if run in the child.
8623             buf[strlen(buf)-1] = 0; // chop \n
8624             if(chdir(buf+3) < 0)
8625                 printf(2, "cannot cd %s\n", buf+3);
8626             continue;
8627         }
8628 #ifdef USE_BUILTINS
8629         if (buf[0]=='_' ) { // assume it is a builtin command
8630             dobuiltin(buf);
8631             continue;
8632         }
8633 #endif
8634         if(fork1() == 0)
8635             runcmd(parsecmd(buf));
8636         wait();
8637     }
8638     exit();
8639 }
8640
8641 void
8642 panic(char *s)
8643 {
8644     printf(2, "%s\n", s);
8645     exit();
8646 }
8647
8648
8649

```

```

8650 int
8651 fork1(void)
8652 {
8653     int pid;
8654
8655     pid = fork();
8656     if(pid == -1)
8657         panic("fork");
8658     return pid;
8659 }
8660
8661 // Constructors
8662
8663 struct cmd*
8664 execcmd(void)
8665 {
8666     struct execcmd *cmd;
8667
8668     cmd = malloc(sizeof(*cmd));
8669     memset(cmd, 0, sizeof(*cmd));
8670     cmd->type = EXEC;
8671     return (struct cmd*)cmd;
8672 }
8673
8674 struct cmd*
8675 redircmd(struct cmd *subcmd, char *file, char *efile, int mode, int fd)
8676 {
8677     struct redircmd *cmd;
8678
8679     cmd = malloc(sizeof(*cmd));
8680     memset(cmd, 0, sizeof(*cmd));
8681     cmd->type = REDIR;
8682     cmd->cmd = subcmd;
8683     cmd->file = file;
8684     cmd->efile = efile;
8685     cmd->mode = mode;
8686     cmd->fd = fd;
8687     return (struct cmd*)cmd;
8688 }
8689
8690
8691
8692
8693
8694
8695
8696
8697
8698
8699

```

```

8700 struct cmd*
8701 pipecmd(struct cmd *left, struct cmd *right)
8702 {
8703     struct pipecmd *cmd;
8704
8705     cmd = malloc(sizeof(*cmd));
8706     memset(cmd, 0, sizeof(*cmd));
8707     cmd->type = PIPE;
8708     cmd->left = left;
8709     cmd->right = right;
8710     return (struct cmd*)cmd;
8711 }
8712
8713 struct cmd*
8714 listcmd(struct cmd *left, struct cmd *right)
8715 {
8716     struct listcmd *cmd;
8717
8718     cmd = malloc(sizeof(*cmd));
8719     memset(cmd, 0, sizeof(*cmd));
8720     cmd->type = LIST;
8721     cmd->left = left;
8722     cmd->right = right;
8723     return (struct cmd*)cmd;
8724 }
8725
8726 struct cmd*
8727 backcmd(struct cmd *subcmd)
8728 {
8729     struct backcmd *cmd;
8730
8731     cmd = malloc(sizeof(*cmd));
8732     memset(cmd, 0, sizeof(*cmd));
8733     cmd->type = BACK;
8734     cmd->cmd = subcmd;
8735     return (struct cmd*)cmd;
8736 }
8737
8738
8739
8740
8741
8742
8743
8744
8745
8746
8747
8748
8749

```

```

8750 // Parsing
8751
8752 char whitespace[] = " \t\r\n\v";
8753 char symbols[] = "<|>&;()";
8754
8755 int
8756 gettoken(char **ps, char *es, char **q, char **eq)
8757 {
8758     char *s;
8759     int ret;
8760
8761     s = *ps;
8762     while(s < es && strchr(whitespace, *s))
8763         s++;
8764     if(q)
8765         *q = s;
8766     ret = *s;
8767     switch(*s){
8768     case 0:
8769         break;
8770     case '|':
8771     case '(':
8772     case ')':
8773     case ';':
8774     case '&':
8775     case '<':
8776         s++;
8777         break;
8778     case '>':
8779         s++;
8780         if(*s == '>'){
8781             ret = '+';
8782             s++;
8783         }
8784         break;
8785     default:
8786         ret = 'a';
8787         while(s < es && !strchr(whitespace, *s) && !strchr(symbols, *s))
8788             s++;
8789         break;
8790     }
8791     if(eq)
8792         *eq = s;
8793
8794     while(s < es && strchr(whitespace, *s))
8795         s++;
8796     *ps = s;
8797     return ret;
8798 }
8799

```



```

8800 int
8801 peek(char **ps, char *es, char *toks)
8802 {
8803     char *s;
8804
8805     s = *ps;
8806     while(s < es && strchr(whitespace, *s))
8807         s++;
8808     *ps = s;
8809     return *s && strchr(toks, *s);
8810 }
8811
8812 struct cmd *parseline(char**, char*);
8813 struct cmd *parsepipe(char**, char*);
8814 struct cmd *parseexec(char**, char*);
8815 struct cmd *nulterminate(struct cmd*);
8816
8817 struct cmd*
8818 parsecmd(char *s)
8819 {
8820     char *es;
8821     struct cmd *cmd;
8822
8823     es = s + strlen(s);
8824     cmd = parseline(&s, es);
8825     peek(&s, es, "");
8826     if(s != es){
8827         printf(2, "leftovers: %s\n", s);
8828         panic("syntax");
8829     }
8830     nulterminate(cmd);
8831     return cmd;
8832 }
8833
8834 struct cmd*
8835 parseline(char **ps, char *es)
8836 {
8837     struct cmd *cmd;
8838
8839     cmd = parsepipe(ps, es);
8840     while(peek(ps, es, "&")){
8841         gettoken(ps, es, 0, 0);
8842         cmd = backcmd(cmd);
8843     }
8844     if(peek(ps, es, ";")){
8845         gettoken(ps, es, 0, 0);
8846         cmd = listcmd(cmd, parseline(ps, es));
8847     }
8848     return cmd;
8849 }

```

```

8850 struct cmd*
8851 parsepipe(char **ps, char *es)
8852 {
8853     struct cmd *cmd;
8854
8855     cmd = parseexec(ps, es);
8856     if(peek(ps, es, "|")){
8857         gettoken(ps, es, 0, 0);
8858         cmd = pipecmd(cmd, parsepipe(ps, es));
8859     }
8860     return cmd;
8861 }
8862
8863 struct cmd*
8864 parseredirs(struct cmd *cmd, char **ps, char *es)
8865 {
8866     int tok;
8867     char *q, *eq;
8868
8869     while(peek(ps, es, "<>")){
8870         tok = gettoken(ps, es, 0, 0);
8871         if(gettoken(ps, es, &q, &eq) != 'a')
8872             panic("missing file for redirection");
8873         switch(tok){
8874             case '<':
8875                 cmd = redircmd(cmd, q, eq, O_RDONLY, 0);
8876                 break;
8877             case '>':
8878                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
8879                 break;
8880             case '+': // >>
8881                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
8882                 break;
8883         }
8884     }
8885     return cmd;
8886 }
8887
8888
8889
8890
8891
8892
8893
8894
8895
8896
8897
8898
8899

```

```

8900 struct cmd*
8901 parseblock(char **ps, char *es)
8902 {
8903     struct cmd *cmd;
8904
8905     if(!peek(ps, es, "("))
8906         panic("parseblock");
8907     gettoken(ps, es, 0, 0);
8908     cmd = parseline(ps, es);
8909     if(!peek(ps, es, " "))
8910         panic("syntax - missing ");
8911     gettoken(ps, es, 0, 0);
8912     cmd = parseredirs(cmd, ps, es);
8913     return cmd;
8914 }
8915
8916 struct cmd*
8917 parseexec(char **ps, char *es)
8918 {
8919     char *q, *eq;
8920     int tok, argc;
8921     struct execcmd *cmd;
8922     struct cmd *ret;
8923
8924     if(peek(ps, es, "("))
8925         return parseblock(ps, es);
8926
8927     ret = execcmd();
8928     cmd = (struct execcmd*)ret;
8929
8930     argc = 0;
8931     ret = parseredirs(ret, ps, es);
8932     while(!peek(ps, es, "|&");){
8933         if((tok=gettoken(ps, es, &q, &eq)) == 0)
8934             break;
8935         if(tok != 'a')
8936             panic("syntax");
8937         cmd->argv[argc] = q;
8938         cmd->eargv[argc] = eq;
8939         argc++;
8940         if(argc >= MAXARGS)
8941             panic("too many args");
8942         ret = parseredirs(ret, ps, es);
8943     }
8944     cmd->argv[argc] = 0;
8945     cmd->eargv[argc] = 0;
8946     return ret;
8947 }
8948
8949

```

```

8950 // NUL-terminate all the counted strings.
8951 struct cmd*
8952 nulterminate(struct cmd *cmd)
8953 {
8954     int i;
8955     struct backcmd *bcmd;
8956     struct execcmd *ecmd;
8957     struct listcmd *lcmd;
8958     struct pipecmd *pcmd;
8959     struct redircmd *rcmd;
8960
8961     if(cmd == 0)
8962         return 0;
8963
8964     switch(cmd->type){
8965     case EXEC:
8966         ecmd = (struct execcmd*)cmd;
8967         for(i=0; ecmd->argv[i]; i++)
8968             *ecmd->eargv[i] = 0;
8969         break;
8970
8971     case REDIR:
8972         rcmd = (struct redircmd*)cmd;
8973         nulterminate(rcmd->cmd);
8974         *rcmd->efile = 0;
8975         break;
8976
8977     case PIPE:
8978         pcmd = (struct pipecmd*)cmd;
8979         nulterminate(pcmd->left);
8980         nulterminate(pcmd->right);
8981         break;
8982
8983     case LIST:
8984         lcmd = (struct listcmd*)cmd;
8985         nulterminate(lcmd->left);
8986         nulterminate(lcmd->right);
8987         break;
8988
8989     case BACK:
8990         bcmd = (struct backcmd*)cmd;
8991         nulterminate(bcmd->cmd);
8992         break;
8993     }
8994     return cmd;
8995 }
8996
8997
8998
8999

```

```

9000 #include "asm.h"
9001 #include "memlayout.h"
9002 #include "mmu.h"
9003
9004 # Start the first CPU: switch to 32-bit protected mode, jump into C.
9005 # The BIOS loads this code from the first sector of the hard disk into
9006 # memory at physical address 0x7c00 and starts executing in real mode
9007 # with %cs=0 %ip=7c00.
9008
9009 .code16                # Assemble for 16-bit mode
9010 .globl start
9011 start:
9012     cli                # BIOS enabled interrupts; disable
9013
9014     # Zero data segment registers DS, ES, and SS.
9015     xorw    %ax,%ax    # Set %ax to zero
9016     movw    %ax,%ds    # -> Data Segment
9017     movw    %ax,%es    # -> Extra Segment
9018     movw    %ax,%ss    # -> Stack Segment
9019
9020     # Physical address line A20 is tied to zero so that the first PCs
9021     # with 2 MB would run software that assumed 1 MB. Undo that.
9022 seta20.1:
9023     inb     $0x64,%al    # Wait for not busy
9024     testb   $0x2,%al
9025     jnz     seta20.1
9026
9027     movb    $0xd1,%al    # 0xd1 -> port 0x64
9028     outb    %al,$0x64
9029
9030 seta20.2:
9031     inb     $0x64,%al    # Wait for not busy
9032     testb   $0x2,%al
9033     jnz     seta20.2
9034
9035     movb    $0xdf,%al    # 0xdf -> port 0x60
9036     outb    %al,$0x60
9037
9038     # Switch from real to protected mode. Use a bootstrap GDT that makes
9039     # virtual addresses map directly to physical addresses so that the
9040     # effective memory map doesn't change during the transition.
9041     lgdt    gdt desc
9042     movl    %cr0,%eax
9043     orl     $CR0_PE,%eax
9044     movl    %eax,%cr0
9045
9046     # Complete transition to 32-bit protected mode by using long jmp
9047     # to reload %cs and %eip. The segment descriptors are set up with no
9048     # translation, so that the mapping is still the identity mapping.
9049     ljmp    $(SEG_KCODE<<3), $start32

```

```

9050 .code32 # Tell assembler to generate 32-bit code now.
9051 start32:
9052     # Set up the protected-mode data segment registers
9053     movw    $(SEG_KDATA<<3), %ax    # Our data segment selector
9054     movw    %ax,%ds                # -> DS: Data Segment
9055     movw    %ax,%es                # -> ES: Extra Segment
9056     movw    %ax,%ss                # -> SS: Stack Segment
9057     movw    $0,%ax                # Zero segments not ready for use
9058     movw    %ax,%fs                # -> FS
9059     movw    %ax,%gs                # -> GS
9060
9061     # Set up the stack pointer and call into C.
9062     movl    $start,%esp
9063     call    bootmain
9064
9065     # If bootmain returns (it shouldn't), trigger a Bochs
9066     # breakpoint if running under Bochs, then loop.
9067     movw    $0x8a00,%ax            # 0x8a00 -> port 0x8a00
9068     movw    %ax,%dx
9069     outw    %ax,%dx
9070     movw    $0x8ae0,%ax            # 0x8ae0 -> port 0x8a00
9071     outw    %ax,%dx
9072 spin:
9073     jmp     spin
9074
9075 # Bootstrap GDT
9076 .p2align 2                # force 4 byte alignment
9077 gdt:
9078     SEG_NULLASM                # null seg
9079     SEG_ASM(STA_X|STA_R, 0x0, 0xffffffff) # code seg
9080     SEG_ASM(STA_W, 0x0, 0xffffffff)       # data seg
9081
9082 gdt desc:
9083     .word    (gdt desc - gdt - 1)        # sizeof(gdt) - 1
9084     .long    gdt                        # address gdt
9085
9086
9087
9088
9089
9090
9091
9092
9093
9094
9095
9096
9097
9098
9099

```

```

9100 // Boot loader.
9101 //
9102 // Part of the boot block, along with bootasm.S, which calls bootmain().
9103 // bootasm.S has put the processor into protected 32-bit mode.
9104 // bootmain() loads an ELF kernel image from the disk starting at
9105 // sector 1 and then jumps to the kernel entry routine.
9106
9107 #include "types.h"
9108 #include "elf.h"
9109 #include "x86.h"
9110 #include "memlayout.h"
9111
9112 #define SECTSIZE 512
9113
9114 void readseg(uchar*, uint, uint);
9115
9116 void
9117 bootmain(void)
9118 {
9119     struct elfhdr *elf;
9120     struct proghdr *ph, *eph;
9121     void (*entry)(void);
9122     uchar* pa;
9123
9124     elf = (struct elfhdr*)0x10000; // scratch space
9125
9126     // Read 1st page off disk
9127     readseg((uchar*)elf, 4096, 0);
9128
9129     // Is this an ELF executable?
9130     if(elf->magic != ELF_MAGIC)
9131         return; // let bootasm.S handle error
9132
9133     // Load each program segment (ignores ph flags).
9134     ph = (struct proghdr*)((uchar*)elf + elf->phoff);
9135     eph = ph + elf->phnum;
9136     for(; ph < eph; ph++){
9137         pa = (uchar*)ph->paddr;
9138         readseg(pa, ph->filesz, ph->off);
9139         if(ph->memsz > ph->filesz)
9140             stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
9141     }
9142
9143     // Call the entry point from the ELF header.
9144     // Does not return!
9145     entry = (void*)(void)(elf->entry);
9146     entry();
9147 }
9148
9149

```

```

9150 void
9151 waitdisk(void)
9152 {
9153     // Wait for disk ready.
9154     while((inb(0x1F7) & 0xC0) != 0x40)
9155         ;
9156 }
9157
9158 // Read a single sector at offset into dst.
9159 void
9160 readsect(void *dst, uint offset)
9161 {
9162     // Issue command.
9163     waitdisk();
9164     outb(0x1F2, 1); // count = 1
9165     outb(0x1F3, offset);
9166     outb(0x1F4, offset >> 8);
9167     outb(0x1F5, offset >> 16);
9168     outb(0x1F6, (offset >> 24) | 0xE0);
9169     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
9170
9171     // Read data.
9172     waitdisk();
9173     insl(0x1F0, dst, SECTSIZE/4);
9174 }
9175
9176 // Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
9177 // Might copy more than asked.
9178 void
9179 readseg(uchar* pa, uint count, uint offset)
9180 {
9181     uchar* epa;
9182
9183     epa = pa + count;
9184
9185     // Round down to sector boundary.
9186     pa -= offset % SECTSIZE;
9187
9188     // Translate from bytes to sectors; kernel starts at sector 1.
9189     offset = (offset / SECTSIZE) + 1;
9190
9191     // If this is too slow, we could read lots of sectors at a time.
9192     // We'd write more to memory than asked, but it doesn't matter --
9193     // we load in increasing order.
9194     for(; pa < epa; pa += SECTSIZE, offset++){
9195         readsect(pa, offset);
9196     }
9197
9198
9199

```

```

9200 #ifdef CS333_P4
9201 // this is an ugly series of if statements but it works
9202 void
9203 print_mode(struct stat* st)
9204 {
9205     switch (st->type) {
9206         case T_DIR: printf(1, "d"); break;
9207         case T_FILE: printf(1, "-"); break;
9208         case T_DEV: printf(1, "c"); break;
9209         default: printf(1, "?");
9210     }
9211
9212     if (st->mode.flags.u_r)
9213         printf(1, "r");
9214     else
9215         printf(1, "-");
9216
9217     if (st->mode.flags.u_w)
9218         printf(1, "w");
9219     else
9220         printf(1, "-");
9221
9222     if ((st->mode.flags.u_x) & (st->mode.flags.setuid))
9223         printf(1, "s");
9224     else if (st->mode.flags.u_x)
9225         printf(1, "x");
9226     else
9227         printf(1, "-");
9228
9229     if (st->mode.flags.g_r)
9230         printf(1, "r");
9231     else
9232         printf(1, "-");
9233
9234     if (st->mode.flags.g_w)
9235         printf(1, "w");
9236     else
9237         printf(1, "-");
9238
9239     if (st->mode.flags.g_x)
9240         printf(1, "x");
9241     else
9242         printf(1, "-");
9243
9244     if (st->mode.flags.o_r)
9245         printf(1, "r");
9246     else
9247         printf(1, "-");
9248
9249

```

```

9250     if (st->mode.flags.o_w)
9251         printf(1, "w");
9252     else
9253         printf(1, "-");
9254
9255     if (st->mode.flags.o_x)
9256         printf(1, "x");
9257     else
9258         printf(1, "-");
9259
9260     return;
9261 }
9262 #endif
9263
9264
9265
9266
9267
9268
9269
9270
9271
9272
9273
9274
9275
9276
9277
9278
9279
9280
9281
9282
9283
9284
9285
9286
9287
9288
9289
9290
9291
9292
9293
9294
9295
9296
9297
9298
9299

```

```
9300 #include "types.h"
9301 #include "user.h"
9302 #include "date.h"
9303
9304
9305 int
9306 main(int argc, char *argv[])
9307 {
9308     struct rtcdate r;
9309     if(date(&r)) {
9310         printf(2, "date failed\n");
9311         exit();
9312     }
9313     printf(1, "Current UTC time is: %d/%d/%d - %d:%d:%d\n", r.year, r.month, r.day, r.hour, r.min, r.sec);
9314
9315     exit();
9316 }
9317
9318
9319
9320
9321
9322
9323
9324
9325
9326
9327
9328
9329
9330
9331
9332
9333
9334
9335
9336
9337
9338
9339
9340
9341
9342
9343
9344
9345
9346
9347
9348
9349
```

```
9350 #define STRMAX 32
9351
9352 struct uproc {
9353     uint pid;
9354     uint uid;
9355     uint gid;
9356     uint ppid;
9357     uint elapsed_ticks;
9358     uint CPU_total_ticks;
9359     char state[STRMAX];
9360     uint size;
9361     char name[STRMAX];
9362 #ifdef CS333_P3
9363     int priority;
9364 #endif
9365 };
9366
9367
9368
9369
9370
9371
9372
9373
9374
9375
9376
9377
9378
9379
9380
9381
9382
9383
9384
9385
9386
9387
9388
9389
9390
9391
9392
9393
9394
9395
9396
9397
9398
9399
```

```

9400 #include "types.h"
9401 #include "user.h"
9402
9403 // Test GID and UID to be in the correct range
9404 #ifdef CS333_P2
9405 int
9406 testgiduid(void)
9407 {
9408     uint uid, gid, ppid;
9409
9410     uid = getuid();
9411     printf(2, "Current UID is : %d\n", uid);
9412     printf(2, "Setting UID to 100\n");
9413     setuid(100);
9414     uid = getuid();
9415     printf(2, "Current UID is : %d\n", uid);
9416
9417     gid = getgid();
9418     printf(2, "Current GID is : %d\n", gid);
9419     printf(2, "Setting GID to 100\n");
9420     setgid(100);
9421     gid = getgid();
9422     printf(2, "Current UID is : %d\n", gid);
9423
9424     ppid = getppid();
9425     printf(2, "My parent process is : %d\n", ppid);
9426     printf(2, "Done!\n");
9427
9428     return 0;
9429 }
9430
9431 int
9432 main(int argc, char *argv[])
9433 {
9434     testgiduid();
9435     exit();
9436 }
9437 #else
9438 int
9439 main(int argc, char *argv[])
9440 {
9441     printf(2, "Please compile with CS333_P2 on to enable this feature.\n");
9442     exit();
9443 }
9444 #endif
9445
9446
9447
9448
9449

```

```

9450 #include "types.h"
9451 #include "uproc.h"
9452 #include "user.h"
9453
9454 #ifdef CS333_P2
9455 int
9456 main(int argc, char *argv[])
9457 {
9458     int ptable_size;
9459     uint display_size;
9460     display_size = 64;
9461     struct uproc* ps;
9462     ps = malloc(sizeof(struct uproc) * display_size);
9463     ptable_size = getprocs(display_size, ps);
9464     if(ptable_size <= 0) {
9465         printf(1, "\nGetting processes information failed\n");
9466         exit();
9467     }
9468     printf(1, "\nNumber of processes is : %d\n", ptable_size);
9469 #ifdef CS333_P3
9470     printf(1, "\nPID      State      Name      UID      GID      PPID      Pri
9471     int i;
9472     for(i=0; i < ptable_size; ++i){
9473         printf(1, "\n%d      %s      %s      %d      %d      %d      %d      %d.%d      %d.%d",
9474             ps->state, \
9475             ps->name, \
9476             ps->uid, \
9477             ps->gid, \
9478             ps->ppid, \
9479             ps->priority, \
9480             ps->elapsed_ticks/100, ps->elapsed_ticks%100, ps->CPU_total_ticks,
9481             ++ps;
9482     }
9483 #else
9484     printf(1, "\nPID      State      Name      UID      GID      PPID      E
9485     int i;
9486     for(i=0; i < ptable_size; ++i){
9487         printf(1, "\n%d      %s      %s      %d      %d      %d      %d.%d      %d.%d      %d",
9488             ps->state, \
9489             ps->name, \
9490             ps->uid, \
9491             ps->gid, \
9492             ps->ppid, \
9493             ps->elapsed_ticks/100, ps->elapsed_ticks%100, ps->CPU_total_ticks,
9494             ++ps;
9495     }
9496 #endif
9497     free(ps);
9498     exit();
9499 }

```

```

9500 #else
9501 int
9502 main(int argc, char *argv[])
9503 {
9504     printf(2, "Please compile with CS333_P2 on to enable this feature.\n");
9505     exit();
9506 }
9507 #endif
9508
9509
9510
9511
9512
9513
9514
9515
9516
9517
9518
9519
9520
9521
9522
9523
9524
9525
9526
9527
9528
9529
9530
9531
9532
9533
9534
9535
9536
9537
9538
9539
9540
9541
9542
9543
9544
9545
9546
9547
9548
9549

```

```

9550 #include "types.h"
9551 #include "user.h"
9552
9553 #ifdef CS333_P2
9554 int
9555 main(int argc, char *argv[])
9556 {
9557     int elapsed_t = 0;
9558     int pid;
9559     int start_t = 0;
9560     int end_t = start_t;
9561     if(argc > 1) {
9562         start_t = uptime();
9563         pid = fork();
9564         if(pid > 0) {
9565             pid = wait();
9566             end_t = uptime();
9567         }
9568         else if(pid == 0) {
9569             //child process running
9570             if(exec(argv[1], argv+1) < 0)
9571                 printf(2, "%s failed to execute.", argv[1]);
9572             exit();
9573         }
9574         else {
9575             // error: fork failed
9576             printf(2, "Error: Fork failed");
9577             exit();
9578         }
9579     }
9580     elapsed_t = end_t - start_t;
9581     char *proc_name = argv[1] ? argv[1] : "";
9582     printf(1, "%s ran in %d.%d seconds\n", proc_name, elapsed_t/100, elapsed_t%100);
9583
9584     exit();
9585 }
9586 #else
9587 int
9588 main(int argc, char *argv[])
9589 {
9590     printf(2, "Please compile with CS333_P2 on to enable this feature.\n");
9591     exit();
9592 }
9593 #endif
9594
9595
9596
9597
9598
9599

```



```
9600 // This program can be freely used to test your scheduler. It is
9601 // by no means a complete test.
9602
9603 #include "types.h"
9604 #include "user.h"
9605
9606 // PrioCount should be set to the nummber of priority levels
9607 #define PrioCount 3
9608 #define numChildren 10
9609
9610 void
9611 countForever(int p)
9612 {
9613     int j;
9614     unsigned long count = 0;
9615
9616     j = getpid();
9617     p = p%PrioCount;
9618     setpriority(j, p);
9619     printf(1, "%d: start prio %d\n", j, p);
9620
9621     while (1) {
9622         count++;
9623         if ((count & 0xFFFFFFFF) == 0) {
9624             p = (p+1) % PrioCount;
9625             setpriority(j, p);
9626             printf(1, "%d: new prio %d\n", j, p);
9627         }
9628     }
9629 }
9630
9631 int
9632 main(void)
9633 {
9634     int i, rc;
9635
9636     for (i=0; i<numChildren; i++) {
9637         rc = fork();
9638         if (!rc) { // child
9639             countForever(i);
9640         }
9641     }
9642     // what the heck, let's have the parent waste time as well!
9643     countForever(1);
9644     exit();
9645 }
9646
9647
9648
9649
```